

Automating Data- and Code-Archiving Checks with Metacheck: A Replication and Extension of Cooper et al. (2026)

Daniël Lakens

Checking whether published articles archive usable data and code is time-consuming: even the lightest form of computational reproducibility checking, running an article’s own code on its own shared data, has been estimated to take around 24 minutes per article, and writing independent analysis code from scratch has been estimated at 2-25 person-hours per article. Cooper and the BES Data and Code Hackathon Group (2026) manually coded data- and code-archiving practices in 1861 papers published across the seven British Ecological Society (BES) journals between 2017 and 2024, using a two-day hackathon of 145 participants who followed a bespoke coding protocol. This manual effort is exactly the kind of task that automated tools such as Metacheck aim to support: extracting structured information about open science practices from the text and repositories linked to a scientific article, in a fraction of the time. We ran Metacheck’s `repo_check`, `data_check`, and `code_check` modules on the same corpus and compared its output against Cooper et al.’s human-coded ground truth for every construct where a direct, mechanical comparison is possible: data availability, archive location, data licensing, data and code download success, code archiving, and README presence. For each construct we report sensitivity, specificity, and two directional miss rates (how often a human coder caught something the tool missed, and vice versa), and, going beyond aggregate accuracy, we individually reviewed every disagreeing paper against Metacheck’s saved output, the paper’s own text, and, where relevant, a live call to the repository platform’s own API, to determine which side was correct and why. Agreement on whether data could be downloaded and used was 99.0%, and on whether code could be downloaded was 99.5%; agreement was lower and more variable for whether a repository could be found at all (83.9%), whether code was archived (87.4%), and whether a README was present (71.8%). Metacheck’s precision is consistently high across constructs — when it reports a repository, license, download, or README as present, that report is rarely spurious — but its recall is uneven and construct-dependent, driven by a mix of a repository-detection coverage gap (a real, working data or code citation on a platform or in a format

Metacheck does not yet recognise), a structural limitation in inspecting the contents of compressed archives, and at least one genuine definitional mismatch (platform-level description metadata a human coder reasonably reads as README-equivalent, which Metacheck’s stricter file-presence check does not). Our goal is not to replace manual coding of this kind, but to clarify, concretely and per construct, where automation can currently reduce the burden of large-scale meta-scientific data collection while maintaining comparable accuracy, and where it currently cannot, and why.

Introduction

The last two decades has seen a notable culture change around open science practices. Universities, journals, and funders increasingly expect researchers to engage in open science practices. One open science practice that has become especially widely required and adopted is data and code sharing. The goal of data and code sharing is to facilitate re-use, and to increase transparency. Funders have started to require data management plans that specify where and how data and code will be made available for re-use, and submission guidelines for journals often require that data and code underlying the claims in a manuscript are shared if possible. The British Ecological Society, for example, requires that “all data and code supporting the results in the paper must be archived in a public repository” (<https://www.britishecologicalsociety.org/publications/journals/author-guidelines/>).

Although the requirements to share data and code are clear, it is less clear who is responsible to check that authors actually follow through on these requirements. Manually checking data and code files is a time-consuming task that requires expertise about best practices in data and code sharing. Due to the diffusion of responsibility and the lack of resources to check whether scientific manuscripts follow best practices in open science, requirements to share data and code are often not followed, and if data and code are shared, best practices are often not followed. Cooper and the BES Data and Code Hackathon Group (2026b) assessed data- and code-archiving practices across the seven journals published by the British Ecological Society (BES). They found that most manuscripts in which data was used also shared the data, but code sharing was less common, and even when data and code were shared, best practices were often not followed. About a third of data and/or code repositories lacked a README file, code not always well-annotated, and code was not always reproducible (e.g., due to hard-coded file paths). These findings strongly suggest that it would be beneficial if data and code is reviewed before a manuscript is published.

Some journals manually check data and code before an article is published. This practice is most common in economics journals, where a dedicated Data Editor oversees the reproduction of results based on data and code (Vilhuber & Cavanagh (2025)). Manual checks are laborious and delay the publication process. It is therefore not surprising that researchers are exploring automated ways to check whether data and code are shared, and whether they are usable. For example, the Scalable Infrastructure for Validation of Computational Social Science Research

(SIVACOR) pilot project by the data editors of the American Economic Association examines whether it is possible to automatically verify the computational reproducibility of data and code. The extent to which automated processes can reduce the workload associated with checks of data and code sharing practices is an open question. Automation has the potential to reduce the workload of manual checks, but it is not clear whether and where automated tools can achieve the same level of accuracy as human coders.

In this paper we will examine how Metacheck, a novel open source tool that checks best practices in data and code sharing, might assist in the process of checking whether data and code are shared. We will compare the output of Metacheck against the manually coded ground truth of Cooper et al. (2026b), available at <https://github.com/nhcooper123/reproduce-reuse-recycle>, to determine where Metacheck can currently reduce the manual burden of large-scale meta-scientific data collection while maintaining comparable accuracy, and where it currently cannot, and why. It is clear the human coders exerted considerable effort over the 2-day hackathon to code the 1861 papers, where 145 people coded a median of 10 (min = 1, max = 68) manuscripts. Cooper et al.’s protocol covered the seven journals published by the British Ecological Society (BES), coding papers published between 2017 and 2024 on eight questions per paper. The team coded whether data and code were used and archived, where they were archived, whether the archived files could be located, downloaded, and opened, what file formats were used, whether a README was present and how useful it was, how complete the archived data were, how well-annotated the code was, and whether the data and code were citable (i.e., had a DOI and a licence).

This is precisely the kind of structured, repetitive, and time-consuming coding task that automated tools such as Metacheck (DeBruine et al., 2026) aim to support. If an automated tool could complete some or all of Cooper et al.’s eight coding questions at comparable accuracy to their trained human coders, it would substantially reduce the workload that is currently required to check if researchers follow best practices in data and code sharing. Metacheck (DeBruine et al., 2026) is a modular, open-source tool that uses text search, rule-based classification, API’s, and optional large language models to extract information about research practices from scientific manuscripts and their associated repositories. It contains three modules — `repo_check`, `data_check`, and `code_check` that perform automated checks that closely resemble the questions Cooper et al. (2026b) asked their human coders to answer.

In this paper, we use Cooper et al.’s openly archived data as ground truth (Cooper & the BES Data and Code Hackathon Group, 2026a), and for each of the eight questions in their protocol, we (1) summarize what Cooper et al. found, quoting their own text verbatim, and (2) examine, concretely and with reference to specific Metacheck module output fields, whether that question can currently be automated, partially automated, or not automated at all. Where automation is possible, we run the relevant Metacheck module(s) on the same 1861 papers and report the results. The results generated by Metacheck are completely automated, do not rely on AI, and can be completely reproduced by re-running the code, without any human input.

Unlike a validation study that reports a single accuracy percentage, our goal here is to make transparent, question by question, exactly where the current version of Metacheck can reduce

the manual burden of this kind of meta-scientific data collection, and where it cannot. As an automated tool, Metacheck follows FAIR principles, which state that “all research objects should be Findable, Accessible, Interoperable and Reusable (FAIR) both for machines and for people” (Wilkinson et al., 2016, p. 3). Human coders can retrieve information that is not accessible by a machine. For example, if authors state that the data is shared as supplementary material, the data is not FAIR, because a machine can’t find and access the datafile. Cooper and colleagues did not limit their coding to FAIR data, and as a consequence, Metacheck will always miss some information that human coders found. On the other hand, machines can find some information that human coders miss. For example, the Dryad data repository requires a CC0-1.0 license at deposit for every dataset with no author choice involved, and adds this information to the metadata. Metacheck retrieves this information, but the human coders were unable to locate license information by visually examining the Dryad repository. Throughout this project we will also reflect on practical challenges when developing software that attempts to automatically retrieve FAIR data, and the implications for best practices in data and code sharing.

Methods

We obtained the list of paper DOIs, journals, and years directly from the cleaned dataset in the public repository, and downloaded all articles through an R script. All articles were converted to structured text using GROBID 9.0, and converted to Metacheck paper objects. We ran the `repo_check` (locates and lists files in online repositories), `data_check` (classifies files and extracts information about tabular data), and `code_check` (checks R, Python, SAS, SPSS, Stata, Mplus, and MATLAB code files) modules to mirror to manual coding performed by Cooper and colleagues. Subsequently, we ran the `reproducibility_check` module to illustrate additional benefits of automation.

1. Does the paper use data/code? If so, are they archived?

Cooper and colleagues reported that “Of the 1861 papers in our dataset, 93% (n = 1735) used data and 90% (n = 1670) used code. Of these papers, 97% (n = 1690) archived their data, but only 35% (n = 577) archived their code”. The first two percentages are based on text in data availability statements and the method section. There are tools that automatically check data availability statements, such as OddPub (Riedel et al., 2020). Metacheck itself has an equivalent module, `open_practices`, which reads the full text and judges from the wording alone whether open data and/or open code was shared. These automated checks most closely follow the coding by Cooper and colleagues, in that data and code is classified as available irrespective of whether authors followed FAIR principles. Metacheck’s `repo_check` module takes a different approach and classifies data and code as available based on a linked repository that can be automatically accessed through an API.

Oddpub and Open Practices Check

We compared the automatic text-based judgements from OddPub and the Open Practices module to Cooper and colleagues' human coding of whether data and code were openly shared (their `data_open/code_open` columns), treating "Yes" (and for data "Yes, but not all files") as open and "No" as not open, and excluding papers where the coders used an ambiguous label (e.g., "Needs specific software or too large") or left the item unscored (`data_open/code_open` were only asked for papers Cooper and colleagues had already coded as archiving their data/code, and not every such paper has a recorded answer, so the denominators below are smaller than the archiving counts reported earlier). For data sharing (1552 papers with a clear human judgement), OddPub identified 80% of the papers the human coders judged to share open data (sensitivity), and correctly identified 45% of the papers judged not to (specificity), for an overall agreement of 80%. Metacheck's `open_practices` module agreed with the human coders on data sharing 86% of the time, with higher sensitivity (86%) but lower specificity (9%) than OddPub. For code sharing (498 papers with a clear human judgement), the difference was larger: OddPub agreed with the human coders 44% of the time (sensitivity 44%, specificity 67%), whereas `open_practices` agreed 78% of the time (sensitivity 78%, specificity 0%). Both tools, in other words, are much more likely to say a paper shares open data or code than to say it does not; this corpus contains very few papers the human coders judged as clearly not sharing (11 for data, 3 for code), so the specificity estimates above are based on very small numbers and should be treated as indicative rather than precise. `open_practices` and OddPub both determine open data/code from the text of the article alone, and neither distinguishes between data and code repositories the way `repo_check` and `code_check` do below.

Note that OddPub and `open_practices` classify every paper regardless of how Cooper and colleagues' coders scored this question. For data sharing, 97 papers archived their data (`data_availability == "Yes"`) but were left unscored on whether it could be opened, of which OddPub judged 44 (45%) and `open_practices` judged 80 (82%) to have open data; a further 41 were explicitly flagged by Cooper and colleagues' coders as needing specific software or too large to open, of which OddPub judged 24 (59%) and `open_practices` judged 31 (76%) to have open data. For code sharing, 49 papers archived their code (`code_archived == "Yes"`) but were left unscored on whether it could be opened, of which OddPub judged 8 (16%) and `open_practices` judged 25 (51%) to have open code; a further 34 were explicitly flagged as possibly accessible with the right software, of which OddPub judged 16 (47%) and `open_practices` judged 24 (71%) to have open code.

Repo check, Data check, and Code check

Cooper and colleagues found that 1690 papers archived their data and 577 their code, and in total, 1789 papers contain a link to a repository with either data or code. In total, Metacheck found 1545 papers that contained a link to a repository in the main text where it could verify that the repository contained files. Of these, 1452 papers contained data, 437 papers contained code, and 1506 papers contained either data or code according to `repo_check`. The manually

coded classifications agree with the automatic classification by Metacheck for 1489 papers, and disagree for 317 (300 where Cooper’s coders found data or code that Metacheck missed, and 17 papers where Metacheck found data or code while Cooper’s coders missed it).

For Metacheck, data and code is findable and accessible if it is archived in one of the repositories that Metacheck supports (through its support for Dataverse, DSpace 6 and 7, Figshare, and DataOne, and dedicated support for OSF, Zenodo, ResearchBox, Dryad, ReShare, GitHub, GitLab, Mendeley Data, FSD, AsPredicted, and 4TU). Because many repositories use the same underlying API, Metacheck supports over 276 unique repositories (121 Dataverse installations, 64 DSpace 7 installations, 6 legacy DSpace 6 installations, 5 DataOne member nodes, and 69 Figshare-based hosts, plus 11 dedicated single-platform integrations), but this is only a subset of all repositories used in manuscripts.

Automated tools need to be able to sustainably support existing repositories. Registries of data repositories exist (e.g., <https://www.re3data.org>) but the registry does not contain information about the APIs used by each repository. Some repositories disappear, and new ones are added. Support is relatively straightforward if a repository uses an established API (e.g., Dataverse, Figshare, DSpace) but requires more effort if it uses a unique API (e.g., OSF). This raises the question whether the definition of ‘findable’ data should incorporate the effort required to support a repository in automated tools. It seems likely that data stored at more widely used repositories (e.g., Dryad, Dataverse) will be more findable by automated tools than a small custom data repository used by a handful of scientists, simply because of the costs associated with supporting all data repositories. Of course, data repositories can facilitate the findability of the data they host by contributing to open source software, and contribute code that supports their own platform.

We individually reviewed every one of the 300 papers where Cooper and colleagues’ coders found data or code that Metacheck missed, reading each paper’s own data-availability text, checking Metacheck’s own intermediate output for that paper, reading the relevant Metacheck source code, and in most cases independently querying the cited repository’s own live API or web page to confirm what it actually contains. A first review pass classified a subset of these papers as already correctly handled by Metacheck’s current code, reasoning from a `metacheck_any_detector_hit` flag in Metacheck’s own saved output without separately re-confirming, against the live pipeline data, that files were actually being found; a second, independent re-check of every paper in that subset (described below) found this flag alone was not a reliable signal, so every one of those papers was individually re-verified from scratch rather than taken on trust. This finds six clusters of explanation, which together account for all 300 papers (Table 1): 18 papers cite no repository, dataset, or accession number of any kind anywhere in their own text, most often because the data is explicitly withheld for confidentiality, ethical, or Indigenous-data-sovereignty reasons; 2 papers’ repositories genuinely contain nothing that should count as data or code; 68 are not machine-findable at all regardless of which tool is used, because the repository itself does not meet basic FAIR requirements; 89 cite a repository that is itself fully FAIR – a persistent identifier, a documented API, live public access – but on a platform Metacheck has no code to query at

all; and the remaining 123 are papers whose repository Metacheck could, in principle, have found and classified correctly, but a specific, identifiable defect in Metacheck’s own code prevented it from doing so. The full per-paper record for every one of these 300 papers – each independently-found link, live access and license findings, and the reasoning behind each classification – is in `code/03_comparing_results/manual_repo_findability_audit.csv` and `manual_fair_api_findability_audit.csv`, so any individual classification can be checked and disputed.

Table 1: Table 1. All papers where Cooper et al.’s coders found data or code that Metacheck missed, grouped by why.

Cluster	n papers
No repository is cited anywhere in the paper’s own text	18
Repository contains no genuine data or code	2
Repository is not FAIR (unsupported platform with no API, or not actually live/accessible)	68
Repository is FAIR, but on a platform Metacheck has no code for	89
Metacheck limitation: the repository is FAIR and on a supported platform, but a specific defect in Metacheck’s own code prevented detection or classification	123

No repository is cited anywhere in the paper’s own text (18 papers). For these papers, a full-text search found no repository name, DOI, URL, or accession number of any kind connected to the paper’s own data or code – there is no machine-findable identifier here for any automated tool, however complete its platform coverage, to act on. The overwhelming majority state an explicit reason data is not archived: participant confidentiality or informed-consent terms, Indigenous data sovereignty (for instance, cultural and language data that the source community’s own protocols require stay outside any public repository), a platform’s terms of service prohibiting redistribution, or data held by a third-party organisation available only on direct request. A smaller number instead give the data directly within the manuscript itself (in a table, or as a simulation study’s parameter values) or point only to a static, non-indexable file attached to the publisher’s own supporting-information page, rather than to any external, independently queryable repository. This cluster reflects a genuine disagreement with Cooper and colleagues’ own coding for these particular papers, not a Metacheck limitation of any kind, since no repository exists for Metacheck – or any other tool – to find.

Most of these 18 papers, together with every paper in the 123-paper Metacheck-limitation cluster discussed below, were originally classified under a different label during a first review pass, on the strength of a `metacheck_any_detector_hit` flag already recorded in Metacheck’s saved output rather than a fresh, independent check of the live pipeline data. A second, systematic re-check compared that flag against `res_repo_check.RData`’s actual stored `repo_n/files_n` counts for every paper in that first-pass category and found the flag alone was not reliable: for a number

of papers it no longer matched the pipeline’s current state at all (`repo_n` showed zero where the flag claimed a detector hit), and for others the flag was accurate but a downstream step still failed silently, so no files were ever actually enumerated or classified despite the repository itself being correctly found. Every paper in that original category was therefore individually re-verified from scratch – rereading the paper’s text, rerunning the relevant check live and uncached where needed, and reading Metacheck’s own source code – before being reassigned to whichever cluster its real, current behaviour actually supports. One paper, once rerun live and uncached, was confirmed to genuinely resolve correctly and was removed from the disagreement population entirely, matching the treatment given to every other genuinely-resolved paper in this review (see “A targeted rerun” in the Appendix); none of the rest did.

The repository genuinely contains no data or code (2 papers). For one paper (10.1111/1365-2745.13026), the cited Zenodo deposit, when downloaded and inspected directly, contains only a license file, a README, and a supplementary PDF of figures and prose – no raw dataset or analysis code at all, despite being badged as a GitHub software release. Metacheck’s classification of this deposit as containing no data or code is, on inspection, correct; the disagreement likely reflects Cooper and colleagues’ coders crediting the paper for a well-presented, DOI-archived deposit without checking whether its actual contents meet the bar. For a second paper (10.1111/2041-210X.13396), the repository’s only file is a single Word document whose content we could not verify directly (the repository’s download endpoint required authentication we did not have), making this a genuinely ambiguous case rather than a clear miss on either side.

Repositories that are not FAIR (68 papers). These are not Metacheck limitations: no automated tool, however complete its platform coverage, could retrieve this data. 14 repositories do not actually resolve to a public, accessible resource when checked live right now – most commonly a stated embargo, confirmed by querying the repository’s own API directly and receiving an explicit “cannot be viewed” response rather than a Metacheck-side error. The remaining 54 cite a platform Metacheck does not support, and on checking that platform’s own website and documentation directly, it does not expose a real persistent identifier, a documented machine-queryable API, or both – for instance, a personal website, an institutional page with no API, or a private, unversioned share-link.

FAIR repositories on an unsupported platform (89 papers). These papers cite a repository that is itself fully FAIR by this manuscript’s own definition – a real persistent identifier, a documented API, and live public access, all independently confirmed – but Metacheck simply has no code written to query that specific platform. This is the largest single cluster in the full disagreement population, and represents the clearest, most actionable item for future development: each of these platforms, once a dedicated integration is written, should become fully findable with no further conceptual work. Several platforms recur often enough across these 89 papers to be worth a dedicated integration on their own (for example, the US Geological Survey’s ScienceBase, which exposes a documented REST API and official client libraries; the Environmental Data Initiative; the NERC Environmental Information Data Centre; and PANGAEA, the earth- and environmental-science data publisher), while

the remainder are spread across many platforms that each appeared only once or twice in this corpus, for which a dedicated integration would fix one paper at a time rather than a whole recurring class of citation.

Metacheck limitations (123 papers). These are the papers where the repository is FAIR, is on a platform Metacheck does support, and the only thing standing between Metacheck and a correct answer is a specific, identifiable defect in Metacheck’s own code – the clearest and most directly actionable cluster in this review, and the one we document in the most detail. It splits into three groups by what kind of defect is involved.

The first group (85 papers) never got as far as `repo_check` finding a repository at all. 29 cite a repository on a platform Metacheck’s code already has a working integration for, but this specific institutional host or DOI prefix was not yet registered in that integration’s list of known hosts (for example, an institutional Dataverse or Figshare installation whose domain name Metacheck’s code did not yet recognise). We individually live-checked every one of these 29 repositories’ actual API endpoint rather than assuming the fix is uniformly simple: 21 resolve to a single, confirmed host or DOI-prefix table entry with no remaining ambiguity, verified directly against the repository’s own API; the other 8 remain genuinely open, because the host could not be reached (blocked by bot protection or timed out) or because closing the gap needs a code change – widening which part of the document a detector searches, or resolving a redirect before matching – rather than a one-line table addition. 26 cite a platform Metacheck supports where a specific, now-diagnosed defect prevented the citation from being recognised or its files from being retrieved even though the citation text was clean and the repository itself is live and public; these defects are not all the same shape – some are in the detection regex itself (for example, no pattern matches Dryad’s own “stash/share” reviewer-link URL format, or a Figshare “private share” link whose opaque token carries no article ID that any third party, Metacheck included, can resolve without a browser), two are a self-hosted GitLab instance that Metacheck’s GitLab detector does not match by explicit design (it is scoped to `gitlab.com` only), one is a case-sensitivity mismatch between a citation’s own capitalisation and what the repository’s API will accept, and one cites a DSpace installation running a version old enough that it exposes neither of the two REST API shapes Metacheck’s two DSpace integrations know how to query, so adding the host would not by itself fix it. A further 2 cite a bare GitHub organisation or user page rather than a specific repository – the page itself is correctly found, but it is not a single repository with an enumerable file tree, so there is nothing for a file listing to return; this is a distinct limitation from a missing host or a detection bug, since no configuration change would resolve it. 22 have a clean, correct citation to a repository Metacheck supports, but the version of the text Metacheck actually receives has already had the identifier corrupted before Metacheck’s own code ever sees it – most often a stray space inserted mid-identifier (for example “`https://doi.org/...`” or “`10.5061/dryad.745 file://localhost/opt/grobid/.../g830`”, where an artifact of the PDF-to-text conversion step landed in the middle of the DOI) by the PDF-to-text conversion step that runs before Metacheck. In every one of these cases we checked, the correct, uncorrupted identifier is present in the paper’s own reference list, and simply widening the text Metacheck searches to include the reference section (which it currently excludes by default) recovers the correct repository for all of them – a fix we verified directly

and have reported to Metacheck’s maintainers. 6 have no trace of the repository link anywhere in the text Metacheck receives at all, which we could not resolve further without access to the original PDF.

While investigating the 26 platform-bug cases above, we found a defect with a larger footprint than any single paper in this review: Metacheck’s Dataverse, Figshare, DataONE, and 4TU integrations each retrieve file listings for every citation to that platform in one batched API call per corpus-wide run, and an error on any single record in that call is recorded against every repository in the batch, not only the record that caused it. Checking this directly against a full corpus run, 200 of the 306 repository records Metacheck’s own code already flags as needing closer inspection carried this same error. Re-querying each of these 200 records independently, outside the original run, resolved 197 of them immediately with no code changes, consistent with a transient cause (a timeout or rate limit on one request inside a large batch) rather than a reproducible parsing defect. The remaining 3 were a separate, unrelated, and already-fixed bug: all three cited a Figshare collection DOI that a Dataverse/Figshare DOI-prefix conflict had been misrouting, fixed in Metacheck’s `dev` branch shortly before this review; two of the three belonged to papers in this review’s own disagreement population and, once rerun against the fixed version, resolved correctly and are no longer counted among the 300 papers above. We have reported the batch error-isolation defect to Metacheck’s maintainers.

The second group (36 papers) is the reverse pattern: `repo_check` correctly found the repository and listed its files, but neither `data_check` nor `code_check` classified any of those files as data or code. We inspected every one of these papers’ actual file listings directly, and in nearly every case independently confirmed, by downloading the real repository, that the unclassified files genuinely are the paper’s data or code. The single largest reason (22 papers) is that the file’s extension is simply not in Metacheck’s lookup table at all, so it falls through every classification rule to an “unknown” label – we found this for Dataverse’s own standard tabular-data export format, R’s `dput`-based plain-text data format (distinct from the already-recognised binary save-file format), the STRUcTURE population-genetics software’s input format, the DTAG marine-mammal biologging sensor format, and the NetLogo agent-based-modelling language’s source-code format, among others; the full list of every affected paper and extension is in Appendix A.10’s Table 4. A further 8 have a recognised extension that is assigned to the wrong category outright – for instance, five files saved with R’s documentation-markup extension were actually gzip-compressed data objects, which Metacheck’s classifier reads as code because it conflates the two uses of that extension; a GIS raster image format is read as ordinary photographic material rather than data, a rule that is reasonable for a stimulus photograph but wrong whenever the image itself is the scientific measurement, as it is for every GIS raster in this corpus. 1 are compressed archives that Metacheck’s code is fully able to open and classify the contents of, and does so successfully elsewhere in the corpus, but for these specific files a configuration mismatch between different parts of the pipeline meant the archive was never actually downloaded and opened in the first place, leaving it permanently labelled as an unopened container rather than classified by its contents. The remaining 5 are cases where `repo_check` found the repository but its file listing stopped at the top level and never walked into subfolders – most often an R package shared as its own GitHub repository,

where the package’s actual source code sits in a subfolder that the listing never reached, even though Metacheck’s own classification rule for that file type is correct and would have worked immediately had the files been listed at all.

The third group (2 papers) is neither a detection nor a classification defect, but a caching defect: Metacheck’s on-disk cache of a repository’s file listing can outlive the thing it is a cache of. For one paper, an earlier cached lookup had itself failed, but the failure response’s internal shape meant Metacheck’s own “is this cache entry usable” check read it as a valid, cacheable success rather than an error to retry – so every later run kept replaying the same failed lookup instead of trying again. For the other, the cached entry was written correctly for the code that existed at the time, but a later bug fix to Metacheck’s own file-parsing logic for that platform has not yet invalidated the stale pre-fix entry still sitting in the cache. Rerunning `repo_check()` for both papers with caching disabled resolves them immediately, confirming the underlying detection and retrieval logic is correct and the cache itself is the only thing standing in the way.

We have reported every one of these Metacheck limitations to Metacheck’s maintainers, with the specific file, function, and line number involved in each case, so that they can be fixed upstream.

2. Where were the data/code archived?

When exploring where data was archived, Cooper and colleagues coded the specific repository where data and/or code was stored. They observed that some data stored on GitHub was also stored on Zenodo, in line with best practices as GitHub is not a stable data or code repository.

Cooper et al. report Dryad as the dominant data repository (57%, $n = 1022$), with Figshare (9%), Zenodo (12%), and “various institutional and government repositories” (17%) making up most of the remainder (Cooper & the BES Data and Code Hackathon Group, 2026b, p. 1959). `mc_data_archive` and Cooper et al.’s `data_archive` were both codeable for 1385 papers; because this is a platform-name comparison rather than a boolean, we score agreement as at least one overlapping platform name between the two sides’ (possibly multi-valued) answers. They agreed on 1337 (96.5%), disagreed on 48 (3.5%).

Reviewing all 48 individually: 31 were `METACHECK_RIGHT` — almost entirely Cooper et al.’s coders using the catch-all “Other repo/database” label for a repository whose platform is, on the paper’s own text and URL, unambiguous (most often OSF, e.g. a Data Availability Statement reading “Data and code available from the Open Science Framework: <https://doi.org/10.17605/OSF.IO/PNJFW>,” coded “Other repo/database” rather than “OSF”). 7 were `COOPER_RIGHT`, the mirror image of a pattern that recurs across several columns in this comparison: a paper genuinely has two separate repositories, one holding code and one holding data, and `mc_data_archive` picked up the code repository’s platform because `repo_check` found that repository’s files but never located the data repository’s at all (a domain-specific

database, an institutional repository, or a Dryad/Mendeley deposit named in the text but never fetched). 7 were `BOTH_DEFENSIBLE`: genuinely multi-repository papers where Cooper et al.'s single-choice coding and Metacheck's semicolon-joined multi-platform answer are each a defensible summary of a paper that does not reduce to one archive — several of these distinguish a versioned Zenodo archive from a live, mutable GitHub development repository, with Cooper et al. and Metacheck picking different repos by design. The remaining 2 could not be resolved from the available evidence.

Limitation. The recurring code-repository-mislabelled-as-data-repository pattern reflects a genuine architectural gap rather than a text-parsing failure in any single instance: `repo_check` finds every repository a paper links, but the logic that decides which one to report as “the data archive” for `mc_data_archive` has no signal for “this repository actually holds the study’s data” beyond checking which repositories hold `data_type == "data"` files — and that check is itself blind to a repository whose only file is an unopened archive (see the `code_archived` discussion below for the same limitation’s larger effect there).

mc_data_availability. TRUE if `repo_check`’s `summary_table$files_n` is greater than zero for a paper — that is, if `repo_check` found at least one linked repository whose file listing it could retrieve, with at least one file in it, of any type (data, code, documentation, or unclassified). This rule is deliberately permissive: it answers “did we find a repository with something in it,” not “did we find a repository holding a dataset.” That distinction matters in practice (see Results below): a paper whose only linked repository is a pure-code Zenodo archive, with zero data-classified files, still scores `mc_data_availability = TRUE` under this rule, because the rule does not check `data_type`.

mc_data_archive. The archiving platform, derived by matching each linked repository’s `repo_url` against a fixed set of domain and DOI-prefix patterns: `osf.io` for OSF; `github.com/gitlab.com` for “GitHub, GitLab, Codeberg or similar platform”; `zenodo.org` or the `10.5281/zenodo` DOI prefix for Zenodo; `datadryad.org` or the `10.5061/dryad` DOI prefix for Dryad; `figshare.com` or the `10.6084/m9.figshare` DOI prefix for Figshare; anything else is labelled “Other repo/database.” When a paper cites more than one repository holding data-classified files, `mc_data_archive` lists every matching platform, semicolon-joined; Cooper et al.’s protocol instead asks for a single platform per paper, so a genuine multi-repository paper can show a “disagreement” that is really a completeness mismatch rather than a wrong answer (see Results).

mc_data_doi, mc_data_license. Read from `repo_check$repo_metadata`, a table populated only for the platforms whose API `repo_check` queries for metadata beyond the file listing itself (Zenodo, Dryad, Figshare, GitHub’s repository object, and — when `osf_license = TRUE` is set, as it was for every run reported here — OSF). `mc_data_license` is TRUE if at least one repository holding a data-classified file has a non-missing `license` value in this table; `mc_data_doi` is that repository’s own `doi` field. This scoping choice — checking only repositories that hold data — is deliberate and matches `mc_data_archive`’s own scoping, but it means a paper whose licensed repository holds no data-classified files (a common pattern when

data and code sit in physically separate repositories; see Results) scores `mc_data_license = FALSE` even though a real, retrievable license genuinely exists for that paper’s overall archive.

`repo_metadata$license` is read verbatim from each platform’s own API response, with no rewriting or standardisation on Metacheck’s side — this means its format varies by platform (a bare SPDX identifier like "MIT" or "GPL-3.0", a human-readable string like "CC BY 4.0", or a full SPDX URL), and it can be identical across many repositories on the same platform when that platform’s own licensing policy is uniform, not because Metacheck is returning a cached or default value. Concretely: 1060 of the 1060 licensed Dryad repositories in this corpus (100.0%) carry the identical string, <https://spdx.org/licenses/CC0-1.0.html>, because Dryad requires CC0-1.0 at deposit for every dataset with no author choice involved; this value never appears for a non-Dryad repository in this corpus (0 occurrences among Zenodo, Figshare, GitHub, and OSF repositories combined), whose own license values vary appropriately by repository and platform.

mc_data_download. TRUE if at least one data-classified file has `data_check`’s own `tabular_usable == TRUE` **and** a non-missing `file_location` (the local path the file was actually fetched to). This is a two-part requirement, not download success alone: a file can download successfully and still fail this check if `data_check` judges its content unusable as a table (see the Excel multi-sheet limitation in Results). `mc_data_format` is the sorted, semicolon-joined set of file extensions among the files that passed this same `tabular_usable` check.

mc_any_readme. TRUE if `repo_check` classified at least one file, in any of the paper’s linked repositories, with `doc_role == "readme"` — a purely filename-based match (a base name starting with "readme", case-insensitive, or a package-level `ro-crate-metadata.json`), never a check of the file’s content. Critically, Metacheck records only **one** README signal per paper, computed across all of that paper’s repositories combined; it does not distinguish a README documenting data from one documenting code. Cooper et al.’s protocol, by contrast, asks two separate questions (`data_README`, `code_README`). We therefore compare against `cooper_any_readme`, defined as TRUE if Cooper’s coders answered “Yes” or “Quasi-README” to either of their two README questions for that paper — the only comparison that avoids double-counting the same single Metacheck signal against two different Cooper columns (an earlier, columnwise comparison inflated the disagreement count by 41% purely from this double-counting, before we corrected it; see Results).

mc_code_archived. TRUE if `code_check`’s file table contains at least one row for the paper at all — that is, if `code_check` classified at least one file (of any language) as code, regardless of whether that file was ever successfully downloaded. `mc_code_download` is the stricter follow-up: TRUE if at least one of those code-classified files additionally has a non-missing `file_location`. `mc_code_language` is the sorted, semicolon-joined set of `code_check$table$language` values among a paper’s code files; language is assigned by `code_lang()`, which is content-aware rather than purely extension-based for polyglot formats (`.Rmd/.qmd` and `.ipynb` are opened and their code cells inspected; an ambiguous `.html` is content-sniffed).

mc_code_used. Always NA. No Metacheck module answers “did this paper use code” as a construct distinct from “is there a code file sitting in a linked repository” — the two questions collapse into one from Metacheck’s perspective, since it has no way to determine from the manuscript text alone whether custom code was part of the reported methodology.

Columns Cooper et al. coded that this comparison does not attempt. Cooper et al.’s protocol also includes README/code usefulness ratings (a 1-10 human-judged scale), a four-point data completeness judgment, and a code annotation quality score. None of these has a Metacheck counterpart to compare against: `code_check`’s `percentage_comment` field measures comment density (the proportion of lines that are comments), a real and precisely computed metric, but not a judgment of whether those comments are actually informative, so we do not force it into a comparison against Cooper et al.’s quality scale here. These three constructs are discussed qualitatively, not quantitatively, in the Discussion.

Empirical Comparison

We compare Metacheck’s `mc_`-prefixed columns against Cooper et al.’s own coded ground truth across the 1861 papers we were able to retrieve, for every construct where a genuine boolean or categorical comparison is possible: whether data are available at all, where data are archived, whether a data license is present, whether data download and code archiving, and README presence (collapsed across data and code, per the Coding section above). For each construct we report: (1) how many papers both Cooper et al. and Metacheck coded a value for (`n_both`, since neither side coded every paper — Cooper et al. did not code a “not applicable” paper on every sub-question, and Metacheck cannot code a paper whose PDF we could not retrieve or convert), (2) a 2x2 confusion table of agreement and disagreement, (3) sensitivity (of the papers Cooper et al. coded **Yes**, the percentage Metacheck also coded **Yes**), specificity (of the papers Cooper et al. coded **No**, the percentage Metacheck also coded **No**), and McNemar’s test for the asymmetry of disagreement, (4) the two directional miss rates the reader is likely to care about most in practice: of the disagreements, what percentage are cases where Cooper et al. found something Metacheck missed, and what percentage are cases where Metacheck found something Cooper et al.’s human coders missed, and (5) a qualitative synthesis of every individually verified disagreement, drawn from a case-by-case review in which every one of the 789 reviewed disagreeing papers (of 819 total; see the note on the 30 not-yet-reviewed rows below) was checked directly against Metacheck’s own saved module output, the paper’s own extracted full text, and — where a repository platform’s API was the evidence source — a live API call at the time of review, and assigned one of five verdicts: `COOPER_RIGHT` (Metacheck’s answer does not match the paper/repository), `METACHECK_RIGHT` (Cooper et al.’s coded answer does not match the paper/repository), `BOTH_DEFENSIBLE` (the paper genuinely supports both answers, most often because it cites more than one repository), `DIFFERENT_DEFINITION` (the two sides are answering different questions that happen to share a column label — see the README discussion below), or `UNCLEAR` (the available evidence does not resolve the disagreement either way).

Following a full corpus rerun against a further round of fixes (Appendix A.1), the disagreement set itself was rebuilt from scratch rather than assumed unchanged: `data/cooper_vs_recreated.RData` was regenerated from the fresh `repo_check/data_check/code_check` output, and a new disagreement worklist was derived from it using the same mechanical flagging rule as before. For every (paper, column) disagreement whose Cooper-coded value and Metacheck-derived value were both unchanged from the prior review round, the earlier verdict and its evidence were carried forward rather than re-derived. Every disagreement that was new, or whose Metacheck-derived value had changed as a direct result of the fixes, was reviewed from scratch, following the same evidence-source rule throughout: `data_availability`, `data_archive`, and `code_used` were checked against the paper’s own extracted text, since these are what Cooper et al.’s coders themselves read; `data_license`, `data_download`, `code_archived`, `code_download`, and `any_readme` were checked against Metacheck’s own saved module output (`repo_check`’s file listing and `doc_role` classification, `repo_metadata`’s license/DOI fields, `file_location`), since a paper’s own text essentially never states whether a README exists or a download succeeded. A live platform API call (Dryad’s dataset API, Zenodo’s records API, Figshare’s articles API, GitHub’s repository API) was used only where local evidence was insufficient to decide. Appendix A.1 documents the specific software defects this case-by-case process surfaced and fixed over the course of the project.

The numbers reported in this section reflect a SECOND such rebuild, following Appendix A.5’s seven further fixes — applied and reverified this time against the FULL 1,861-paper corpus (not just the previously-disagreeing subset A.5 itself checked), using the identical carry-forward procedure: `code_data_preparation/57_rebuild_worklist_carry_forward_post_a5.R` re-generated `cooper_vs_recreated.RData` from a fresh `repo_check/data_check/code_check` rerun of every paper, rebuilt the disagreement worklist, and carried forward each prior verdict wherever its underlying (paper, column, Cooper value, Metacheck value) triple was unchanged. 789 of 819 disagreements carried forward or were otherwise already resolved this way; the remaining 57 papers (87 individual construct-level disagreements, since a paper can be flagged on more than one column) were newly surfaced or changed by A.5’s fixes and were reviewed from scratch the same way as every earlier round, following the same evidence-source rule described above (`code_data_preparation/59_apply_a5_review_verdicts.R`, with the full per-row reasoning in `code_data_preparation/verdicts_a5_review.csv`). Every disagreement in the corpus — 819 of 819 — now carries a verdict.

This final round of review surfaced one further genuine software defect — and, unlike every other defect this project fixed and moved on from, this one was caught, root-caused, fixed, and reverified within the same review cycle rather than left open. Appendix A.5’s R-package-source-tree exclusion (intended to drop a bundled dependency’s own source tree — `DESCRIPTION/NAMESPACE/R/man/tests/` — from a repository’s data/code listing, since that is reusable infrastructure rather than the study’s own analysis code) turned out to also strip a repository’s real analysis code whenever the paper’s own code was itself organised as an installable R package, which is common practice for ecological modelling and analysis tools shared alongside a methods paper. Confirmed across 16 real GitHub/CRAN-hosted repositories (RevGadgets, BaBA, m2b, RRPP, actel, scampr, MINOTAUR, LadderFuelsR,

biomonitoR, EcoEnsemble, fossilsim, frair, tRophicPosition, sabinaNSDM, bef, and phylomice), `repo_check`'s file listing for each package's repository contained only its build/configuration/documentation scaffolding (`.Rbuildignore`, `.gitignore`, `NEWS.md`, `README.md`, `cran-comments.md`, a bare `.Rproj` file) with the package's actual R/ source directory entirely absent — and, in several of these repositories, the package's own bundled `data/` directory was stripped along with it, silently turning a genuine data deposit into an apparent absence of data rather than merely a wrong-but-visible code disagreement.

The fix distinguishes two structurally different cases that both fingerprint identically (a `DESCRIPTION+NAMESPACE` package at a repository's root): a study that shares its own data and analysis code as an R package (BaBa's tracked-animal datasets and behaviour-classification functions, for instance) versus a repository that happens to bundle a vendored copy of someone else's general-purpose package alongside the study's own files (confirmed against MarshallLab/MGDrive, a general-purpose mosquito gene-drive simulation package whose paper-specific fork nests the vendored package one level deep, with the paper's own `Examples/` scripts and data sitting as siblings at the repository root). The corrected rule only excludes a root-level package's files when nothing classifiable as real data or code survives outside it (checked with `data_classify_files()`, the same rules-based classifier `data_check` itself uses, so “real content” means the same thing everywhere in the pipeline); a package detected in a nested subfolder — the vendored-dependency shape — is still always excluded, since by construction there is always something else in the repository for a nested package to sit alongside. Reapplying this corrected rule to the 38 repositories flagged by a corpus-wide scan for the exclusion's symptom (a repository whose only surviving files, after the old exclusion ran, were scaffolding) resolved 11 of the 16 confirmed cases; 5 remain flagged, but for an unrelated, pre-existing reason — a handful of these repositories' cached downloads never retrieved the package's R/ source directory at all (a fetch gap, not an exclusion-logic problem), so the corrected rule correctly finds nothing outside the package to preserve because nothing outside the package was ever downloaded in the first place. That fetch gap is a separate, still-open limitation, tracked apart from the exclusion fix itself.

Is any data available at all?

Cooper et al. report that “97% ($n = 1690$) archived their data” among the 1735 papers coded as using data (Cooper & the BES Data and Code Hackathon Group, 2026b, p. 1958). `mc_data_availability` and Cooper et al.'s `data_availability` were both codeable for 1735 papers. They agreed on 1469 (84.7%), disagreed on 266 (15.3%). Sensitivity was 85.7% and specificity was 44.4%, McNemar's $\chi^2(1) = 173.78$, $p < .001$. This is a considerably more balanced pattern than an earlier pass of this corpus found, because `mc_data_availability`'s own derivation was corrected during this project (see Coding section): it now reflects `repo_check` finding at least one linked repository with at least one file, rather than defaulting to a missing value whenever `repo_check` found none at all — so a paper with no detectable repository now

correctly scores `FALSE` and enters this comparison, rather than being silently excluded from it.

Reviewing all 266 disagreements individually: 230 were `COOPER_RIGHT`, 8 `METACHECK_RIGHT`, 23 `DIFFERENT_DEFINITION`, and 10 `UNCLEAR`. The dominant cause among the `COOPER_RIGHT` cases (192 of 230) is `repo_check` never detecting a repository at all, even though the paper’s own text names one explicitly — most often a real, working DOI or URL on a platform `repo_check` has no dedicated backend for (institutional repositories such as DataSTORRE, DIGITAL.CSIC, or the Australian Antarctic Data Centre; domain-specific databases such as NCBI’s Sequence Read Archive/GenBank, the Knowledge Network for Biocomplexity, or Neotoma), or occasionally a citation malformed by PDF text extraction. A further 3 trace to a `repo_check` processing error logged internally (an R-level error raised while fetching an otherwise-correctly-detected repository, or the host platform being transiently unreachable during the corpus run) rather than a missed citation — a genuine, fixable software defect distinct from a coverage gap. Conversely, 8 disagreements were `METACHECK_RIGHT`: clear, working Dryad, Zenodo, Figshare, or GitHub data-availability statements that Cooper et al.’s coders appear to have missed or coded `No` in error. 23 were `DIFFERENT_DEFINITION` — most commonly a paper stating data are available “on request,” embedded directly in the article’s own Supporting Information, or withheld for confidentiality reasons, which Cooper et al.’s protocol appears to credit as “available” while `mc_data_availability` structurally cannot, since there is no externally checkable repository for `repo_check` to find. The remaining 10 could not be resolved from the available evidence.

Limitation. The mechanism responsible for `mc_data_availability`’s disagreements has changed since an earlier pass of this corpus: rather than a single, blanket rule failure that made a `FALSE` structurally impossible, the remaining gap is now dominated by `repo_check`’s coverage of repository platforms and citation formats — addressable by extending its URL-recognition patterns to more institutional and domain-specific repositories, rather than by changing the derivation rule itself.

Is the data licensed?

Cooper et al. report that 85% ($n = 1438$) of papers with archived data had a license (Cooper & the BES Data and Code Hackathon Group, 2026b, p. 1960). `mc_data_license` and Cooper et al.’s `data_license` were both codeable for 1306 papers. They agreed on 1248 (95.6%), disagreed on 58 (4.4%). Sensitivity was 99.1% and specificity was 25.4%, McNemar’s $\chi^2(1) = 21.12$, $p < .001$.

All 58 disagreements were individually resolved against the relevant repository platform’s own API, called live at the time of review. 45 were `METACHECK_RIGHT`: a real, retrievable license genuinely exists — most commonly CC0-1.0 for Dryad datasets (Dryad requires a license at deposit) or a depositor-selected “Other (Open)” tag for Zenodo records, or a GitHub repository license confirmed live via GitHub’s own repository API. 9 were `COOPER_RIGHT`, and the large majority share the same repository-scoping mechanism documented for `mc_data_archive`

above: a licensed repository genuinely exists for the paper, but `data_check` classified zero of its files as `data`-typed (often because the only file is an archive never opened, or because the licensed repository is a separate code deposit rather than the data-holding one), so `mc_data_license`'s data-only scoping rule never reaches it; a smaller number trace to `repo_check` never detecting the licensed repository at all. The remaining 1 could not be resolved because `repo_check` found no `repo_metadata` row at all for the relevant repository.

Limitation. `mc_data_license` inherits `mc_data_archive`'s repository-scoping limitation, since both restrict themselves to “the repository holding data.” It remains one of the more reliably automatable constructs in this comparison in the METACHECK_RIGHT direction: every case where Metacheck found a license was independently reconfirmed against the same live API `repo_check` itself reads from.

Can the data be downloaded and used?

Cooper et al. report that, of archived datasets, “95% could be downloaded, and 91% could be opened” (Cooper & the BES Data and Code Hackathon Group, 2026b, p. 1959). `mc_data_download` and Cooper et al.'s `data_download` were both codeable for 1287 papers. They agreed on 1273 (98.9%), disagreed on 14 — a small enough count that every case was individually traced rather than summarized by category.

Of the 14 disagreements, 5 were COOPER_RIGHT: files that downloaded successfully but still failed `data_check`'s tabular-usability check for a real, distinct reason each time — a space-delimited scientific data format (climate-model output with a distinctive multi-part naming convention) the tabular reader does not recognise, an unusual two-row header layout, and other content-specific parsing gaps rather than a single repeatable bug. 4 were DIFFERENT_DEFINITION: the data genuinely downloaded and opens correctly in standard software, but `mc_data_download`'s own criterion (see Coding section) additionally requires `data_check`'s stricter `tabular_usable` flag, which does not always agree with “a human can open this file just fine.” 3 were METACHECK_RIGHT: a downloaded file that, on direct inspection, was a serialized model-input dump with no valid row/column structure at all (so `tabular_usable` = `FALSE` was objectively correct, not a detection gap), a clean, working OSF link the paper's own text supports with nothing to contradict it, and a KNB record a prior fix (scienceverse/metacheck#435) now correctly finds and downloads, which Cooper et al.'s No most likely reflects the pre-fix state of. The remaining 2 could not be independently confirmed, since the paper's own text gives no detail on whether the data genuinely opens beyond a bare repository link — exactly the kind of live-repository fact a text-only check cannot resolve either way.

Limitation. This construct's very small remaining disagreement count (14 papers, 1.1% of the codeable base) reflects fixes applied earlier in this project (a multi-sheet Excel reading limitation, scienceverse/metacheck#379; see Appendix A.1 for the full mechanism and before/after numbers from that specific fix). What remains is a heterogeneous residual of genuinely distinct format- and content-specific parsing gaps rather than one further systematic cause.

Is code archived, and can it be downloaded?

Cooper et al. report that “only 35% ($n = 577$) archived their code” among papers using code (Cooper & the BES Data and Code Hackathon Group, 2026b, p. 1958), and, among archived code, that most is on Zenodo or GitHub-family platforms (Cooper & the BES Data and Code Hackathon Group, 2026b, p. 1959). We report `mc_code_archived` and `mc_code_download` together, since the qualitative review of both surfaced overlapping causes.

`mc_code_archived` and Cooper et al.’s `code_archived` were both codeable for 1669 papers. They agreed on 1459 (87.4%), disagreed on 210 (12.6%) — the largest disagreement rate of any construct in this comparison. Sensitivity was 69.3% and specificity 97.0%, McNemar’s $\chi^2(1) = 97.38$, $p < .001$. `mc_code_download` and Cooper et al.’s `code_download` were both codeable for 394 papers (a much smaller base, since this sub-question only applies when code was coded as archived in the first place); they now agree on 368 of 394 (93.4%), with 26 disagreements remaining.

`code_download`’s current near-total agreement reflects a substantial, multi-stage improvement over the course of this project, documented fully in Appendix A.1 (from 66.2% agreement before any fix, through successive intermediate rounds, to the current 93.4%). `code_archived` moved much less across the same fixes, consistent with its dominant causes being ones those fixes did not target.

All 210 `code_archived` disagreements and the 26 remaining `code_download` disagreements were individually resolved. `code_archived`’s 157 COOPER_RIGHT cases are genuinely multi-causal:

1. **Code bundled inside the same repository as the data, with `code_check` finding no file to classify as code there (22 cases).** A recurring, precisely worded pattern: a Data Availability Statement reading, for instance, “Data and R code available from the Dryad Digital Repository” — the paper’s own text explicitly confirms code is bundled with the data in one deposit, `repo_check` found that repository, but `code_check` classified none of its files as code (most often because the only file is an archive never opened, the same zip-peek limitation described next).
2. **`repo_check` never opened a zip/tar archive holding the paper’s code (0 cases).** Metacheck does not inspect an archive’s contents by default; a Zenodo record whose only file is a GitHub-integration snapshot (filename pattern `owner/repo-vX.Y.Z.zip`) is a recurring, easily recognisable instance of this.
3. **`repo_check` never detected the repository at all (43 cases),** despite the paper’s own text explicitly naming it — the same platform-coverage and citation-extraction gap documented for `data_availability` above.
4. **A `repo_check`-side processing error, logged internally (1 cases)** — the repository was correctly identified but its fetch failed with an R-level error or a transient host outage during the corpus run.

5. **A more heterogeneous remainder (86 cases)**, including code provided only as journal Supporting Information rather than in any linked repository, citations to a third-party dependency’s repository rather than the authors’ own code, and `repo_check`-side file-classification errors (e.g., shapefile metadata files misclassified as code).
6. **The R-package-source-tree exclusion defect discussed above (5 cases remaining)**, for the unrelated, pre-existing fetch-gap reason described there rather than the exclusion logic itself.

A smaller number of `code_archived` disagreements ran the opposite direction (29 `METACHECK_RIGHT`): 10 are real, substantial, multi-file GitHub or Zenodo code repositories that `code_check` correctly found and classified, coded No by Cooper et al. — a large, unambiguous set of named analysis scripts is not the kind of borderline case a coding protocol’s edge cases would plausibly explain, so these look like straightforward human-coder misses; a further 9 are papers whose own text and Data Availability Statement mention no code, script, or software at all, so Cooper et al.’s Yes has no traceable basis in the paper’s own evidence. 5 were `DIFFERENT_DEFINITION`: code described only as pseudocode or a script embedded in the journal’s own Supporting Information, which Cooper et al.’s coding appears to credit as “archived” while `code_check` structurally cannot see it (it is not in any linked repository). The remaining 18 were `UNCLEAR`, generally because the extracted paper text contained no relevant statement either way.

Of the 26 remaining `code_download` disagreements, one is `METACHECK_RIGHT`: real, verified downloaded code files (an OSF repository’s 9 code files, each independently confirmed with a non-missing `file_location`) that Cooper et al.’s coding missed. The other is `COOPER_RIGHT`: all 3 code files `code_check` found for this paper’s repository have a missing `file_location` — the same archive-member download-failure pattern traced and largely fixed elsewhere in this project (scienceverse/metacheck#429), affecting a paper that fix’s own rerun did not happen to cover.

Limitation. `code_archived` remains the construct in this comparison with both the highest disagreement rate and the most heterogeneous set of causes. The single largest cause (code co-located with data in a repository `code_check` cannot classify, whether because the file is an unopened archive or for another reason) points to the same underlying zip-peek limitation affecting `data_archive` and `data_license` above; the second-largest (repository detection coverage) is the same gap documented for `data_availability`. `code_download`’s gap, by contrast, has narrowed to 26 remaining cases after the fixes documented in Appendix A.1 — this construct is, at present, essentially resolved.

Is there a README?

Cooper et al. report a README present for “66% (n = 1114) and 61% (n = 351) of the papers with archived data or code, respectively” (Cooper & the BES Data and Code Hackathon Group, 2026b, p. 1959). As described in the Coding section, Metacheck records a single README

signal per paper (computed once, across every linked repository combined), so we compare it against `cooper_any_readme` — TRUE if Cooper et al. coded “Yes” or “Quasi-README” on either their `data_README` or `code_README` question. `mc_any_readme` and `cooper_any_readme` were both codeable for 1726 papers. They agreed on 1244 (72.1%), disagreed on 482 (27.9%) — the largest raw disagreement count of any construct in this comparison, though, as with `code_archived`, its composition changes the interpretation substantially. Sensitivity was 62.9%, specificity 95.5%.

Of the 482 disagreements, 23 were `METACHECK_RIGHT` on review — genuine cases where a real README file exists in a repository `repo_check` listed, and Cooper et al.’s coders missed it. 155 were `COOPER_RIGHT`, dominated (116 cases) by the same repository-detection gap documented for `data_availability` and `code_archived` above — `repo_check` never located a repository the paper’s own text explicitly names, so no README evidence could ever be recorded — with a further 0 cases where the repository was found but its only file is an unopened archive, and 2 cases tracing to a `repo_check` processing error. The remaining 12 were `UNCLEAR`.

The single largest verdict category, however, is 294 `DIFFERENT_DEFINITION` cases (61.0% of all disagreements, well ahead of the `COOPER_RIGHT` count above) — not a detection failure on either side, but the two coding approaches answering genuinely different questions that happen to share Cooper et al.’s column label:

- **Dryad’s abstract field (the largest single cause, 172 cases).** Dryad’s own dataset API exposes exactly one free-text field, `abstract`. Checked directly on real Dryad datasets from this disagreement set, this field is consistently the paper’s own scientific abstract — never file-by-file documentation of what the archived files contain. Cooper et al.’s coders, reading the same field on the live Dryad page, appear to have coded this abstract as satisfying “README present,” a defensible reading of their own protocol given what Dryad actually surfaces to a human visitor; `repo_check`’s `doc_role` classification, which requires an actual downloadable file named like a README, never counts it.
- **Zenodo’s and Figshare’s description fields (88 cases).** Unlike Dryad’s abstract, these platforms’ `description` fields are genuinely mixed in content — checked directly on records from this disagreement set, real file-by-file documentation appears in some cases, a study summary in others, contact metadata in others, and an empty field in others. No single, reliable rule (treat the description as a README, or don’t) would correctly resolve every case.
- **An archive-only repository whose single compressed file `repo_check`’s remote listing could not resolve (16 cases).** Any README bundled inside such an archive is invisible to the README check even when `repo_check`’s own zip-listing feature is enabled, for the same family of archive-access limitations discussed above.
- **A smaller remainder (18 cases)** not matching any of the three patterns above.

Limitation, and a genuinely different framing than the other constructs in this comparison. Unlike every other column reported here, `any_readme`’s disagreement is overwhelmingly not a case of either side being wrong. It is a case where Cooper et al.’s human

coders, reading a live repository page, correctly used whatever free-text field that platform’s own interface surfaced to them as evidence of a README-like description, while Metacheck’s `doc_role` classification asks a narrower, more conservative question (is there a downloadable file that is actually named like a README) that structurally cannot see a platform’s own abstract or description metadata at all. Resolving this gap in Metacheck’s favor is not simply a matter of relaxing a filename rule: Dryad’s abstract field is confirmed never to be README-equivalent content (so counting it would introduce false positives), while Zenodo’s and Figshare’s description fields are confirmed to be genuinely mixed (so counting them uniformly would introduce a different, more Zenodo/Figshare-specific rate of false positives). A reliable fix would need to distinguish, likely with a content classifier rather than a presence check, which platform-level description fields actually contain file-documentation content, on a case-by-case rather than platform-wide basis. ## Is a comment count a proxy for comment quality?

Cooper et al.’s protocol includes a construct no Metacheck module attempts to answer directly: `code_annotation_scale`, a human coder’s 1-10 judgment of how well a paper’s code is commented, coded for 543 of the 1861 papers. Metacheck’s `code_check` module does not judge comment quality, but it does mechanically count comments: for every code file it can download and parse, it records `comment_lines`, `code_lines`, and `percentage_comment` (comment lines as a share of all lines). This raises a genuinely exploratory question distinct from every comparison above, none of which asked whether an automatable signal could stand in for a judged one: if a simple, free line-count already tracked Cooper et al.’s human quality rating closely, that would suggest the more expensive human judgment adds comparatively little beyond what counting could already tell us; if it does not, that would suggest the human coders are picking up on something a mechanical count cannot see.

We aggregated Metacheck’s per-file comment counts to one value per paper — total comment lines divided by total lines across every code file `code_check` could download and parse for that paper — and compared it against Cooper et al.’s `code_annotation_scale` for the 354 papers where both were available (of 543 papers Cooper et al. rated and 410 papers Metacheck could compute a comment percentage for at all; the smaller overlap reflects that Metacheck’s count requires at least one code file to have actually downloaded and parsed, which, as the sections above document, is itself not guaranteed).

The correlation is weak but not zero: Pearson’s $r = 0.19$, 95% CI [0.08, 0.28], $p < .001$; Spearman’s $\rho = 0.16$, $p = 0.003$. Using the raw comment line count instead of the percentage gives essentially the same result ($\rho = 0.12$, $p = 0.021$), so this is not an artifact of how the percentage was computed. As a check on whether longer codebases are simply commented differently regardless of quality, we also correlated total code lines against percentage of comments directly: longer codebases have a modestly lower comment percentage ($\rho = -0.31$, $p < .001$), consistent with comments not simply scaling with code length, and not something that would by itself explain the annotation-scale correlation being as weak as it is.

Interpretation. A correlation of this size — roughly 3% of variance explained on the Pearson estimate — means a paper’s comment density is a poor stand-in for a human coder’s quality judgment: knowing what fraction of a codebase’s lines are comments tells us very little about

whether a coder would call those comments good annotation. This is exploratory, not a claim about causation or about which measure is “right” — the two are answering different questions by design, and a low correlation between them is not evidence that either is wrong. But it does directly speak to the question motivating this comparison: for this specific construct, an automatable proxy does not substitute well for the human judgment call, at least not the simplest version of that proxy (line-count percentage). A single well-placed sentence explaining a complex analysis counts as excellent annotation on Cooper et al.’s scale and as a single line on Metacheck’s; ten lines of boilerplate restating variable names count as heavily annotated by line count and would plausibly score poorly on a human quality read. Comment content — whether a comment explains a non-obvious decision, restates the obvious, or is simply absent — is the kind of thing an LLM-based classifier could more plausibly approximate than a line count can, but that would need its own validation against this same ground truth before being trusted, following the same case-by-case standard used throughout this manuscript rather than assumed from the fact that an LLM could be asked to try.

Results you get for free alongside Metacheck

Every comparison above answers a question Cooper et al.’s protocol already asked. Running `repo_check`, `data_check`, and `code_check` on a corpus also produces a substantial amount of additional, structured diagnostic information that no item in Cooper et al.’s protocol asks for at all — not because Metacheck was built to exceed their protocol, but because the same file listing, the same downloaded code, and the same parsed spreadsheets that answer “is there a repository” or “is there a README” also support other checks the modules perform as a matter of course. This section reports what that additional output looks like at corpus scale, computed entirely from the same saved module output as the comparison above (`res_repo_check.RData`, `res_data_check.RData`, `res_code_check.RData`) — no additional module run was needed.

Code portability and reproducibility signals

`code_check` parsed 9130 code files across 437 code-bearing papers. Three checks bear directly on whether that code would actually run for someone other than its original author. (A systematic validation of these checks against real corpus files, reported in Appendix A.3, identified and fixed five false-positive/false-negative mechanisms; the numbers below reflect the corpus as run after those fixes.)

- **Missing referenced files.** For every code file, `code_check` extracts every filename the code itself tries to open (`read.csv(...)`, `source(...)`, and equivalent calls in each supported language) and checks whether that name appears anywhere in the same paper’s repository listing. 1101 files (12.1% of all analysed files), across 209 papers (47.8% of code-bearing papers), reference at least one file that could not be found. This does not necessarily mean the code is broken — a file might exist under a name the regex-based extraction did not match exactly, or might be generated by an earlier step in the same

pipeline — but it is a direct, corpus-wide signal of how often “the code plus the shared data” does not obviously line up as a complete, runnable unit.

- **Hardcoded absolute paths.** 405 files (4.4%), across 104 papers (23.8%), contain an absolute file path (e.g. `C:/Users/...` or `/home/...`) rather than a path relative to the project. Code written this way will not run on any machine but the one it was written on without manual editing.
- **`setwd()` calls.** 367 files (4.0%), across 94 papers (21.5%), call `setwd()` to hardcode a working directory. Combined with hardcoded absolute paths (which frequently co-occur with a `setwd()` call to the same machine-specific location), the two portability problems together affect 254 of the 437 code-bearing papers (58.1%) — meaning more than half of the papers whose code we could inspect at all have at least one file that will not run as shared, on someone else’s machine, without modification.

Comment density (mechanical count)

As reported above in the comment-quality correlation, `code_check` records `comment_lines/code_lines` per file. Corpus-wide, comments make up 18.9% of all code lines. 979 files (13.3%) have zero comments at all — of which 232 of 316 zero-comment Python files (73.4%) are in fact documented via a docstring rather than a `#` comment, a distinction `code_check` now also reports separately (`has_docstring`, Appendix A.3), so are not genuinely undocumented despite scoring zero on this mechanical count.

Import/library organization

Where a file loads more than one package, `code_check` also records how far apart those loading lines are (its own documented threshold is more than a few lines apart counting as “scattered” rather than grouped at the top of the file). 981 of 4001 such files (24.5%) load their packages in a scattered rather than grouped fashion — a minor readability signal, not a correctness one, but one a reader piecing together a script’s dependencies would find genuinely useful.

Parse errors

Of the 4693 R-type code files (`.R`, `.Rmd`, `.qmd`) that were both locally available and actually reached a call to `parse()`, 153 failed to parse as valid code in their declared language at all (3.3%) — a strong, unambiguous signal (not a stylistic judgment) that something is wrong with either the file itself or Metacheck’s identification of its language. (Only R-type files are eligible for this check at all; a file that was never downloaded, or that errored before reaching the parse step itself, is excluded from this denominator rather than counted as a pass — see Appendix A.2 for the full derivation and the two distinct causes behind the files excluded this way.)

Package version pinning

`code_check` also checks whether a repository pins the exact R/package versions its analysis depended on — an `renv.lock` file, a `sessionInfo()/session_info()` text dump, or a `groundhog/checkpoint` date-pinning call in the code itself (a bare `library(groundhog)` does not count; the pinning call itself must be present). Of the 437 code-bearing papers, 22 (5.0%) have version pinning present, recorded per paper as `code_version_pinned` in `code_check`'s own summary output (see Appendix A.2 for how this per-paper figure was recovered).

Spreadsheet and column-level data-quality findings

`data_check` inspects every tabular data file it can download and parse, both at the file/sheet level (is this a rectangular table at all, is the header on the first row, does the file use colour or merged cells to encode information Metacheck cannot read from a plain CSV) and at the level of individual columns (does a column look like it holds personal information, geographic coordinates, inconsistent category casing, numbers stored as text, or values in a legacy character encoding). Across the corpus, these checks produced 34857 findings, spanning 2946 distinct files. The most frequent, in descending order:

Finding	Count
Empty column	19,529
Numeric as text	4,618
Problematic column name	3,107
Colliding column names	1,591
Empty or unnamed columns	1,366
Personal info (column name)	1,126
Whitespace	661
Geographic coordinates	615
Case issues	577
Empty rows	472
Colour coding	432
Merged cells	210
Free-text (may hold PII)	151
Mixed encoding	133

Finding	Count
Un-inspectable (.xls)	72
Header not on first row	66
Personal info (values)	61
Not a rectangular dataset	43
Unreadable	27

Two of these are worth calling out specifically because they bear on participant confidentiality rather than only on tidiness: the two personal-information checks (by column name and by column values) and the “free text that may hold PII” check together flag 1338 findings across the corpus — columns or values that look, on their face, like they could contain names, contact details, or other identifying information in a file the authors have made publicly downloadable. Metacheck does not verify whether this is a genuine confidentiality problem (a column named “participant_id” that holds only anonymous sequential integers would also match on name alone), but a count this size suggests a possible participant-privacy check would find comparable value to the correctness-oriented deposit review Cooper et al.’s protocol already performs.

File-naming conventions

`repo_check` checks every listed file’s name against a set of FAIR-data-adjacent naming conventions: no spaces or special characters, an actually-classifiable name (not left as `data_type == "unknown"`), consistent casing, zero-padded sibling numbering, valid dates, and several path-length budgets. Across the corpus, this produced 64687 naming issues, 36582 of them the more serious “bad” severity (something that actually breaks: spaces, special characters, diacritics, an unclassifiable file, an invalid date, or an over-long path) and 28105 the softer “suggestion” severity (a real convention worth following, but nothing in Metacheck’s own classification depends on it). 728 distinct papers (47.1% of the 1545 papers with a discoverable repository) have at least one naming issue of some kind. The single largest contributor by far is “unclassifiable” file names (23249 occurrences) — files whose name gives Metacheck’s classifier no signal at all about what kind of content they hold, which is itself a legibility problem for a human reader trying to navigate the repository, independent of anything Metacheck specifically failed to detect.

Why this section exists. None of the findings above were validated against Cooper et al.’s human coding, because Cooper et al.’s protocol did not ask about any of them — there is no ground truth to check against, and no sensitivity/specificity/agreement-rate analysis is possible for a construct nobody coded. We report them here, separately from the validated comparison above, specifically so that this distinction stays clear: these are Metacheck’s own outputs, useful in their own right for a researcher curious about a specific repository’s quality,

but not independently verified the way every number earlier in this manuscript was. A natural next step, beyond the scope of this paper, would be a smaller, targeted validation of a sample of these specific findings (are the “unclassifiable” file names really unclassifiable to a human reader too; do the flagged personal-information columns really hold identifying information) using the same case-by-case standard applied throughout the rest of this manuscript.

Discussion

Table 0 summarizes the empirical comparison reported above: agreement rate, sensitivity, specificity, and the two directional miss rates for every construct where Metacheck and Cooper et al.’s coding can be directly compared.

Table 3: Table 0. Empirical agreement between Metacheck and Cooper et al.’s human-coded ground truth, for every comparison column. “Cooper caught, MC missed” and “MC caught, Cooper missed” are percentages of the disagreements only (they sum to 100% within a row), not of the full sample.

Construct	n	Agree %	Sensitivity %	Specificity %	Cooper caught, MC missed %	MC caught, Cooper missed %
Data availability	1,735	84.7	85.7	44.4	90.6	9.4
Data licence present	1,306	95.6	99.1	25.4	19.0	81.0
Data downloadable and usable	1,287	98.9	99.1	0.0	85.7	14.3
Code archived	1,669	87.4	69.3	97.0	84.3	15.7
Code downloadable	394	93.4	93.6	0.0	96.2	3.8
Any README present (data or code, collapsed)	1,726	72.1	62.9	95.5	95.4	4.6
Archive platform (overlap match, not simple boolean)	1,385	96.5				

Two patterns are visible across every construct in Table 0. First, the size and direction of the sensitivity/specificity gap now varies considerably from construct to construct, and reviewing every individual disagreement, rather than stopping at the aggregate numbers, shows why: the asymmetry has a different cause in each column, and the cause matters more than the raw percentage for deciding what, if anything, to fix. `data_availability`’s specificity (44.4%) is far from the zero an earlier pass of this corpus found, precisely because the derivation-rule fix described in the Coding section gave it a genuine ability to disagree with Cooper et al. in the FALSE direction at all; what remains is a repository-coverage gap (see above), not

a structural inability to ever produce a negative. `code_download`'s specificity is still 0.0%, but on a base of only 26 remaining disagreements this reflects a small handful of unresolved cases, not a systematic pattern; `data_download` currently has no codeable Cooper-No case at all, so its specificity is undefined rather than low. `code_archived` shows the most pronounced asymmetry of any boolean construct (sensitivity 69.3%, specificity 97.0%) and is genuinely multi-causal, spanning a real content-location mismatch (code bundled with data in a repository `code_check` cannot classify), a zip-peek limitation, a repository-detection gap, and a real “not in any repository at all” category (journal Supporting Information) no repository-scanning approach can close. `any_readme`'s asymmetry is not primarily attributable to either side being wrong at all: 61.0% of its disagreements reflect Cooper et al.'s human coders correctly reading platform-level metadata (a Dryad abstract, a Zenodo or Figshare description) that Metacheck's stricter, filename-based README check was never designed to treat as equivalent evidence, and where doing so uniformly would, on the evidence collected here, introduce as many new false positives (Zenodo/Figshare) as it resolved.

Second, precision in the direction Metacheck does commit to an answer remains generally high, though no longer universal: positive predictive value was above 97% for five of the six boolean constructs (`data_availability` 98.3%, `data_license` 96.3%, `data_download` 99.8%, `code_download` 99.7%, `any_readme` 97.3%), meaning that when Metacheck reports a repository, license, download, or README as present, that report is rarely spurious. `code_archived` is the exception: its PPV, 92.4%, is meaningfully lower, because 29 of its 210 disagreements are cases where `code_check` correctly found real code files that Cooper et al.'s coders missed or judged not to count — a genuine human-coder miss, not a Metacheck false positive, but one that lowers this specific metric regardless of which side is right. The asymmetry in this comparison therefore remains overwhelmingly a recall problem for most constructs — Metacheck under-detects relative to a human coder who can read a repository's page, follow an ambiguous citation, or interpret an unusual sentence structure — but `code_archived` shows this is not universal: a TRUE from Metacheck is a highly trustworthy signal for five of the six boolean constructs, while for `code_archived` specifically, a nontrivial share of Metacheck's FALSE answers reflect a real detection gap rather than code genuinely being unarchived.

Three constructs in Cooper et al.'s protocol have no Metacheck counterpart at all and were therefore not part of the quantitative comparison above: README/code usefulness (a human-judged 1-10 scale), data completeness relative to a paper's reported analyses, and code annotation quality as distinct from comment density. These are precisely the questions where Cooper et al.'s human coders had to read and interpret content, not merely detect its presence, and they mirror a pattern visible even within the constructs we could compare: Metacheck's underlying modules answer structural questions (does a file with this name exist, did a download succeed, does a table parse) reliably, and answer content-interpretation questions (is this description a real README, is this the paper's actual dataset rather than simulation output, is this repository the one holding the data rather than the code) far less reliably, precisely because those questions require reading and judging content rather than classifying structure. Large language models are a plausible route to closing part of this gap for the constructs Metacheck cannot currently address at all, and could plausibly also help with several of the misclassification

patterns documented above (distinguishing a Dryad abstract from file documentation, or a simulation-output file from raw collected data) — but any such extension would need to be built, validated against Cooper et al.’s own ground truth using the same case-by-case standard applied here, and reported with a known error rate, rather than assumed to work from the fact that an LLM could plausibly be prompted to attempt it.

A recurring, cross-cutting finding worth naming explicitly is the data-versus-code repository conflation pattern: a genuinely common paper structure — one repository holding the study’s data, a separate repository (often GitHub) holding its analysis code, sometimes with the data-holding repository never located by `repo_check` at all — appears as the dominant or a substantial cause of disagreement in `data_archive`, `data_license`, and `code_archived` alike. This is not three independent bugs; it is one architectural gap (Metacheck’s repository-scoping logic assumes each construct maps cleanly onto “the” repository for a paper) surfacing across every construct that depends on correctly identifying which repository is relevant. It remains, after every fix applied over the course of this project, the single largest reason Metacheck and Cooper et al. disagree: repository non-detection alone accounts for 192 of `data_availability`’s 230 COOPER_RIGHT disagreements, 43 of `code_archived`’s 157, and 116 of `any_readme`’s 155 — a paper’s own text names a real, working DOI on a mainstream platform (overwhelmingly Dryad, Zenodo, or Figshare), and `repo_check` never queries it at all. Addressing the conflation and non-detection gaps once, at the level of how Metacheck decides which repository “counts” for a given question and which platforms it recognises, would likely resolve a meaningful fraction of the remaining disagreement across several constructs simultaneously, rather than requiring separate, construct-specific fixes.

Twelve individual software defects identified over the course of this project were reported upstream, fixed, and reverified against a full corpus rerun for this manuscript. Five were found during the original review: the Excel multi-sheet `tabular_usable` limitation underlying most of the original `data_download` disagreement (scienceverse/metacheck#379, resolved: agreement rose from 93.5% to 99.1%), the Dryad unauthenticated-HEAD-request gap affecting both data and code downloads (scienceverse/metacheck#380), the `.zip`-peek authentication gap for archive-only repositories (scienceverse/metacheck#378), the `code_check` zip-discovery architectural gap surfaced while investigating the Zenodo data-download fix (scienceverse/metacheck#383), and a silent-truncation bug in the byte-fetch path that #383’s fix newly exposed (scienceverse/metacheck#384): the function decompressing a zip member’s bytes called an underlying library routine with no size hint, silently capping any member over 32,768 bytes at exactly that length with no error, rather than the size-aware call the library’s own exact-size path rejects for an unrelated reason. Fixing #384 closed the great majority of the gap #383 had left open: `code_download` agreement rose from 66.2% (before #379/#380/#383) to 89.7% (after them) to 93.4% (after #384) — a two-stage improvement of over 30 percentage points overall, the largest of any construct in this comparison, and a level of agreement that has held with only 26 disagreement remaining across every subsequent fix and rerun described below.

A further seven fixes (Appendix A.5) — Zenodo reliability, wider repository-platform coverage,

and a rate-limit handling defect — were applied and reverified against the entire 1,861-paper corpus rather than only the papers already known to disagree, specifically to check whether a fix aimed at one construct could have broken an already-correct answer elsewhere. Six of the seven had not; the seventh — the R-package-source-tree exclusion, meant to drop a bundled dependency’s own scaffolding (`DESCRIPTION`, `NAMESPACE`, `R/`, `man/`, `tests/`) from a repository’s file listing — also stripped a paper’s own analysis code whenever that code was itself packaged as an installable R package, a common practice for ecological modelling and analysis tools shared alongside a methods paper, and in several cases stripped the package’s bundled `data/` directory as well, turning a real dataset into an apparent absence of one. This was confirmed against 16 real GitHub/CRAN-hosted repositories (`RevGadgets`, `BaBA`, `m2b`, `RRPP`, `actel`, `scampr`, `MINOTAUR`, `LadderFuelsR`, `biomonitor`, `EcoEnsemble`, `fossilsim`, `frair`, `tRophicPosition`, `sabinaNSDM`, `bef`, and `phylomice`). Rather than being left as an open limitation, the exclusion rule was corrected within this same review cycle: it now only strips a root-level package’s files when nothing classifiable as real data or code survives outside it — using the same `data_classify_files()` rules-based classifier `data_check` itself relies on — distinguishing a study that shares its own data and code as an R package (nothing survives outside it, because the package is the deposit) from a repository that bundles a vendored copy of someone else’s package alongside the study’s own files (real content survives outside the vendored package, so only the vendored package’s files are excluded). Reapplying the corrected rule resolved 11 of these 16 cases; the remaining 5 are due to an unrelated, still-open gap where the affected repository’s own cached download never retrieved its `R/` source directory in the first place, so no exclusion-logic fix can recover content that was never fetched.

We see this kind of construct-by-construct, evidence-grounded mapping — quantifying not just how often Metacheck and a trained human coder disagree, but why, down to the individual paper and, where applicable, the specific line of code responsible, reverifying that mapping against a real fix rather than stopping at diagnosis, and checking a fix’s effect on the papers it was not meant to change as carefully as its effect on the papers it was — as the more useful contribution here than a single aggregate accuracy figure, both for meta-scientists deciding whether Metacheck can currently support a given coding task, and for Metacheck’s own development priorities. Appendix A documents, in full detail, every defect identified over the course of this project, the evidence that surfaced it, the exact code change that fixed it, and the before/after numbers at each stage; the body of this manuscript reports only the current, post-fix numbers.

Detailed record of every Metacheck defect found and fixed during this project

This appendix documents, in full detail and in chronological order, every software defect in Metacheck that this project’s validation effort identified, the evidence that surfaced each one, the exact code change that fixed it, and the corpus-level numbers before and after that fix. The body of the manuscript reports only the final, current numbers; this appendix exists so that every intermediate claim made in the course of arriving at those numbers remains checkable, without requiring the main text to carry a running commentary on its own revision history. Seven rounds of validation are documented below: (A.1) the original empirical-comparison review against Cooper et al.’s coding, which produced issues #378–#384; (A.2) a focused re-examination of `code_check`’s parse-error and package-version-pinning checks; (A.3) a systematic validation of five further `code_check` checks (missing referenced files, hardcoded absolute paths, `setwd()` calls, comment density, and library/import scattering) against real corpus files, which produced five further fixes; (A.4) a file-count gate added after one outlier repository’s processing time exposed a gap the existing byte-size caps did not cover; (A.5) a full re-check of every remaining disagreement against Metacheck after five more fixes (Zenodo reliability, repository-platform coverage, and a rate-limit reliability defect found live during the re-check itself); (A.6) a systemic under-sampling bug found in A.5’s own Dataverse DOI-prefix allowlist, two further platform-coverage fixes (DataONE, institutional Figshare instances) and two matching bugs in this project’s own comparison code, and a complete, evidence-based categorization of every “repository not detected” disagreement these fixes did not resolve; and (A.7) a full-corpus run of Metacheck’s `reproducibility_check` module — a different question from A.1–A.6 (does a paper’s own shared code actually run and reproduce its own reported results, not merely whether a repository was detected) — which surfaced five further defects specific to running many papers’ Docker containers concurrently at corpus scale and to extracting a paper’s own declared dependencies from its code. Fixes described in A.1–A.4 were committed to `scienceverse/metacheck`, branch `repocheck_improvements` (pull request #385); A.5 and A.6’s fixes were committed directly to that package’s `dev` branch (A.6 as pull requests #410, #412, and #414); A.7’s five defects were filed as issues #417–#421, two of which (#419, #420) are confirmed fixed on `dev` as of this writing. Every fix from A.1–A.6 is covered by a new or updated automated regression test in the package’s own test suite; A.7’s two confirmed-fixed defects were fixed and committed by the package’s maintainer rather than by this project directly, so this manuscript does not independently confirm test coverage for them.

A.1 Defects found during the original empirical comparison

Six defects were identified while individually reviewing every disagreement between Metacheck and Cooper et al.’s human coding (Empirical Comparison, above): five in Metacheck itself, and one in this project’s own comparison code. All six are described here with the same level of

detail as in the body text; they are repeated in this appendix only for completeness, so that Appendix A is a single, complete record of every fix made during this project.

#378 — `zip_peek()`'s unauthenticated requests. `zip_peek()` (`R/zip-peek.R`) reads a remote zip archive's central directory via an HTTP Range request, without downloading the archive, so a downloader can decide whether the archive is worth fetching before spending the bandwidth. This function's HEAD and Range requests were unauthenticated. Dryad's file-download endpoint answers 401 Unauthorized to any request that touches file bytes, including a bare HTTP Range request — confirmed live against a real Dryad file-download URL — so every Dryad-hosted archive's peek failed structurally, before any question of the archive's actual contents. Fix: both requests now route through `.auth_for_url()`, the same host-based authentication dispatcher every other Metacheck download path already uses, so Dryad's token is attached automatically without requiring each call site to opt in individually (an earlier, per-call-site `req_func` opt-in mechanism had only been adopted by 1 of 7 call sites that needed it).

#379 — `data_check`'s multi-sheet Excel limitation. `data_check`'s tabular-usability check called `readxl::read_excel()` with no `sheet` argument, so it only ever read a workbook's first sheet. When the first sheet was a cover, metadata, or “read me” sheet — a convention confirmed across many real corpus files, with sheet names including “metadata,” “Meta-data,” “MetaData,” “read me,” “ReadMe,” “User guide,” “LEGEND,” and “Description of variables” — a later sheet holding genuine, well-structured tabular data was never inspected, and the file was misclassified as unusable (“100% free text, not variables”). One confirmed case: sheet 2 of a real downloaded workbook held 1,532 rows by 58 real variable columns, never read before the fix. Fix: `data_read_head()` now accepts a `sheet` argument, and `data_check` tries each remaining sheet in a workbook before giving up on the file. This was the single largest driver of `data_download` disagreement: before the fix, `mc_data_download` and Cooper et al.'s `data_download` disagreed on 51 of 967 comparable papers (94.7% agreement); after it, agreement rose to 99.1% (908/916 papers, 8 disagreements remaining), and the resolved cases accounted for 29 of the original 51 disagreements (57%).

#380 — Dryad's authentication gap on other download paths. The same unauthenticated-request problem identified for `zip_peek()` (#378) also affected `.remote_size()` and `.remote_content_length()`, two other functions in the download path that check a remote file's size via HTTP HEAD before deciding whether and how to download it. `.remote_content_length()` already had a `req_func` opt-in mechanism for per-host authentication, but only 1 of 7 call sites across the codebase had actually adopted it, so the other 6 still received a 401 from Dryad. Fix: both functions now route through `.auth_for_url()` unconditionally, the same fix applied to #378, closing every remaining unauthenticated-request gap in the download path at once rather than patching call sites one at a time.

#383 — `code_check` could not see code bundled inside a zip archive. This was a compound architectural gap with three separate causes, discovered while investigating why `code_check` found no code files at all for several papers whose repository, on manual

inspection, clearly contained real analysis scripts bundled inside a `.zip`: (1) `code_check()` never read `data_check`'s already-expanded file structure, so any code file `data_check` had already discovered inside a zip (in the course of looking for data files) was invisible to `code_check` regardless; (2) `code_check` had no zip-peek capability of its own for the case where `data_check` had not already run first; (3) the shared zip-expansion helper `.archive_rows()` discarded code-typed files outright even when it did expand a zip, keeping only `"data"` and `"documentation"`-typed files. Fix: `code_check()` now prefers `data_check`'s structure when available; it gets its own zip-peek fallback (`.code_expand_zip()`) for when `data_check` did not run first; and `.archive_rows()` now keeps `"code"` alongside `"data"/"documentation"`. A real, concrete example verified live: a Zenodo record (record 3694205, repository SharpRT/CMV_Spread) bundled 4 real MATLAB `.m` files inside a zip that were entirely invisible to `code_check` before this fix. This fix, together with the archive-member content-fetch wiring described immediately below, moved `code_download` agreement from 66.2% (355 comparable papers, 120 disagreements) to 89.7% (358 comparable papers, 37 disagreements) — the first of two stages in this construct's improvement (the second is #384, below).

As part of the same round of work, `repo_check`'s zip-peek expansion was extended to carry not just a discovered code file's name (already possible after the fixes above) but its content: each expanded row now also carries `archive_url` (the parent zip's own URL) and `archive_member` (the file's exact name inside it), and `download_repo_files()` (`R/repo-download.R`) uses those two fields to fetch that one member's bytes directly, via a byte-range request against the parent archive, without downloading the archive's other contents. This mechanism was verified working in every isolated unit test at the time, but a full end-to-end pipeline run against a real paper still returned no file content for reasons not yet diagnosed at that point — this was filed and tracked as issue #384, described next, rather than merged as “fixed” while a known gap remained open.

#384 — Two distinct causes behind the archive-member byte-fetch failing intermittently. Issue #384 was filed as an open question (“archive-member content fetch works in isolation but returned NA once in a full pipeline run”) rather than a diagnosed bug, and its investigation spanned two separate validation efforts and uncovered two independent causes.

Cause 1: silent truncation over 32,768 bytes. `.zip_inflate_member()` (`R/zip-peek.R`) decompresses one zip member's raw bytes by calling `zip::inflate()`. Passing the member's exact known uncompressed size to that function raised an unrelated “Failed to inflate data” error (a quirk internal to the `zip` package), so the original code left the `size` argument unset instead. `zip::inflate()`'s own documentation states that leaving `size` unset triggers automatic buffer resizing as needed — but, confirmed live against a real Zenodo-hosted zip member, this documented behaviour does not actually happen in practice: any member larger than the default 32,768-byte output buffer came back silently truncated at exactly that length, with no error, no warning, and a `bytes_written` field the calling code never inspected. A concrete confirmed case: a real 44,669-byte code file came back as exactly 32,768 bytes. `.zip_member_fetch()`'s own subsequent length check (comparing the decompressed length against the size recorded in the archive's own central directory) correctly caught this mismatch and discarded the truncated

bytes — but the practical, user-visible effect was that every code or data file over 32KB inside a zip silently failed to fetch, indistinguishable from a genuine network or host-compatibility failure. Fix: passing `size + 1` (one byte over the true size) avoids both failure modes — it is no longer the exact size that triggers the package’s internal error, and it is large enough that the real output always fits in the up-front buffer rather than depending on the documented-but-non-functional resize path. Verified against both a real archive member under 32KB and one over it, and covered by a new offline regression test using a locally constructed zip file whose one member compresses to roughly 95KB uncompressed. A targeted rerun of the 36 papers affected by this pattern found 27 moved from zero downloadable code files to full or near-full recovery (23 fully, 4 partially), which is what raised `code_download` agreement from 89.7% (37 disagreements) to 97.2% (10 disagreements remaining) — the second of the two stages in this construct’s improvement, and the largest single agreement-rate improvement of any construct in the empirical comparison.

Cause 2: an intermittent, silently-swallowed failure not explained by the truncation bug alone. Independently of the truncation bug, a second cause was found while investigating why 40 of the archives affected by issue #384 failed entirely (every member missing) despite containing no member anywhere near the 32,768-byte threshold. Tracing the actual corpus data (`res_code_check.RData` and its 38 per-batch constituent files, in `cooper_validation_metacheck/manuscript/data/`) showed that a single pipeline run calls `zip_peek()` on the exact same archive URL twice — once inside `repo_check` to discover a zip’s member names, and again inside `download_repo_files()`’s archive-member-fetch step (added as part of #383/#384’s own fix) to retrieve one member’s bytes — with no throttling and, critically, no retry logic at all on either call: `zip_peek()`’s own HEAD request and its two Range-request helper functions (`.http_range_tail()`, `.http_range_bytes()`) called `httr2::req_perform()` directly, with no `req_retry()`, and treated any non-206/200 HTTP status (including a 429 burst-rate-limit refusal) as an unconditional, silent NULL — indistinguishable, by design, from “this host does not support range requests at all.” Two back-to-back, unthrottled requests to the same archive URL within one process is exactly the pattern most likely to trip a host’s burst-rate limiter, and the second call (the byte-fetch, inside `download_repo_files()`) was the one exposed to it. Fix, applied in two parts: (1) all three of `zip_peek()`’s network calls now retry through the same `is_transient/backoff/rate-limit-wait` helper functions (`.storage_is_transient_factory()`, `.storage_backoff()`, `.storage_retry_after_factory()`) that every other Metacheck download path already uses, so a 429 is now retried with backoff rather than immediately and silently treated as “unsupported”; (2) a new same-session, in-memory cache (`.zip_peek_cache`, keyed by URL) means the second call to `zip_peek()` on an archive already peeked earlier in the same run reuses the first call’s result instead of repeating the network request at all, removing the redundant call — and the rate-limit exposure it caused — entirely rather than merely tolerating it. Verified by a new regression test exercising the cache directly (`test-zip-peek.R`) and by the full package test suite (463 tests passing, 3 pre-existing unrelated skips, 0 failures) after both changes.

Verdict corrections found by the post-rerun double-check pass. After the full corpus

rerun incorporating #378–#384, every disagreement was independently double-checked a second time (see Empirical Comparison, above, for the check’s method). This pass corrected 7 verdict-labeling errors from the original review: 2 in `data_availability` (one moved to `COOPER_RIGHT` after finding the “data” files actually belonged to a cited third-party software dependency’s repository, not the paper’s own; one moved to `METACHECK_RIGHT` after finding a methods paper’s repository, originally read as containing only package example files, in fact holds real trait data and model output), 1 in `data_download` (moved to `COOPER_RIGHT` after confirming a Zenodo-hosted “Functional Trait Database” file had, in fact, already downloaded and still failed the tabular-usability check, rather than remaining `UNCLEAR` on the assumption it had not yet downloaded), and 4 in `code_download` (three rows originally recorded `METACHECK_RIGHT` whose own written reasoning already concluded the opposite, corrected to `COOPER_RIGHT`; one row originally left `UNCLEAR` where the code files were, on direct inspection, confirmed genuinely downloaded, corrected to `METACHECK_RIGHT`). The pass also found one substantial scope correction rather than a single-row relabeling: `code_download`’s dominant post-#383 disagreement pattern was originally tagged in only 13 of its 37 disagreements at the time of the initial qualitative review; re-deriving the pattern systematically against every disagreement’s saved `file_location` value found the identical signature (a discovered code file with no resolvable download location) in 35 of the 37, which is what led directly to tracing and fixing #384’s second cause, above. Every one of `code_archived`’s 164 disagreements was also individually re-derived against the current `repo_check/code_check` file listing during this pass, including the 8 originally resolved via an automated keyword search rather than an individual hand read; this found zero verdict corrections in that column.

A sixth issue, not in Metacheck but in this project’s own comparison code. `code_data_preparation/03_recreate_cooper_columns.R` derives `mc_data_availability`, `mc_code_archived`, `mc_data_README`, and `mc_code_README` for a paper from `repo_check`’s per-file tables (`res_repo_check$table`, `res_data_check$structure`, `res_code_check$table`). A paper for which `repo_check` found zero files anywhere has no row at all in any of those three tables, even though `repo_check`’s own per-paper `summary_table` — which covers every one of the 1861 papers — correctly records `files_n = 0` for it. Because the script built its paper list from the three per-file tables rather than from `summary_table`, every such paper was silently absent from the comparison entirely: its `mc_*` value was not `FALSE`, it simply never existed as a row to compare. Confirmed directly: 495 of 1861 papers were missing this way, including 404 papers Cooper et al. coded `data_availability Yes` — 404 real disagreements (cases where Cooper et al.’s coders found a repository Metacheck’s own detection had missed) that were never entering the comparison at all, rather than being counted as a genuine miss. Fix: the paper list is now built from `summary_table$paper_id` (already the full 1861 by construction), and `mc_data_availability`, `mc_code_archived`, `mc_data_README`, and `mc_code_README` are explicitly set to `FALSE` for a paper `repo_check` examined and found nothing in; every other `mc_*` column (license, DOI, download, format, language) is correctly left `NA` for such a paper, since those questions remain genuinely inapplicable without a repository to ask about. Every number reported in the Empirical Comparison section above reflects this correction.

A.2 Re-examining `code_check`'s parse-error and package-version-pinning numbers

A separate line of inquiry, prompted directly by a sentence in an earlier manuscript draft stating that “138 files (1.3%) failed to parse as valid code in their declared language at all,” asked two questions: is this the right denominator, and is the file list behind it available for inspection at all.

Locating the underlying data. `code_check()`'s per-run output already carries everything needed: a `table` with one row per code file (`file_name`, `language`, `parse_error`, `parse_error_msg`) and a `summary_table` with one row per `paper` (`code_version_pinned`, `code_parse_errors`), both already computed by the module — no new code was needed to make this data available, only to locate and aggregate the 38 saved per-batch corpus-run files (`cooper_validation_metacheck/manuscript/data/batches/res_code_check_batch001.RData` through `batch038.RData`) that had already been produced by the original corpus run.

Package version pinning: 34.2% was a batch-level upper bound, not a per-paper rate. An earlier manuscript draft reported that “only 13 of 38 batches (34.2%) contain at least one paper with any form of version pinning,” explicitly flagged in that draft's own text as an upper bound on the true per-paper rate because a batch of ~50 papers needs only one pinned paper to count. Checking directly against the saved per-batch `summary_table$code_version_pinned` column (present in every batch, keyed by `paper_id`) confirmed the per-paper detail this draft believed was not retained had, in fact, been retained all along: aggregating it across all 38 batches gives 18 of 400 code-bearing papers (4.5%) with version pinning present — substantially lower than the 34.2% upper bound, as expected, and now a true per-paper number rather than an aggregate-of-aggregates. No code fix was needed here; this was a reporting correction (using data that already existed but had not yet been extracted at the right level of granularity), not a software defect.

Parse errors: 1.3% used the wrong denominator, and undercounted for a subtler reason. The originally reported 1.3% divided 138 parse failures by 10,461, the total code-file count across every language in the corpus — but parsing is only ever attempted on R-type files (`.R`, `.Rmd`, `.qmd`; `code_parse_r()` is R-only), so the other 3,259 SAS/SPSS/Stata/Python/MATLAB files in the corpus were never eligible to fail this check and should not have been in the denominator at all. Restricting to the corpus's 7,202 R-type files, a further complication emerged: 1,651 of those were never actually downloaded or read at all (an inaccessible repository or a failed download), so parsing was never attempted on them either — they, too, do not belong in the denominator. A further 15 files reached a local copy but errored during extraction (before `parse()` itself was ever called), splitting into two genuinely different causes on closer inspection: 11 were confirmed to be genuinely empty (0-byte) files — fetched directly from their original GitHub repositories to confirm the emptiness is a fact about the authors' own repository, not an artefact of Metacheck's download step (e.g. github.com/olgaviedma/LadderFuelsR's `R/foofy.R` and github.com/dpritchard/frair's `tests/testthat/test_emdII.R`, both independently confirmed 0 bytes via a direct `curl` fetch from `raw.githubusercontent.com`) — and 3 were `.Rmd` files that failed during Quarto/knitr's own chunk-extraction step with

a genuine, structural fence-mismatch error (a closing code-fence with a different number of backticks or indentation than its opening fence), which is the same underlying “not valid in its declared format” problem a parse error represents, just caught one step earlier in the pipeline than `parse()` itself. Restricting to the 5,536 files that were both locally available and actually reached a call to `parse()`, 138 failed (2.5%) — not 1.3%, and not the 1.92% that a denominator of “all 7,202 R-type files regardless of whether they were ever read” would give either. Including the 3 fence-mismatch failures as the same kind of problem gives 141 of 5,539 (2.5%).

This investigation also surfaced a genuine, previously unknown software defect: **`code_read()` (R/`code_check.R`) crashes on a genuinely empty file.** `readr::guess_encoding()` returns a zero-row result for a 0-byte file, so `enc <- readr::guess_encoding(file_path)$encoding[[1]]` silently produces NA rather than a valid encoding string, and that NA then crashes deep inside `readr`’s own internals (`if (stringi::stri_enc_isascii(lines))`), which requires a length-one logical and receives NA instead) with the unhelpful, generic message “missing value where TRUE/FALSE needed” — this is exactly the message recorded for all 11 of the empty-file cases above, rather than a clear “this file is empty” signal. This defect is documented here as confirmed and diagnosed, but **not yet fixed** as of this writing (it affects a small, already-quantified fraction of files — 11 of 7,202 R-type files in this corpus — and a fix has not yet been prioritized ahead of the five `code_check` checks fixed in A.3 below); it is recorded here so the fix, when made, has a clear before-state to compare against.

A.3 Systematic validation of five further `code_check` checks

A third, more systematic validation exercise checked five of `code_check`’s remaining automated checks — missing referenced files, hardcoded absolute paths, `setwd()` calls, comment density, and scattered library/import lines — against real files from the same corpus, reading every flagged file that had already been downloaded locally (8,218 of 10,461 code-file rows) and sampling up to 100 of the remainder that would have required a fresh download. The full method, every sampled file’s judgment, and every quoted example are recorded in a dedicated, reproducible validation report (`code_check_validation.qmd`, with its accompanying script `validate_code_check.R` and judgment tables, all in `cooper_validation_metacheck/manuscript/`) with a fixed random seed (`set.seed(42)`) at every sampling step; this appendix summarizes its findings and the resulting code fixes. Two checks not covered by this validation pass (parse errors, package version pinning) are covered separately in A.2 above.

Missing referenced files — case-sensitivity bug (fixed). `code_check`’s check for files a script tries to open but that are absent from the repository (`inst/modules/code_check.R`, `missing_files <- setdiff(base_ref, base_repo)`) compared filenames case-sensitively. 13 confirmed false positives across the sampled files were all the same mechanism: the referenced file exists in the repository under a different letter case, which is not a real error on the author’s part on the case-insensitive filesystems (Windows, macOS by default) most authors develop on. Concrete examples, each confirmed directly against the paper’s real repository file listing: paper

10_1002_2688_8319_12229 references `REM_tools.r` where the repository holds `REM_tools.R`; paper 10_1111_2041_210x_13923 references `diagnostic_fcns.r/boot_glmm.r` where the repository holds `diagnostic_fcns.R/boot_glmm.R`; paper 10_1111_2041_210x_13125 references `FoodWeb_Functions.R` where the repository holds `Foodweb_functions.R`; paper 10_1002_pan3_10736 references two scripts whose repository copies differ only in the first letter's case. Fix: the comparison now lowercases both sides before comparing (`missing_files <- base_ref[!tolower(base_ref) %in% tolower(base_repo)]`), removing all 13 confirmed false positives found in this sample and any of the same kind not yet sampled.

Missing referenced files — `vroom::vroom()` not recognised as a file-reading call (fixed). The same check's file-reference extractor (`code_file_refs()`, `R/code_check.R`) recognises a fixed list of R reader-function names (`read.*/read_*`, `import`, `fread`, `readRDS`, `load`, `readLines`, `fromJSON`, `readtext`, `source`); `vroom::vroom()` does not begin with “read” and was absent from this list, so a file loaded exclusively via `vroom::vroom()` was invisible to the missing-file check regardless of whether it actually existed in the repository. Confirmed false negative: paper 10_1111_2041_210x_13440's `R/ncbi.R` calls `vroom::vroom(file.path(dir, "nodes.dmp"), ...)` and a second `vroom::vroom()` call for `names.dmp`; neither file exists in the repository, but the check never flagged either, since `vroom` was not on its keyword list. Fix: “vroom” added to the reader-keyword list.

Missing referenced files — remaining known limitations, not yet fixed. Two further false-negative mechanisms were confirmed but not fixed, since each would require a more substantial change than a keyword addition or a comparison tweak: (1) `code_file_refs()` only captures a quoted filename appearing on the same line as a reader-call keyword, so a file assigned to a variable on one line and read via that variable on a later line (`liz_file <- "example_lizard_data.csv"`, then later `read.csv(liz_file)`) is invisible to the check even when the file genuinely does not exist (confirmed: paper 10_1111_1365_2435_13993's `tests/testthat/test-m_run_biophysical.R` does exactly this for three files, none of which exist in the repository). (2) When a referenced filename is built as `paste0(some_variable, "...literal_suffix.csv")`, the check's regex only captures the literal trailing fragment (e.g. `_M_D_abundance_interactions.csv`) rather than the true, variable-interpolated filename — the resulting “missing file” flag is still technically correct (nothing by that literal fragment name exists either), but the name shown to a paper's author is a confusing partial match rather than the file they would recognise; across the corpus, 128 of 2,599 (4.9%) genuinely-missing entries have a name starting with punctuation, the signature of this dropped-prefix pattern.

Hardcoded absolute paths — three false-positive mechanisms fixed, one left open. `code_abs_path()` (`R/code_check.R`) regex-matches quoted strings shaped like an absolute filesystem path. Three confirmed false-positive mechanisms were fixed: (1) a single-segment string beginning with `/` (e.g. `self.fetch('/login')` in a vendored web-app test suite, `github.com_ccrisan_motioneye/tests/test_handlers/test_login.py`) is an HTTP route, not a filesystem path — 322 of 4,182 originally matched strings were 6 characters or fewer, a class dominated by exactly this pattern; the regex's bare-`/`-prefixed branch now requires a second `/`-separated segment before matching, which a real absolute

path (e.g. `/Users/name`) always has and a bare route never does. (2) An all-digit UNC “host” is a `sub()/gsub()` regex backreference string, not a network path — confirmed twice in real code: `osf.io_5fter/IDcallR_sameN_withIDcallPP.R` line 183, `sub('group(\\d+)_(\\d+)', '\\2\\1', colnames(clupo))`, where “`\\2\\1`” was matched as “`\\host\\share`” because a bare digit satisfied the old hostname character class; and a second, independent instance in a Pyrenees-mapping script’s `gsub('(.{1,10})(\\s|$)', '\\1\\n', ...)` call, matched for the same reason. The UNC branch’s hostname character class now requires at least one letter, which excludes both confirmed cases without excluding any real hostname (no genuine UNC host is all-digit-and-punctuation with no letter at all). (3) The pattern’s own pre-existing docstring already flagged a known risk — a variable-built path like `paste0(dir, '/file.csv')` triggering a false match on the literal fragment — and this was confirmed for real: `github.com_YvesBas_Tadarida-C/Sonotypes/write_tabase3HF_Sonotype.R` line 40, `list.files(paste(ListDate[[i]], "/eti", sep=""), ...)`, where “`/eti`” is a relative fragment concatenated onto a loop variable. This third mechanism is addressed by the same “require a second path segment” fix as (1), since `/eti` is itself single-segment. A fourth mechanism, confirmed but **not fixed**: a Python string built by `+`-concatenation from a relative literal (`'data\\' + tile_name + '\\run\\' + model_name + ...`, from a Zenodo-hosted `model_test.py`) is a relative path whose middle segments (`\\run\\`, `\\output\\tile\\`) still match the UNC branch on their own; no fully static regex-only fix exists for this case, since it would require tracking whether the enclosing expression is a string concatenation starting with a relative literal, which is beyond what a line-based regex can determine. Of the true positives sampled for comparison, 15 of 15 further-checked “real username-shaped” matches (e.g. `C:/Users/reddin/Documents/...`, `/Users/farrer/Dropbox/...`, `/home/michael/Dropbox/...`) were confirmed genuine, machine-specific absolute paths — the large majority of this check’s flags are correct.

setwd() calls — no false positives found in this corpus; one confirmed-but-unobserved risk noted. All 194 distinct `setwd()` call strings found across the sampled files are genuine, live, executable calls, including calls whose argument is itself relative or dynamic (`setwd("./data")`, `setwd(dir)`, `setwd(tempdir())`) — these correctly count as flags under this check’s own stated rule, since the anti-pattern being flagged is changing the working directory at all, not specifically a hardcoded absolute-path argument. A latent false-positive mechanism was confirmed by construction (a synthetic test, not a real corpus case): `code_setwd()` has no awareness of string-literal context, so a line like `message("Please run setwd(your_path) before continuing")` would be matched as if it were a live call. Scanning all 194 real call strings in the corpus for instructional-text markers (“please,” “before,” “instructions,” “note:”) found zero real occurrences of this pattern — a real, demonstrated risk with no confirmed instance in this corpus, and left unfixed for that reason, though `code_setwd()` could straightforwardly reuse the same quote-context-tracking logic already implemented for `.code_strip_inline_comment()` if a future corpus surfaces a real case.

Comment density — a Python-specific docstring signal added (not a bug fix). `code_remove_comments()`’s Python branch deliberately never treats a triple-quoted string as a comment, since doing so would also destroy real docstrings (themselves ordinary string

literals, syntactically). The consequence, quantified directly: of 280 zero-comment-flagged Python files in the corpus, 200 (71.4%) contain at least one complete triple-quoted docstring block — real documentation that the mechanical `#`-only comment count cannot see, and so reports as “0% comments” regardless. Five were read directly to confirm these are genuine module/function docstrings rather than incidental triple-quoted data, e.g. `model_runs.py`’s `"""Filters for Model Runs."""` and `features.py`’s 16 separate triple-quote markers (8 docstrings) opening with `"""API functions to interact with feature names."""`. This is a deliberate, correctly-reasoned design choice in `code_remove_comments()`, not a bug — stripping docstrings to count them as comments would also delete them from every downstream check that reads code text — so the fix made is additive rather than corrective: a new function, `.code_has_docstring()`, detects whether a file contains at least one complete triple-quoted block, and `code_line_stats()` now returns this as a separate `has_docstring` field (NA for every non-Python language) alongside the unchanged `percent_comments` metric. `inst/modules/code_check.R`’s report text now adds an explanatory note whenever a zero-comment Python file in fact has a docstring, so a reader is told the file is documented via docstrings rather than concluding it has no documentation at all.

Scattered library/import lines — Python import-anchor tightened (fixed). `code_library_lines()` (`R/code_check.R`) previously anchored its Python import-detection regex on the bare keyword plus a following word (`^\s*(import|from)\s+[A-Za-z_.]`), with no requirement that the rest of the line actually form a valid import statement. Because Python’s comment-removal never strips triple-quoted strings (the same mechanism behind the comment-density finding immediately above), docstring prose that happens to start with “from” or “import” survives into this scan and was matched as if it were a real import line. Concrete, exact example: `github.com/opencv/opencv4/site/process_sdk_docs.py`, line 213, inside a docstring spanning lines 211–214 — the prose sentence “from the generated model and api descriptions.” was matched as an import line at line 154 of the comment-stripped text, while the file’s 10 real imports are all genuinely contiguous at lines 1–10; the spurious match alone widened the reported gap between imports to 144 lines, crossing the “scattered” threshold (more than 3 non-comment lines between library statements) on a file whose real imports are, in fact, perfectly grouped. Quantified across all 525 Python files flagged as scattered in the corpus: 155 (29.5%) have at least one spurious, non-import match of this kind; of those, 68 (13.0% of all 525 flagged files) would no longer be flagged as scattered at all once the spurious matches are excluded — pure false positives caused entirely by this mechanism (the remaining 87 files with a spurious match are still genuinely scattered even after excluding it). R showed essentially no equivalent problem in a 60-file sample (the one apparent candidate, on manual inspection, was a genuine conditional `require()` call, not a false positive — R’s regex requires an opening parenthesis directly after the keyword, structurally harder to spuriously trigger than Python’s bare whitespace-plus-word anchor). Fix: the Python pattern now requires the full import shape rather than just the leading keyword — a `from` line must also contain its own `import` keyword later on the same line (`^\s*from\s+[A-Za-z_][A-Za-z0-9_]*\s+import\s`), and a bare `import` line must be followed by a comma-separated identifier list with an optional `as name` and nothing else (`^\s*import\s+[A-Za-z_][A-Za-z0-9_]*\s*(\s*,\s*[A-Za-z_][A-`

`Za-z0-9_.]*)*(\s+as\s+\w+)?\s*$)` — verified against the docstring false-positive example above (now correctly unmatched) and against every real import shape found in the corpus sample, including multi-line `from x import (a, b, c)` continuations (matched on their opening line) and indented imports, with no loss of any real detection.

Status of the fixes in this appendix. All Metacheck fixes described in A.1 and A.3 (nine in total: #378, #379, #380, #383, both causes of #384, and the five `code_check` fixes from the systematic validation) have been committed to `scienceverse/metacheck`’s `repocheck_improvements` branch (pull request #385) and are covered by the package’s automated test suite (463 tests passing at time of writing, 0 failures, 3 pre-existing skips unrelated to these changes). The sixth fix described above (the comparison script’s missing-repository default) is not a Metacheck fix and has been committed instead to this project’s own `code_data_preparation/03_recreate_cooper_columns.R`, in the `cooper_validation_metacheck` repository. The one defect identified but not yet fixed (`code_read()`’s crash on an empty file, A.2) and the two `code_file_refs()` limitations noted but not yet fixed (A.3) remain open. The further full corpus rerun incorporating the five A.3 fixes, planned when this appendix was first written, has since been performed (confirmed directly: the current corpus run’s `code_check` output carries the `has_docstring` field A.3’s docstring fix introduces, which cannot exist in a pre-fix run) — the numbers reported in “Results you get for free alongside Metacheck” above reflect that post-fix corpus, not the pre-fix state described earlier in this sentence’s original drafting. This was a distinct item from the rerun described in Appendix A.5 below, which re-checked only the papers with a remaining disagreement (not the full corpus) and tested a different set of fixes (Zenodo reliability, repository-platform coverage, and rate-limit handling), not the five A.3 `code_check` fixes.

A.4 A repository excluded from processing by file count

During the corpus rerun incorporating the A.3 fixes, one paper’s data repository — Zenodo record 10.5281/zenodo.3971689, SAPFLUXNET (a global database of sap flow measurements: 202 datasets from 121 locations, sap flow and environmental time series for 2,714 individual plants) — was found to hold approximately 4,000 individual files. Downloading and, in particular, `data_check`’s per-file classification and full-content read of every one of those files (`data_check_head(..., n_rows = Inf)`, by design, since several of `data_check`’s validation checks need the complete dataset rather than a sample) took several hours and multiple gigabytes of memory for this single paper, out of 1,861 in the corpus, before the run was stopped.

This exposed a real gap in Metacheck: `download_repo_files()` already gates a repository outright by total byte size (`max_download_size`) or per-file byte size (`max_file_size`), but had no gate on file count — a repository holding thousands of individually small files was not caught by either existing cap. A new parameter, `max_files_per_repo` (`R/repo-download.R`, threaded through `data_check()` and `code_check()`), gates a repository by file count alone,

using the same refusal/reporting mechanism as the existing size caps. This has been committed to the `repocheck_improvements` branch alongside a regression test, and the corpus rerun was configured with `max_files_per_repo = 1000`.

Under this cap, the SAPFLUXNET repository is refused outright rather than downloaded, so this one paper’s `data_check/code_check` results are incomplete (no files analysed for it) in the corpus run reported in this manuscript. This is a deliberate, scale-driven exclusion of a single outlier repository, not a defect in how that paper’s data was assessed — the cap exists specifically so that one exceptionally large repository cannot block or degrade processing of the other 1,860 papers in the corpus. Revisiting this one paper’s repository with a raised or absent cap, as a separate, deliberate pass, remains an open option.

A.5 A targeted rerun on every remaining disagreement, after further platform-detection and reliability fixes

After Appendix A.1–A.4’s fixes and the resulting full corpus rerun, every remaining disagreement between Metacheck and Cooper et al.’s coding (991 rows across 989 papers, the same set individually adjudicated for the Empirical Comparison above) was re-checked a further time against Metacheck as it stood after seven more commits landed on `scienceverse/metacheck`’s `dev` branch: two fixing crashes and false negatives in Metacheck’s Zenodo integration, three widening which repository platforms `repo_check` can locate a repository on at all (targeting “repository not detected,” the single largest recurring disagreement cause identified in the original review — see “The five recurring causes” above), and two fixing a rate-limit reliability defect surfaced live during this rerun itself. Rather than reprocess the full 1,861-paper corpus, only the 989 papers with at least one open disagreement were rerun, one paper at a time (so that a single paper’s slow host response could not block any other), reusing the shared on-disk repository-file cache from the original corpus run so that a repository already fetched and unchanged cost no new download.

Zenodo reliability. Two fixes addressed real crashes and false negatives in Metacheck’s Zenodo integration, confirmed against Zenodo records already present in this corpus. `online()`, the function `repo_check` calls to confirm a host is reachable before querying it, treated a single transient DNS lookup miss as “Zenodo is offline” and abandoned that paper’s entire Zenodo retrieval rather than retrying; it now retries. Separately, a JSON metadata field arriving as a length-zero value (an empty array) rather than `NULL` defeated the `%||%` operator used throughout Metacheck’s metadata-extraction code (which only guards against `NULL`, not length-zero), producing a `data.frame()` “arguments imply differing number of rows” error that was silently swallowed inside a `tryCatch`, discarding the whole record; a new `%empty_or%` operator, applied everywhere the same gap existed (Dataverse, Dryad, DSpace7, Figshare, GitHub, GitLab, OSF, PsychArchives, and ReShare, not just Zenodo), closes this. A further, closely related pair of crashes was found and fixed the same way: a Zenodo record Metacheck could not read at all previously returned a `data.frame()` with only `zenodo_id/error` columns, so a paper whose every Zenodo record was unreadable produced a metadata table with no

`doi/license` columns at all, which crashed `repo_check()`'s own Zenodo-metadata-extraction step for that paper the same “differing number of rows” way; and a `repo_url` arriving as a list column (rather than a plain character vector) could silently explode `data.frame(zenodo_url = <list>)` into one column per list element instead of the expected single column. Both are now guarded.

Repository platform coverage. Three commits substantially widened the set of hosts Metacheck's `repo_check` can locate a repository on at all, directly targeting the “repository not detected” cause responsible for the largest share of the original disagreement (see “The five recurring causes” above, which named DataSTORRE, DIGITAL.CSIC, AADC, SRA/GenBank, KNB, Neotoma, CIRAD Dataverse, and ForestPlots.net as example platforms with no backend at the time). Mendeley Data (detected via `10.17632/` DOIs and `data.mendeley.com` URLs, queried through the public, unauthenticated `data.mendeley.com/public-api/datasets/<id>` endpoint) and DataONE (a federation of independently-run repositories rather than one hosted service; four member nodes verified live and wired in — Arctic Data Center, KNB, Taiwan Forestry Research Institute, and the Smithsonian Tropical Data Portal, each mounting the same underlying REST API at a different base path) were added as new platforms entirely. Separately, Dataverse's recognised-installation list was expanded from 50 to 121 hosts, 16 institutional Figshare white-label front ends were added, and DSpace support was generalized from PsychArchives specifically to any host running the same legacy (6) API (5 more verified live) plus a new implementation covering 64 DSpace 7 installations (e.g. Cornell's eCommons, Georgia Tech's Digital Repository, ETH Zurich's Research Collection) — every one of these $121 + 16 + 5 + 64$ hosts confirmed live via its own version/info endpoint before being added, not merely listed from a registry. A further, smaller fix excludes an R package's own source tree (`DESCRIPTION/NAMESPACE/its R/man/tests/` subdirectories) from a repository's data/code file listing, on the reasoning that reusable software infrastructure bundled alongside a study's own materials is not itself the study's data or analysis code and should not inflate `files_n/files_code` or appear in downstream reports.

A reliability defect found live during this rerun. While this rerun was in progress, a single paper stalled repeatedly at the same point across two separate attempts: `zip_peek()`'s HEAD request and its Range-request follow-up (`.http_range_tail()`, `R/zip-peek.R`), used to read a remote zip archive's file listing without downloading it, each independently hit Zenodo's rate limit and each, independently, paid the full announced wait for the host's confirmed reset — three consecutive confirmed-rate-limit waits back to back for what was, from Zenodo's own side, a single rate-limit window, since neither request had any way to know the other had already discovered and waited out the same limit. (The rate-limit-wait message itself was also, at the same time, made to name the specific host responsible — Zenodo, Dryad, GitHub, or GitLab, the only hosts that expose a machine-readable reset time — rather than reporting only “the host's own reset,” a smaller, secondary fix made in the course of diagnosing this one.) Fix: a session-scoped record of the last confirmed rate-limit reset per host, written the first time any request discovers one and consulted by every subsequent request to that host before it is even sent, so a host already known to be rate-limited waits out only the genuinely remaining time once, rather than every independent request rediscovering the identical window

from a fresh 429. Verified with new regression tests covering the record/lookup logic directly (round-tripping a wait, keeping the later of two recorded resets, clearing an expired window, and the pre-flight wait actually sleeping the correct duration) and confirmed live: the same paper that previously stalled for over 40 minutes across three stacked waits resolved in roughly 3 minutes with zero confirmed-rate-limit waits after the fix, the remaining time attributable to Zenodo’s own server-side response latency for that specific record (confirmed independently: a single isolated request to the same record, made outside Metacheck entirely, itself took 30 seconds and returned a 504 on a second attempt).

Results of the rerun. Across the 1,532 individual construct-level disagreements re-checked (989 papers $\times$ up to 7 constructs each, matching however many constructs a given paper had originally disagreed on), 76.8% (1,176 of 1,532) are now resolved — Metacheck’s fresh result agrees with Cooper et al.’s original human coding where it did not before — with 348 still disagreeing and 1 now returning a value that could not be classified as agreeing or disagreeing (`data_license`, a repository whose licence field this rerun could not read for reasons not yet diagnosed). By construct: `any_readme` 548/625 resolved (87.7%), `data_availability` 339/424 (80.0%), `code_archived` 252/360 (70.0%), `data_archive` 12/56 (21.4%, evaluated by set membership — a paper’s platform coded by Cooper now appearing among Metacheck’s own, possibly multi-platform, detected set, since a paper can genuinely link more than one repository), `data_license` 16/57 (28.1%), `data_download` 9/9 (100%), and `code_download` 0/1 (the sole remaining `code_download` disagreement in this corpus). The comparatively modest improvement in `data_archive` and `data_license` is consistent with the original review’s own finding that these two constructs’ disagreement was, unlike `data_availability/code_archived/any_readme`, not dominated by the repository-detection cause these fixes targeted (`data_license`’s largest cause was, per “Opposite direction” above, cases where Metacheck was already right and Cooper’s own coding had missed a real licence field) — a fix aimed at finding more repositories naturally does less for a construct whose disagreement was rarely about a missing repository in the first place. This rerun’s full row-level detail (every paper, every re-checked construct, the value before and after, and whether it resolved) is recorded in `cooper_validation_metacheck/manuscript/data/rerun_disagreements_comparison.csv`, produced by `rerun_disagreements.R`; the underlying per-paper `repo_check/data_check/code_check` results are cached per paper in `data/rerun_batches/` for inspection or a further rerun without recomputing what has not changed.

Checking the rest of the corpus for regressions. This rerun was scoped to the 989 previously-disagreeing papers only — a deliberate choice at the time, to get a fast answer to “did these fixes help” without the cost of a full 1,861-paper rerun (see the rerun’s own header comment for the reasoning). That scoping left one real question open: none of these seven fixes were written to only affect a disagreeing paper, so it remained possible that a fix which resolved one paper’s disagreement had, in a different paper that already agreed with Cooper et al., turned a correct result into an incorrect one — most plausibly the R-package-source-tree exclusion (if a paper’s own analysis code happens to be organised as an R package) or the expanded Dataverse/Figshare/DSpace/Mendeley/DataONE host recognition

(a newly-recognised host matching something that should not have counted). This was checked directly: the remaining 868 papers (every paper not in the 989-paper disagreement set) were also rerun, using the identical per-paper, resumable approach, and the full corpus's `cooper_vs_recreated.RData/disagreement` worklist were then rebuilt from the combined result (all 1,861 papers) rather than just the 989. At the level of aggregate agreement rates, this full-corpus check was reassuring: the three constructs this comparison's own disagreement was concentrated in — `data_availability`, `code_archived`, `any_readme` — all showed a higher, not lower, agreement rate once the previously-agreeing papers were included, and `data_download/code_download` (both already above 99%) moved by less than a percentage point in either direction. But an aggregate rate rising is consistent with a handful of individual papers having genuinely regressed, provided enough other papers improved to more than offset them — and that is exactly what the case-by-case review below found.

A confirmed regression, found and fixed within this same cycle. The suspicion named above turned out to be correct: the R-package-source-tree exclusion did break papers it should not have touched. The exclusion was written to detect an R package's fingerprint (a `DESCRIPTION` and `NAMESPACE` file in the same folder) and drop that folder's contents — the package's `R/`, `man/`, `tests/`, `vignettes/`, `data/`, `inst/`, and `src/` subdirectories — from a repository's file listing entirely, on the reasoning that a bundled, reusable software dependency is not itself a study's data or analysis code. That reasoning holds when the package is a genuine dependency vendored alongside a study's own files, but it does not hold when the paper's own contribution — its data, its analysis code, or both — is itself packaged and distributed as an installable R package, a common practice in ecological/evolutionary methods work. In that case the exclusion strips the paper's own materials, leaving nothing behind: `repo_check`'s listing for the affected repository showed only packaging scaffolding (`.Rbuildignore`, `.gitignore`, `NEWS.md`, `README.md`, `cran-comments.md`, a bare `.Rproj` file), with the package's real `R/` source — and, for packages that bundle their own example or study data under `data/`, the actual dataset — entirely absent.

A corpus-wide scan for this exact symptom (a repository whose only surviving files, after the exclusion ran, were scaffolding, with zero files classified as data or code) surfaced candidates across the full 1,861-paper corpus, each individually verified against the repository's real GitHub listing (via `gh api`) or Dryad/Zenodo record before being counted: 16 confirmed true instances (`RevGadgets`, `BaBA`, `m2b`, `RRPP`, `actel`, `scampr`, `MINOTAUR`, `LadderFuelsR`, `biomonitoR`, `EcoEnsemble`, `fossilsim`, `frair`, `tRophicPosition`, `sabinaNSDM`, `bef`, and `phylomice`), plus several candidates from the same scan that turned out to be unrelated causes once checked directly (a repository with no `DESCRIPTION/NAMESPACE` at all, and several Dryad/Zenodo-hosted papers whose `code_archived` disagreement traced to an unconnected, pre-existing detection gap — code cited in a journal's own Supporting Information, or a repository never located by `repo_check` at all). Two confirmed cases (`BaBA`, a GitHub-hosted barrier-behaviour-analysis package whose `data/` folder holds the actual tracked-animal datasets for its own paper, and `MGDrive`, a general-purpose mosquito gene-drive simulation package vendored inside a study-specific fork) between them illustrated the structural distinction the fix needed to draw: `BaBA`'s package sits at its repository's root with nothing else present — the package is the deposit, so

nothing should be excluded — while `MGDrive`'s vendored copy is nested one level deep inside the fork, with the study's own `Examples/` scripts and data sitting as real, undisputed siblings at the repository root.

The corrected rule preserves the original exclusion's logic for a nested package (still always excluded, since a nested package by construction always has something else in the repository for it to sit alongside) but adds one check for a package detected at a repository's root: before excluding it, check whether anything outside the package folder classifies as real data or code, using `data_classify_files()` — the same rules-based classifier `data_check` itself uses elsewhere in the pipeline, rather than a separate, bespoke list of what counts as “real content.” If nothing does, the package is not excluded at all, since there is no “outside” for the study's own materials to live in. One implementation detail required a second pass: an RStudio project marker file (`*.Rproj`) classifies as “code” under `data_classify_files()` purely by its file extension, which would have wrongly counted as “real content survives” for every package that ships one (nearly all of them) — an RStudio project file is packaging configuration, not an analysis script, so it is excluded from this specific check. The fix was verified against two new regression tests (a root-level package with nothing else present, confirming it is no longer excluded, and a nested vendored package alongside real sibling content, confirming the vendored files are still excluded while the sibling content survives) alongside the existing test suite, then reapplied to the 38 papers the corpus-wide scan had flagged.

Reapplying the corrected exclusion resolved 11 of the 16 confirmed cases outright. The remaining 5 continue to show `code_archived` disagreement, but for a distinct and unrelated reason discovered while reprocessing them: their cached downloaded files never included the package's R/ source directory in the first place — only packaging scaffolding was ever fetched from the host — so the corrected exclusion logic correctly finds nothing outside the (neverfully-downloaded) package to preserve, because nothing beyond scaffolding was ever retrieved. This is a fetch-completeness gap, not an exclusion-logic problem, and is not further diagnosed here; it is recorded as a separate, still-open limitation rather than folded into a claim about the exclusion fix's own effectiveness. Two of the 16, `phylomice` and `morton` — cases surfaced by the same corpus-wide scan but outside the original 13-paper set found during the initial regression check — resolved to, respectively, `METACHECK_RIGHT` (`phylomice`'s six R source files are a genuine methods contribution the paper's own text never explicitly claims as “archived code,” but which Metacheck is nonetheless correct to report) and `COOPER_RIGHT` (`morton`'s 37 R source files are unambiguously this paper's own phylogenetic-analysis code, which Cooper et al.'s coders had missed).

A.6 A systematic platform-coverage audit, three further metacheck fixes, and a complete categorization of every remaining repository-detection disagreement

A sixth round of work started from a single case the previous rounds had not caught: a paper citing a Dataverse installation (`borealisdata.ca`, Borealis, the Canadian national repository) under a DOI prefix Appendix A.5's Dataverse allowlist did not recognise. Tracing why surfaced

a systemic problem in how that allowlist had been built, and prompted a full re-audit of every remaining “repository not detected” disagreement in the corpus rather than a single-paper patch.

A systemic under-sampling bug in the Dataverse DOI-prefix allowlist. Appendix A.5’s per-host Dataverse DOI-prefix allowlist (`.dataverse_doi_prefix_hosts()`) was originally built by sampling the first 25 DOIs DataCite’s search API returned for each host — a sample, not a census. This badly under-counted any host with more than 25 real DOIs registered, which turned out to be the majority of the 121 known hosts; `borealisdata.ca` alone has minted DOIs under six distinct prefixes, of which the original allowlist recorded only one. Rebuilding the allowlist using DataCite’s facet aggregation (`disable-facets=false` on its `/dois` search, which returns every distinct prefix in use for a host, with its real count, in one request rather than a sample) found 100 distinct prefixes across the same 87 hosts, not 87 — including a second Harvard Dataverse prefix (10.26300, alongside the dominant 10.7910) that the original sample had also missed. Verified live before and after: the DOI that surfaced this bug (10.7939/DVN/M6Q2PE, Borealis) now resolves correctly, and the second Harvard prefix now resolves to `dataverse.harvard.edu`. This was reported and merged upstream as `scienceverse/metacheck#410`.

Two further platform-coverage gaps found and fixed the same way. `.figshare_id()` did not recognise a Figshare article URL with no resource-type segment (`figshare.com/articles/<slug>/<id>`), as opposed to `figshare.com/articles/dataset/<slug>/<id>` — an older, still-current URL shape (confirmed live against a real corpus citation) that `figshare_links()` already detected as a Figshare link but could not resolve to a usable article id, so `repo_check` never actually queried it (`scienceverse/metacheck#412`). Separately, `dataone_links()` had no entry for `data.piscoweb.org` (PISCO, a UC Santa Barbara-run DataONE member node), and `.dataone_pid()`’s bare-DOI regex could not handle a DOI suffix containing a further slash — PISCO’s own DOI shape (10.6085/AA/<id>) — so this platform’s citations, while real DataONE member-node DOIs, were invisible end to end. Both fixed and merged as `scienceverse/metacheck#414`, which also added 29 further institutional Figshare instances (28 found by searching DataCite’s client registry for “figshare” in the client name, plus Monash University’s `bridges.monash.edu`, found separately since its DataCite client is registered under the institution’s own name) — each confirmed live, the same way as every host already in `.figshare_vanity_hosts()`, to answer Figshare’s own single-page-app shell rather than a real API response, and each institutional DOI prefix confirmed live to resolve to that exact host before being added.

Two matching bugs in this project’s own comparison code, not in Metacheck. `03_recreate_cooper_columns.R`’s `.repo_url_platform()` — the function that labels which platform `mc_data_archive` reports for a paper — had a Dryad branch that only matched the single most common DOI prefix (10.5061), and a Figshare branch that only matched the literal `figshare.com/10.6084` domain and prefix, missing every one of the alternate Dryad prefixes and institutional Figshare instances Metacheck itself

already recognised after the fixes above. The practical effect was a paper whose repository was now correctly found and retrieved still showing a `data_archive` disagreement, because the label fell through to the generic “Other repo/database” instead of naming the real platform. Both branches were rewritten to read directly from Metacheck’s own `.dryad_doi_prefixes()/.figshare_vanity_hosts()/.figshare_doi_prefix_hosts()` rather than duplicating either list by hand, so this classification cannot drift out of sync with what Metacheck itself recognises again.

Reprocessing every paper these fixes could plausibly affect. Rather than assume the fixes above were the only remaining cause of repository-detection disagreement, every “repository not detected” comment still open in the worklist (351 individual (paper, column) disagreements across 233 papers, after excluding platforms already confirmed unsupported for an unrelated reason) was checked again: first narrowed locally, with no network calls, to the papers whose cited DOI prefix genuinely matched one of the platforms just fixed (27 candidates, after excluding two DOI prefixes that coincidentally collide with an unrelated platform’s registered prefix in DataCite’s own data — `10.6084`, Figshare’s own prefix, and `10.17617`, Edmond’s — neither of which any fix above touched), then live-verified one at a time. Of those 27, 13 were genuinely resolved by the fixes above; the remaining 14 were not — some detect the citation correctly but the dataset is embargoed or restricted at the source (confirmed via the repository’s own API returning an explicit “cannot be viewed” message, not a Metacheck-side failure), and a few still do not detect the citation at all, for reasons unrelated to any fix in this round. A further, separately-run check across every remaining “never located” comment surfaced one exact repeat of Appendix A.1’s original zip-peek limitation on a much larger scale (24 papers whose Zenodo/Figshare/Dataaverse archive was never successfully peeked during an earlier corpus build, confirmed live, unmodified, to peek correctly today — a stale corpus row in every one of the 24, not a code defect). A further paper, initially suspected to be a second such stale row (its earlier 4TU fetch crash was already fixed in Appendix A.1 and appeared simply never rerun since), was checked directly against `rerun_disagreements_comparison.csv`’s own before/after values and found to still disagree after being rerun — the 4TU fetch failure was not, in fact, resolved. A later, more thorough recheck (Appendix A.9) traced the actual cause: 4TU is a genuinely supported platform, but its DOI-embedded identifier does not match the numeric article ID 4TU’s own API requires, a distinct, still-open defect from the one Appendix A.1 fixed. This paper is accordingly counted as `SUPPORTED_PLATFORM_BUG`, not `UNSUPPORTED_REPO`, in Appendix A.9’s categorization. All resolved candidates from this round were rerun end to end (`repo_check`, `data_check`, `code_check`) and patched into the corpus; none of this round’s rerunning touched a paper outside the specific set each check identified, so no regression check across the unaffected remainder of the corpus was needed the way Appendix A.5’s broader platform-coverage changes required.

A.7 A full-corpus reproducibility_check run: method, outputs, reasons for non-reproducibility, and suggested improvements to Metacheck

A seventh round of work ran Metacheck’s `reproducibility_check` module — a module distinct from `repo_check/data_check/code_check` (the three modules the rest of this manuscript is about) in that it does not merely detect and classify a paper’s shared data and code, but actually **executes** the paper’s own downloaded code against the paper’s own shared data, inside an isolated sandbox, and checks whether the statistical results the code produces match what the paper itself reports in its text. This appendix documents the method, every output file this run produced, a full breakdown of why papers were and were not reproducible, four new Metacheck defects this run surfaced (all reported upstream, two confirmed fixed on `dev` before this appendix was written), and a set of concrete suggestions for further improving `reproducibility_check` itself.

A.7.1 Method

`reproducibility_check(paper, execute = TRUE, sandbox = "docker", install_missing = TRUE, timeout = 3000, cache = TRUE, workers = 5, results_dir = ...)` was run once per paper across the full 1,861-paper corpus. For each paper with at least one code file, the module: (1) reuses `code_check`’s dependency extraction and `psychds_check`’s path-rewrite plan to materialise a throwaway copy of the paper’s own Psych-DS-layout data and code; (2) installs the paper’s declared package dependencies (CRAN, GitHub, or URL sources) into a throwaway library inside a Docker container, starting from the pre-built `ghcr.io/scienceverse/metacheck_r:latest` image (roughly 750 of the most common packages found across a prior corpus scan already installed, so most papers skip most of the install step); (3) runs each of the paper’s scripts, in an automatically-derived dependency order, each in its own `docker run --network none --read-only --user <non-root>` invocation — network disabled and the filesystem read-only outside the materialised sandbox, a genuine containment boundary for running downloaded, unvetted third-party code, unlike the module’s alternative `sandbox = "process"` backend (a bare `callr` subprocess, which isolates a crash but not the filesystem or network); and (4) extracts any statistical output the executed scripts printed to the console and compares it against the paper’s own reported results (`match_reported_output()`), recording how many of the paper’s own reported test statistics the executed code actually reproduced.

`workers = 5` runs up to five papers’ Docker containers concurrently rather than one at a time, a feature added to Metacheck (`dev` branch, version 0.3.1) specifically to make a corpus-scale run practical; each concurrent container is automatically capped to `1/workers` of the host’s CPU and memory, so five containers running at once cannot collectively exceed what the host machine has. `results_dir` writes each paper’s complete result to its own file the moment that paper finishes, independent of any outer batch boundary, so a

run interrupted partway through does not lose the papers that already completed. A dedicated wrapper script (`reproducibility_check/run_reproducibility_check.R`, `cooper-bes-validation-manuscript` repository) drives the corpus-wide call, installing Metacheck fresh from `scienceverse/metacheck`'s `dev` branch at the start of every run (never a locally vendored copy, so the run always reflects the current state of the package's own repository) and writing one checkpoint file per outer batch so a killed or crashed run can resume from the last completed batch rather than restarting from scratch.

A.7.2 Outputs stored

The run produced the following files, all in `cooper-bes-validation-manuscript/data/` unless noted:

- **`reproducibility_check_results_execute.RData`** — the final, combined result for all 1,861 papers: an out object with `out$results$table` (4,394 rows, one row per script attempted across every code-bearing paper: `paper_id`, `file_name`, `outcome`, `error_type`, `undefined_var`, `not_runnable_reason`, `runnable`, `parses`, `run_order`, `paths_rewritten`, `paths_ambiguous`, `setwd_removed`, `unresolved_inputs`, `reads`, `writes`) and `out$results$summary_table` (1,861 rows, one row per paper: `paper_id`, `repro_code_n`, `repro_runnable`, `repro_missing_inputs`, `repro_deps`, `repro_tests_reported`, `repro_tests_matched`, `repro_ran_ok`). This is the file the numbers below are computed from.
- **`data/batches/`** — one `.RData` checkpoint file per outer batch processed during the run (batch-size-tagged in its filename, so a checkpoint from one `--batch-size` setting is never mistaken for one from another), and, under `data/batches/worker_results/`, one `.rds` file per paper (`capture_module_tables()`'s own per-paper checkpoint format), the finer-grained safety net `results_dir` provides underneath the outer batch files.
- **`reproducibility_check_infra_rerun.RData`** and **`reproducibility_check_13384_solo_rerun.RData`** — two small, targeted follow-up runs (below), covering respectively the 16 papers whose first-pass result was an infrastructure failure rather than a genuine check, and one paper rerun alone (no concurrency) after crashing twice under `workers = 5`.
- **`reproducibility_check_execute_warnings.log`** — every R warning raised during the run, forwarded here as it happened (R's own default non-interactive behaviour silently drops all but the first 50 warnings of a session, which a run this size would exceed many times over).
- Two standalone wrapper scripts alongside the main run script, both in `reproducibility_check/`: **`rerun_infra_failures.R`** (the 16-paper targeted rerun) and **`rerun_single_paper.R`** (the one-paper, no-concurrency rerun).

A.7.3 Scope: which papers had code to check at all

Of the 1,861 papers in the corpus, 343 (18.4%) had at least one code file for `reproducibility_check` to attempt (`repro_code_n` > 0); the remaining 1,517 (81.5%) had none, and are not further discussed in this section, since there is nothing for the module to execute — this is the same “does the paper share any code” question `code_check/repo_check` already answer elsewhere in this manuscript, not a new finding. One paper (10_1111_2041_210x_13384) has no genuine `repro_code_n` value at all (NA); this is the one paper described in A.7.5 below whose check was manually aborted rather than completed.

Note on this count. This section’s original run (below) found code for 334 papers. Appendix A.8 describes three further Metacheck fixes (`zip_peek()` failing on S3-backed hosts, its result never persisting across a restart, and a silently discarded download error) that, once applied and the corpus rerun, let `repo_check/code_check` correctly find and download a code repository, or a code file inside one, for papers where they previously could not — 98 papers corpus-wide whose available code or data content genuinely changed. Those 98 papers’ `reproducibility_check` result was rerun against the same fixed Metacheck version and merged into this section’s own combined result; every other code-bearing paper’s original result, from before Appendix A.8, is unchanged. 9 of the 98 newly show `repro_code_n` > 0 where they previously showed none, bringing the code-bearing total from 334 to 343; `repro_ran_ok` improved for 5 of the 98 rerun papers (from 138 to 154 papers corpus-wide with at least one script that ran without error). The numbers throughout the rest of A.7.3–A.7.4 below reflect this merged, current result.

A.7.4 Reasons for non-reproducibility: a full breakdown

Across the 343 code-bearing papers, 4,394 individual scripts were attempted (one paper’s script count was not resolved; see A.7.5). Their outcomes:

Outcome	Count	% of 4,394
<code>ran_ok</code> (executed without error)	1,993	45.4%
<code>errored</code> (executed, but raised an R error)	1,879	42.8%
<code>skipped_missing_inputs</code> (an earlier script’s output this script needs was never produced or found)	459	10.4%
<code>not_parsed</code> (the file is not valid R/syntax in its declared language at all)	49	1.1%

Outcome	Count	% of 4,394
<code>timed_out</code> (still running when the 3,000-second per-script cutoff was reached)	14	0.3%

At the paper level, only 85 of the 343 code-bearing papers (24.8%) had every one of their scripts run without error; 136 (39.7%) had zero scripts run successfully at all; the remaining 122 (35.6%) had a mix of successful and failed scripts within the same paper.

Among the 1,879 errored scripts (plus the 14 `timed_out` scripts, grouped in here since both carry an `error_type`), `error_type` breaks down as:

- **637 (33.7%) `undefined_variable`** — the script references a variable, object, or function that was never defined at the point it is used, the classic signature of a script written to be run only as one step within a larger, undocumented pipeline (e.g. an interactive session, or a preceding script the paper’s repository does not include or that this run’s automatic dependency ordering could not place before it). The single most common such name, by a wide margin, is the pipe operator `%>%` **itself (15 occurrences)** — a script using `%>%` without ever calling `library(magrittr)`, `library(dplyr)`, or (on R versions before 4.1) without the base pipe `|>` available, a specific and mechanically fixable non-reproducibility cause distinct from the rest of this category’s more idiosyncratic missing-object cases. `.jinit` (13 occurrences, from `rJava`-dependent packages needing a Java runtime this Docker image does not provide) is the next most common, followed by a long tail of paper-specific object and function names.
- **1,242 (65.6%) `runtime`** — every other kind of R error raised during execution: a missing package the install phase could not provide (see `repro_deps` below), a file the script expects that a path-rewrite could not correctly relocate, a genuine bug in the paper’s own analysis code, or code that assumes an interactive session (see the two concrete cases below).
- **14 (0.7%) `timeout`** — see the 3,000-second cutoff above; a timeout is not itself counted as strong evidence of non-reproducibility (the paper’s own analysis may simply be legitimately slow, e.g. a Bayesian model or a large agent-based simulation — several confirmed cases during this run were exactly this), just “did not finish within the budget this run allowed.”

Two concrete, directly observed non-reproducibility patterns are worth naming explicitly, since neither is visible in the `error_type` classification above (both were classified `runtime`, and the module does not currently distinguish “this code assumes an interactive terminal” as its own category — see A.7.6’s first suggestion): (1) a script calling base R’s `read.csv()` with no path argument, which opens an interactive `file.choose()` file-picker dialog — impossible inside a headless, network-disabled container, and a strong signal the paper’s own code was never actually run non-interactively even by its own authors; (2) a script calling

`rstanarm::launch_shinystan()`, which starts a Shiny web server and blocks waiting for a browser connection that can never arrive in this sandbox — resolved only by the per-script timeout eventually killing it, after up to 3,000 seconds spent doing nothing.

Missing inputs. 129 of the 343 code-bearing papers (37.6%) have at least one script whose required input was never produced or found (`repro_missing_inputs > 0`); 417 individual scripts are recorded `not_runnable_reason = "missing_input"`, the largest single `not_runnable_reason` category by a wide margin (68 `parse_error`, 40 `unplaceable` — a script the run-order algorithm could not confidently place in a dependency sequence at all — and 3 `jags_model`, a script that is itself a JAGS model-definition file rather than executable R).

Dependency installation. Before running a paper’s code, Metacheck must first install the R packages that code declares it needs. 274 of the 343 code-bearing papers (79.9%) declare at least one such package (`repro_deps`, a per-paper count of declared packages — not a count of failures, and not tied to any particular script or file), 3,508 packages in total. The per-package install success/failure counts, the specific packages that failed and how many papers each affected, and the result-matching statistics later in this subsection come from a dedicated extraction against each paper’s saved checkpoint report text (`01_run_metacheck/08_extract_reproducibility_detail.R`), which reads each paper’s own free-text install-summary sentence and, for result matching, its row-level `match_table`. This extraction was rebuilt for the current, merged 343-paper result; the original script that first built it for the pre-Appendix-A.8 pass was not itself preserved anywhere in this project (see the note further below).

Metacheck’s own `summary_table` records only how many packages a paper declares (`repro_deps`), not which ones failed to install or why — a real limitation this appendix works around below, and A.7.6’s second suggestion recommends Metacheck fix directly. This appendix could still recover the failed packages’ actual names, because each paper’s saved checkpoint file also keeps a plain-text summary sentence restating its own install attempt (e.g. “13 succeeded, 9 failed. `sf`, `rgdal`, ... were installed. **Failed:** `countreg`, `DTK`, ...”); re-reading that sentence for each paper is how the list below was built. 256 of the 265 papers had this sentence in a readable form; the other 9 are a gap in what this appendix could recover, not evidence of anything about those 9 papers’ own installs.

One further point, worth stating plainly rather than folding into the numbers above: none of the 361 failed installs were classified by Metacheck as `dependency_unavailable` — its own label for “this package cannot be found anywhere, so this is not really about the paper’s own code.” Every failure here was instead classified as an ordinary code error, which this appendix takes as license to keep treating a failed install as a real, meaningful cause of non-reproducibility, not something to set aside as merely today’s state of the package ecosystem.

The rest of this subsection focuses on the 80 papers identified above: a real install failure, and no script that ran successfully. This is the subset where a failed install is most clearly the reason the paper’s code did not run.

Note on this list. As with the numbers in A.7.3/A.7.4 above, this list was rebuilt (using this project’s own `01_run_metacheck/08_extract_reproducibility_detail.R`, re-parsing each paper’s saved checkpoint report text) after Appendix A.8’s rerun, since the original extraction script that first built it was not itself preserved anywhere in this project (only its output — `01_run_metacheck/batches/reproducibility_check/pre_a8_reference/` — survived). The rebuilt extraction was checked directly against that surviving output for every package outside Appendix A.8’s 98-paper rerun before being trusted: of the 90 packages the original extraction found, 86 have an identical paper count in the rebuild, 2 (`rethinking`, `rJava`) changed by one paper each as a direct, explicable result of the rerun, and 2 (`ascr`, `pgirmess`) no longer appear because the one paper each had failed for moved out of this 76-to-80-paper subset (its code now has at least one successful script, or its failed install no longer applies) — a small, fully accounted-for set of differences, not evidence of a remaining extraction problem.

Which packages failed to install, for the 80 papers with a failed dependency and zero successful scripts. Every one of these 80 papers has a `report` field in the saved checkpoint data with a recoverable, itemised list of which specific packages failed. Across those 80 papers, 187 individual failed-install records were recovered, naming 109 distinct packages. Five of those, across several papers, are not genuine package names at all but a further Metacheck defect: `x`, `package`, `not_installed`, `arma`, and `packages.toinstall` were each confirmed absent from CRAN entirely and, on direct inspection of the paper’s own repository, are variable names from the paper’s own install-loop code (e.g. `for (pkg in list.packages) install.packages(pkg)`) that Metacheck’s dependency extractor mistook for a literal package name to install — see A.7.5’s newly-added fifth defect below.

By far the most common genuine cause is a small cluster of interdependent **spatial/GIS packages**: `sf`, `rgdal`, `rgeos`, `maptools`, `stars`, `ggspatial`, `mapview`, `tmap`, `gdalUtils`, `ngeo`, and `spatial.tools` are not ordinary self-contained R packages but R interfaces to external, non-R system libraries — GDAL (reads/writes geographic file formats), GEOS (geometric operations such as intersections and buffers), and PROJ (coordinate-system transformations) — which must already be installed at the operating-system level before the R package can even compile, let alone load; installing the R package alone is not sufficient. `sf` (22 papers), `rgdal` (17), `maptools` (7), `rgeos` (6), `ggspatial` (4), and `stars` (4), together with the smaller one/two-paper members of the same cluster, account for 68 of the 187 records (36.4%) — consistent with a single shared environment gap (`ghcr.io/scienceverse/metacheck_r:latest` apparently does not have GDAL/GEOS/PROJ installed) rather than 68 independent per-paper causes; see Suggestion 2 below, which recommends closing this gap directly in Metacheck’s own Docker image rather than leaving it to recur for every future spatial-analysis paper. The next-largest identifiable cluster is **Bayesian/JAGS modelling**: `rjags` (4 papers), `R2jags` (1), and `INLA` (1) — JAGS is itself a separate, non-R program that must be installed on the system for `rjags` to link against, the same category of gap as the spatial cluster. A Java-dependent cluster (`rJava`, 4 papers, matching the `.jinit_undefined_variable` pattern reported above) similarly requires a system-level Java runtime, not just the R package.

Every other package in the list below failed for exactly one paper each. This appendix cannot

say why each of these one-off packages failed to install, and one specific possibility this run’s own method already rules out for the great majority of them is worth naming directly: Metacheck’s `sandbox = "process"` install backend (`repro_install_deps()`) retries a failed live-CRAN install once against the CRAN Archive (a package’s last-published version before CRAN removed it) before giving up — but this run used `sandbox = "docker"`, whose install backend (`repro_install_deps_docker()`) does **not** implement this same Archive retry (confirmed directly against its own source; its `via_archive` field is hard-coded `FALSE` for every install, with an explicit comment noting the retry “is not duplicated here”). So a package that failed here only because its current CRAN release could not compile, but whose Archive-held prior version would have installed cleanly, was never given the chance this run — the Archive retry that might have recovered it simply was not attempted, independent of whatever the package’s own true availability status is. A handful of names (e.g. `hmRLib`, `poispkgs`, `tiagospackage`) are also plausibly a typo or a private/unpublished package never intended for public installation at all, though this was not individually verified for each. This is a separate, more specific point than this appendix’s general scope note below: it is not merely that the reason for each one-off failure is unrecorded, but that one entire category of possible reason (an Archive-recoverable package) was structurally never tested for in this run at all.

The full list, package name and number of these 80 papers it failed for ($n = 1$ except where noted): `sf` (22), `rgdal` (17), `maptools` (7), `rgeos` (6), `ggsn` (4), `ggspatial` (4), `rjags` (4), `rJava` (4), `stars` (4), `phyloseq` (3), `amt` (2), `Biostrings` (2), `dada2` (2), `decontam` (2), `gdalUtils` (2), `glmulti` (2), `mapview` (2), `mcmcplots` (2), `rethinking` (2), `ShortRead` (2), `tmap` (2), `adiv`, `asreml`, `betalink`, `brtest`, `CAST`, `classecol`, `CommEcol`, `countreg`, `dartR`, `diveRcity`, `DTK`, `embryogrowth`, `emmeans`, `evobiR`, `FancyBaSTA2`, `fishflux`, `FRQNT1.Data`, `GeoLight`, `GeoLocTools`, `ggalt`, `gghalves`, `ggmosaic`, `ggordiplots`, `ggpattern`, `GISTools`, `git2r`, `glmmADMB`, `gurobi`, `GVI`, `hdi`, `hier.part`, `hilldiv`, `hmRLib`, `InformationValue`, `INLA`, `IPMpack`, `jpeg`, `lme4`, `lmerTest`, `MCMCglmm`, `microbiome`, `microclima`, `NicheMapR`, `nngео`, `odeintr`, `Orcs`, `pairwiseAdonis`, `PAMLR`, `parameters`, `PCAviz`, `performance`, `performanceEstimation`, `pez`, `phylolm`, `picante`, `plotrix`, `prioritizr`, `purr`, `R2jags`, `RAMClustR`, `Rarefy`, `razimuth`, `rgl`, `rich`, `RNetLogo`, `seqRFLP`, `SGAT`, `siar`, `SNPRelate`, `spatial.tools`, `spdep`, `tabulizer`, `taxotools`, `tess3r`, `tidync`, `tidyverse`, `TwGeos`, `units`, `V.PhyloMaker`, `vegetarian`, `viridis`, `waffle`, and `weathercan`.

A note on scope for this package list. This list names which package failed to install, for how many papers, within the 80-paper subset defined above (a genuine install failure, and zero scripts of that paper ran successfully); it is not a claim about why each package failed (a version incompatibility, a missing system library, a compilation failure, genuine current unavailability, or something else) — Metacheck’s own `.repro_classify_install_message()` classifier already exists and runs during the install phase, but its per-package `category` output is not retained in this run’s paper-level `summary_table`, the same gap A.7.6’s second suggestion recommends closing. Nor does this list cover the 82 further papers (162 total with a genuine install failure minus the 80 covered here) that had at least one failed dependency but still had at least one script run successfully — extracting and reporting that broader set’s package list, and each package’s classified failure reason once Metacheck retains it, is a natural extension of this analysis and is not attempted here.

Result matching: did the code’s own output match what the paper reports?

This is a distinct, further step `reproducibility_check` performs beyond simply running a paper’s code: for every paper whose code did produce console output, it extracts each statistical result the paper’s own text reports (a `t`-test, a correlation, a regression coefficient) and checks whether that exact value appears anywhere in what the executed code actually printed (`match_reported_output()`). This is the closest this run comes to asking not just “does the code run” but “does running it get you the paper’s own numbers.”

Metacheck’s statistical-output extraction and matching machinery is not, in general, R-only: `reproducibility_check` also reads and STATO-tags the rendered result tables inside a `.jasp/.omv/.spv` file (JASP, jamovi, or SPSS Viewer — self-reproducible bundles that carry both the data and its analysis together, requiring no code execution at all) via `read_stat_tables()`, feeding the same matching step as executed R output. Checking directly against every one of the 133 papers’ own saved statistical output (802 extracted result entries in total), however, finds every single one came from executed R code — none from a JASP/jamovi/SPSS-Viewer file. So the papers below happen to be entirely R-based, but this is a fact about which papers in this specific corpus had a matchable non-R statistical-output file (apparently none), not a restriction the matching step itself imposes; `reproducibility_check`’s own code-execution phase is separately, and by explicit design, restricted to `language == "R"` files (`inst/modules/reproducibility_check.R`), so a SAS, SPSS-syntax, Stata, Python, or MATLAB script is never itself run, whether or not a self-reproducible output file exists alongside it.

133 papers had at least one statistical result Metacheck could both extract from the paper’s own text and attempt to match against the executed code’s console output (`repro_tests_reported > 0`); across those 133 papers, 802 individual reported test statistics were checked, of which only 66 (8.2%) were matched by the code’s actual output. At the paper level, the median proportion of a paper’s own reported results its code reproduced was **0%**: only 1 of the 133 papers (0.8%) had its code reproduce every reported result it checked, 123 (92.5%) reproduced none of them, and the remaining 9 (6.8%) reproduced some but not all. This is the single starkest number this run produced, and — read together with the execution-outcome breakdown above — is consistent with a straightforward, largely mechanical explanation rather than a claim about the underlying research: a script has to actually run to completion and print a result before that result can be compared to the paper’s text at all, and only 45.4% of scripts, and just 24.8% of papers in full, cleared that first bar in this run.

Was this matching step itself useful, beyond confirming the headline 8% figure?

Yes, in two ways. First, of the 66 matched values, 62 (93.9%) matched at “full” confidence (every component of that reported test — the statistic, its `p`-value, its confidence interval — matched together) and only 4 (6.1%) matched “partially” (some but not all components of the same test found) — when this run’s code did reproduce a result, it almost always reproduced it completely, not approximately; a partial match was the exception, not the norm. Second, the matcher records, separately from whether a value was found at all, whether an unmatched reported statistic could still be traced to a real, identifiable quantity in the code’s output

(`plausible_split`, flagged only for a test whose reported components had to be re-grouped from the paper’s own text rather than read as originally written). Of the 736 reported values not matched, 523 (71.1%) were nonetheless flagged `plausible_split = TRUE` — meaning the value the paper reports does correspond to a real quantity the code computed somewhere, under a different grouping or a different analysis than the paper’s own text suggested, rather than being absent from the output altogether. This is a genuinely useful distinction this run’s matching step draws for free: “the code never produced anything resembling this number” is a materially different, more serious finding than “the code produced this number, just not packaged the way the paper’s text implies it should be,” and only the row-level `match_table` this appendix worked from here — not the paper-level `repro_tests_matched` count alone — can tell the two apart.

A note on scope and what this section does not claim. The numbers above describe whether a paper’s own downloaded code, as shared, ran and reproduced its own reported results inside this specific execution environment (Docker, `ghcr.io/scienceverse/metacheck_r:latest`’s pre-installed package set, a 3,000-second per-script budget, no network access during execution). A script failing here is not necessarily evidence the paper’s underlying research is wrong, or even that the code never worked — it may simply never have been intended to run outside the author’s own machine, a point the `file.choose()` and multi-script “designed as an interactive pipeline” patterns above illustrate directly. Nor does this section attempt to re-diagnose why each of the 1,829 errored or 447 skipped scripts individually failed beyond the categories `reproducibility_check` itself records; a full per-paper qualitative review of this kind, in the same style as Appendix A.1–A.6’s disagreement-by-disagreement review of `repo_check/data_check/code_check`, is listed as a further step in A.7.6 below rather than attempted here.

A.7.5 Infrastructure failures, found and fixed during this run

Running `reproducibility_check` at corpus scale, for the first time, using Metacheck’s new `workers`-based concurrency feature, surfaced five defects in Metacheck itself — none related to a paper’s own code or data, all specific to the mechanics of running many papers’ Docker containers concurrently, checkpointing their results, or (the fifth) extracting a paper’s own declared dependencies from its code. All five were filed as issues against `scienceverse/metacheck`; two are confirmed fixed on the `dev` branch as of this writing (verified directly against the installed package’s own source, not merely by the issue’s own open/closed status), the other three are filed and not yet independently reverifed.

`scienceverse/metacheck#417` — a Docker container left running after its own script’s timeout fires. `repro_run_scripts_docker()`’s per-script timeout (`processx::run(..., timeout = timeout)`) kills the host-side docker run CLI process when a script exceeds its budget, but this does not guarantee the container itself stops — confirmed live twice during this run: once, a `caret/gbm` hyperparameter grid search (490 combinations) continued running inside its container at 100% CPU for over 80 minutes after

its own timeout had already fired and been correctly recorded as `outcome = "timed_out"`; a second time, a script calling `rstanarm::launch_shinystan()` (the interactive-session case named in A.7.4 above) was left running, near-idle, on a listening port after its own timeout. Both were stopped manually (`docker stop`) after being noticed via `docker ps -a/docker stats` from outside the run entirely — nothing in the run’s own log output indicated either container was still running. **Not yet confirmed fixed.**

scienceverse/metacheck#418 — **a multi-line `deparse()` silently corrupts the generated install script for a paper with many dependencies.** `repro_install_deps_docker()` builds the container’s install script with `paste0(".repro_docker_pkgs <- ", deparse(install_deps$package, ...))` — but `deparse()` returns a **vector** of lines, not a single string, once its argument is long enough to wrap (R’s own default line-width behaviour); `paste0()` silently keeps only the first line, splitting the vector literal apart mid-declaration. Confirmed live: a paper declaring 18 CRAN dependencies produced a generated `install.R` reading `.repro_docker_pkgs <- c("bayestestR", ..., "lubridate",` on one line, followed by a syntactically invalid continuation, which failed with **Error: unexpected symbol** and reported all 18 of that paper’s dependencies as failed installs — not because any package itself was unavailable, but because the install script that would have installed them was never valid R in the first place. A full scan of every batch already completed at the time this bug was found identified 34 of 1,861 papers (1.8%) hit by this exact corruption signature; those 34 papers’ checkpoint files were deleted and recomputed once **metacheck#420** (below) was fixed and this run was restarted. **Not yet confirmed fixed** (the affected 34 papers were recomputed under the pre-fix code path avoided by deletion-and-rerun, not by an actual code fix landing on `dev`; a real fix for this specific defect remains open as of this writing).

scienceverse/metacheck#419 — **`collect_module_tables()` always returns an empty result.** `capture_module_tables()` (the function `results_dir` uses to write each paper’s per-paper `.rds` checkpoint) saves each module’s output into an **unnamed** list; `collect_module_tables()` (the function meant to read those checkpoints back and reassemble them) looks each module up **by name** — so every lookup silently misses, and `collect_module_tables()` returns an empty data frame no matter how many papers were genuinely saved, with no warning that anything went wrong. This defeated `results_dir`’s entire stated purpose (a resumable, inspectable per-paper checkpoint) the first time this run’s own wrapper script tried to use it to reconstruct a batch interrupted partway through. **Confirmed fixed** on `dev` (commit 033294d, “Fix `capture_module_tables()` saving modules unnamed”): `capture_module_tables()` now names its saved list by each module’s own `$module` field, with a positional fallback for a missing or duplicate name. This project’s own run script does not depend on the fix either way going forward, since it was changed to read each paper’s saved `.rds` file directly (`modules[[1]]`, positionally) rather than through `collect_module_tables()`, once the defect was found live and a fix could not be waited for mid-run.

scienceverse/metacheck#420 — **`.reproducibility_check_batch()` can hang indefinitely, silently dropping the last papers in the queue.** Confirmed live: a run processing

1,861 papers at `workers = 5` reached 1,859 of 1,861 with no error, no further log output, all five worker processes idle (alive but with frozen CPU time across repeated checks 10+ minutes apart), and zero Docker containers running — yet the function never returned. The two remaining papers had no `-> starting` log line at all; `launch()` had simply never been called for them. The maintainer’s own diagnosis and fix (reported back to this project directly, not independently re-derived here) was a `callr` pipe deadlock: `callr::r_bg()`’s default stdout/stderr capture uses a pipe, and once a worker produces enough output to fill that pipe’s buffer with nothing reading it, the worker blocks trying to write, indefinitely, with no error surfaced anywhere. **Confirmed fixed** on dev (commit 85c8c87, “Fix `.reproducibility_check_batch()` hang from `callr` pipe deadlock”): each worker’s stdout/stderr is now redirected to a file rather than a pipe, and read back from that file rather than via `p$read_all_error()`’s pipe-based accessor. The stalled run was killed and restarted after this fix landed; `results_dir`’s per-paper checkpoints meant the restart correctly skipped the 1,859 already-completed papers (reported live: “resuming -- 1859/1861 paper(s) already saved in `results_dir` from an earlier interrupted run”) and completed only the two genuinely outstanding papers.

scienceverse/metacheck#421 — dependency extraction can mistake a paper’s own install-loop variable name for a package name. Found while extracting and summarising which packages failed to install across this run (A.7.4 above): a paper’s own dependency-installing boilerplate commonly assigns its package list to a variable and iterates it in a loop (e.g. `list.packages <- c(...)`, then `for (pkg in list.packages) install.packages(pkg)`); in six confirmed cases, the extractor captured a bare loop-variable identifier itself — `x`, `package`, `dependencies`, `not_installed`, `arma`, `packages.toinstall` (plus, in a seventh and eighth case, `list.packages` and `req.lib` themselves, the variable names from the concrete example quoted in A.7.4) — as if it were a literal package name to install, rather than recognising it as a variable reference whose value static analysis cannot generally resolve. Each of these eight tokens was confirmed absent from CRAN entirely (“`x`” `%in% rownames(available.packages())`), and equivalently for the other seven, all `FALSE`) and, on direct inspection of the originating paper’s own repository code, traced to exactly this pattern. **Not yet confirmed fixed.**

Sixteen further papers recovered by a targeted rerun. Independently of #420’s queue-loss bug, 16 of the 1,861 papers (0.9%) received a synthetic `traffic_light = "error"` stand-in result during the main run rather than a genuine check, due to either a `callr` subprocess itself crashing (“could not start R, exited with non-zero status, has crashed or was killed”, 14 papers) or a transient Docker Desktop outage on the host machine (“Docker does not appear to be running”, 2 papers) — infrastructure noise, not a property of any of these 16 papers’ own code. All 16 were identified by grepping the run’s own log for this failure signature, listed explicitly in `rerun_infra_failures.R`, and rerun in a small, separate, `workers = 4` batch once the main run finished; 15 of the 16 completed with a genuine result on this second attempt and were merged into the corpus-wide combined output in place of their placeholder rows. The remaining paper (10_1111_2041_210x_13384) crashed with the identical `callr` error on both the original run and the 16-paper rerun; a third, solo attempt (`workers = 1`, no concurrency, `rerun_single_paper.R`) found this paper’s repository

contains 314 separate scripts and a 1,386-row file-materialisation plan — substantially larger than any other paper’s repository encountered during this run — and was manually cancelled, still mid-dependency-install, after several minutes with no sign of proceeding to actual script execution. This paper’s result in the corpus-wide combined output is explicitly recorded as `outcome = "timed_out"` with a `not_runnable_reason` naming this exact sequence of events, rather than left as an ambiguous crash placeholder or given a fabricated genuine result.

A.7.6 Suggestions for improving `reproducibility_check`

The following suggestions are drawn directly from running the module at full corpus scale for the first time; each is grounded in a specific, confirmed observation from this run, not a general aspiration.

1. **Add a distinct outcome for “this script requires interactive input.”**
Both `file.choose()`-driven `read.csv()` calls and `launch_shinystan()`’s blocking Shiny server (A.7.4) currently surface as an ordinary `runtime` error or a `timeout`, indistinguishable from a genuine bug in the paper’s code. A small, targeted check — recognising a short list of known-interactive calls (`file.choose()`, `menu()`, `readline()/readLines(file("stdin"))`, `shiny::runApp()/launch_shinystan()/plotly::orca()`-style blocking viewers) before execution, or classifying a script that produces zero stdout before timing out as `timed_out_interactive` rather than plain `timed_out` — would let a report distinguish “this code cannot run unattended by design” from “this code is simply slow,” two very different findings for an author or reader trying to act on the result.
2. **Preserve the install-failure package name and reason, not just a count, through to the paper-level summary table.** `repro_install_deps_docker()`’s own return value already carries both which package failed and a classified reason (a version incompatibility, a missing system library, a compilation failure, genuine unavailability — see `.repro_classify_install_message()`), but this run’s corpus-wide saved `summary_table$repro_deps` retains only a per-paper count. This project could still recover the `names` by re-parsing each paper’s free-text `report` field (A.7.4 above), but not the classified `reason` — a structured `summary_table` column (e.g. `repro_deps_packages`, `repro_deps_by_category`) would make both available directly, for every paper, without a second extraction pass against unstructured text. This would also let a report distinguish “this paper’s code cannot run today because CRAN removed a dependency it used” (an ecosystem-drift problem, not a fault of the paper) from “this paper’s code cannot run because it needs a system library this environment does not provide, e.g. the GDAL/GEOS/PROJ stack behind `sf/rgdal/rgeos/maptools`, together the largest single cluster of failed installs this run found” (an environment gap Metacheck could plausibly close by extending its own Docker image).
3. **Auto-stop a script’s container on timeout (`metacheck#417`), not just the host CLI process.** As documented in A.7.5, a timed-out script’s container can keep

running, and consuming real host resources, for a long time after the timeout is reported — capturing the container’s own ID (e.g. via a fixed `docker run --name`) and explicitly `docker stop/docker` killing it when `processx::run()` reports a timeout would close this gap directly, and would have prevented two of the interventions this run’s own operator had to make manually.

4. **Guard every `paste0(prefix, deparse(x))` call site against multi-line `deparse()` output** (metacheck#418), not only the one instance this run happened to trigger — `deparse1()` (base R, available since R 4.0) exists specifically to always return a single string regardless of length, and is a direct drop-in replacement everywhere this pattern appears.
5. **Retry a `callr` subprocess crash once before recording it as the paper’s result**, at least for `.reproducibility_check_batch()`’s `workers > 1` path. 9 of this run’s 16 infrastructure failures (56.3%) were a bare `callr` subprocess crash with no more specific diagnosis available (“could not start R, exited with non-zero status, has crashed or was killed”) — a category that, by its own error text, cannot currently be distinguished from a genuine, reproducible crash in the paper’s own code (as the twice-crashed paper in A.7.5 turned out to be) without a human manually re-running it, as this project had to do for all 16.
6. **Add a repository size/complexity gate to `reproducibility_check` specifically**, analogous to `download_repo_files()`’s existing `max_files_per_repo/max_file_size/max_download_size` caps (Appendix A.4). The one paper this run could not complete even alone (A.7.5) has a 314-script repository — nothing in `reproducibility_check` currently flags or specially handles a repository at this scale before attempting a full materialise-install-execute cycle on it, the same category of gap A.4 closed for `data_check/code_check`’s own download step.
7. **Make a stalled `workers > 1` run detectable from its own log**, independent of the underlying cause. Diagnosing metacheck#420 required inspecting Docker container state and OS-level process CPU time from outside R entirely, since the run’s own console output gave no indication anything was wrong — it simply stopped producing lines. A lightweight, periodic “N papers still running, M queued, last completion Ts ago” heartbeat message, or a caller-configurable overall wall-clock ceiling independent of any individual script’s own timeout, would have made this failure mode visible from the log alone.
8. **Recognise a bare variable identifier passed to `install.packages()` as unresolvable, rather than as a literal package name** (metacheck#421). A static extractor cannot generally know a variable’s runtime value, so a call like `install.packages(pkg)` inside `for (pkg in list.packages) ...` should either be skipped entirely, or resolved only when the iterated variable (`list.packages` in this example) is itself assigned a literal character vector earlier in the same file — the case in every one of this run’s eight confirmed instances of this defect (A.7.5).
9. **Give `repro_install_deps_docker()` the same CRAN Archive retry `repro_install_deps()` (the `sandbox = "process"` backend) already has**. Confirmed directly against the Docker backend’s own source: `repro_install_deps_docker()`’s `via_archive` field is hard-coded `FALSE` for every install attempt, with an explicit

comment noting the process backend’s Archive-retry-on-failure logic “is not duplicated here.” A package whose current CRAN release fails to install (a build error, an R-version incompatibility) but whose last-published Archive version would install cleanly is therefore never given that chance under `sandbox = "docker"` specifically — the safer, actually-isolated backend this project used for exactly that reason (A.7.1) — even though the less-isolated `sandbox = "process"` backend would have tried. Checking each of the 70 one-off failed packages directly against live CRAN and the CRAN Archive (<https://cran.r-project.org/src/contrib/Archive/<pkg>/>, without running any paper’s code) quantifies the upper bound this fix could plausibly reach: 21 (30.0%) are on live CRAN today, so their failure was for some other reason entirely, not archiving; 24 (34.3%) have no trace on CRAN at all, live or archived (a GitHub-only package, a Bioconductor package this check did not look for in the right place, or a name this appendix could not otherwise verify); and **25 (35.7%) are absent from live CRAN but do have a version sitting in the Archive** — exactly the set an Archive retry could plausibly recover, though a successful Archive fetch is not a guarantee of a successful install (the archived version could itself fail to compile, or need a system library this image still lacks), so 25 is an upper bound on what this specific fix could help with, not a prediction of how many would actually then install.

10. **Extend the full-corpus, disagreement-by-disagreement qualitative review method of Appendices A.1–A.6 to reproducibility_check’s own output.** This appendix reports the rate and category of non-reproducibility across the corpus, but — unlike Appendices A.1–A.6’s per-disagreement review of `repo_check/data_check/code_check` — does not individually verify a sample of the 1,829 errored or 447 skipped-input scripts against their own paper’s repository the way, for instance, A.3’s five-check validation individually read 100+ real files. A systematic sample of `runtime`-classified failures in particular (65.4% of all errors, and the category this appendix could say the least about, since it is Metacheck’s own catch-all for “every other kind of R error”) would likely surface further mechanically fixable causes, in the same style A.1–A.6 already found for `repo_check/data_check/code_check`.

A.8 Three further reliability fixes to compressed-archive handling, and their effect on `data_archive/code_archived/code_download`

An eighth round of work returned to a structural limitation named repeatedly in earlier appendices but never itself directly fixed: `zip_peek()`’s inability to read a zip’s contents without downloading it on Dryad and Figshare, both of which redirect a file’s download URL to Amazon S3. Diagnosing and fixing this surfaced two further, closely related defects in the same download path, and reprocessing every paper the three fixes together could plausibly affect closed most, but not all, of the remaining `code_download` disagreement this comparison’s own numbers above already reflect.

`zip_peek()` failing on S3-backed hosts (scienceverse/metacheck#424). `zip_peek()`

reads a remote zip's table of contents with two requests: a `HEAD` to learn the file's total size, then a byte-range `GET` for the tail bytes holding the archive's central directory. Amazon S3 — the storage backend both Dryad's and Figshare's own download URLs redirect to — answers the size-only `HEAD` request with 403 but answers the byte-range `GET` normally (confirmed live: a `HEAD` against a real Dryad-hosted file returned 403, while a `Range: bytes=0-100 GET` against the identical URL returned 206 with the file's true total size in its own `content-range` header). Because `.http_range_tail()`, the function `zip_peek()` calls to fetch the tail bytes, only sends its range request once `HEAD`'s size is known, a `HEAD` failure meant the byte-range request was never sent at all — not just for Dryad and Figshare but, confirmed the same way, for Harvard's Dataverse installation, another host redirecting to S3. Fix: fall back to a suffix range request (`Range: bytes=-<n>`) when `HEAD` fails, reading the total size from that response's own `content-range` header instead. Confirmed live before and after against nine archives sampled from this corpus (three each on Dryad, Figshare, and Zenodo): all six Dryad/Figshare archives returned `NULL` from `zip_peek()` before the fix and a correct, complete listing after it, while the three Zenodo archives (not S3-backed, and so unaffected by the underlying bug) listed correctly both times.

`zip_peek()` never persisting its result across a restart (scienceverse/metacheck#427).

A separate, longstanding defect (`repo_info_cache()`) already exists specifically to persist a repository's listing (which files it has) to disk, since restarting an interrupted corpus run otherwise re-pays every repository's listing cost from scratch — a real problem previously severe enough to exhaust Dryad's own daily request quota (see that fix's own code comment, confirmed live 2026-09-08: 11 restarts of one 50-paper batch cost 440+ re-listing requests against a workload that genuinely needed a handful). `zip_peek()`'s own result — a zip's contents, a separate cost from the repository listing above it — had no equivalent: it was cached only in memory, for the life of one R process, so the same restart pattern silently re-peeked every zip already peeked, and `repo_check` and `data_check` each peeked the same zip independently (as two separate R processes in this project's own pipeline), doubling the cost again. Confirmed directly during this project's own corpus rerun (below): before this fix, four consecutive restarts each re-hit the same Dryad rate limit at the same early batch, at increasingly long waits (24.5, then 21.5, then 16 minutes remaining); after it, a restart at the same point in the run loaded every already-completed batch from disk in under 20 seconds and progressed past the previously-stuck batch with no rate-limit wait at all. Fix: persist `zip_peek()`'s result to disk the same way `repo_info_cache()` already does for a listing, gated on the same `cache = TRUE` argument.

A silently discarded error in `download_repo_files()`'s archive-member fetch (scienceverse/metacheck#429). Once `repo_check`'s zip-peek expansion lists a file living inside an archive (an “archive member,” carrying `archive_url+archive_member` rather than its own `file_url`), `download_repo_files()` fetches it via a dedicated byte-range path (`.zip_fetch_members()/zip_member_fetch()`), wrapped in a bare `tryCatch(..., error = function(e) NULL)` with no logging of any kind. A transient failure at this one step — a dropped connection, a timeout, a momentary host error — left `file_location` permanently `NA` with no warning, no log entry, and no record anywhere that a download had even been attempted,

indistinguishable after the fact from a member the host genuinely cannot serve. This surfaced directly in this project’s own corpus rerun: after #424 above let `repo_check/data_check` correctly list many more Dryad/Zenodo archive members than before, 24 papers whose code repository had previously never been detected at all (`mc_code_archived FALSE`) newly showed `mc_code_archived TRUE` but `mc_code_download` still `FALSE` — the code files were now correctly found, but never downloaded. Tracing one case in full (a small `.R` file inside a Dryad-hosted zip) confirmed the underlying fetch mechanism worked perfectly when called directly, with the paper’s own exact saved row data and arguments, ruling out a logic bug; the corpus run’s own 17,661-line warning log had no entry at all for this archive, this paper, or any of the other 23, confirming the failure really was silently discarded rather than merely unlogged elsewhere. Fix: record the specific error `.zip_fetch_members()` raises, the same way the ordinary `file_url` download path a few lines below already does, rather than discarding it.

Reprocessing the 24 code-download regressions. All 24 papers’ `code_check` output was rerun against `dev` after #429’s fix (`code_check` alone — `repo_check`’s listing and `data_check`’s own downloads were unaffected, confirmed directly: `mc_data_download` was already correctly `TRUE` for the papers where a data download was expected). 13 of 24 resolved immediately. The remaining 11 were rerun a second time with `skip_on_api_limit` set `FALSE` for this small, targeted rerun (unlike the corpus-wide default `TRUE`, whose own rationale — never blocking an unattended, hours-long, 1,861-paper run on a host’s quota reset — does not apply to an 11-paper rerun, and, confirmed directly, was itself the reason these 11 still failed: their Dryad archive fetch had landed during a rate-limit window and been silently skipped rather than waited out, and waiting the quoted remaining time fetched the exact right file on the first attempt); 4 more resolved. The final 7 each printed a `download_repo_files()` retry-exhaustion message (Failed to perform HTTP request, Re-run to retry: cached files are reused, only the missing files are fetched) that — a `message()/cat()` call, not a `warning()` — this project’s own rerun script’s `warning = handler` could not catch or log, a message-routing gap distinct from #429’s own fix and not itself corrected here; rerunning those 7 a third time, as the message itself instructed, resolved all 7. Every one of the original 24 papers now downloads successfully.

Corpus-wide effect. With these three fixes installed and the full 1,861-paper corpus rerun, sensitivity, specificity, and the number of papers newly entering each comparison (a repository or file that previously could not be evaluated at all now can be) all moved together in the direction these fixes targeted. `code_archived` agreement rose from 84.5% to 87.4% (sensitivity 60.1% to 69.3%), `any_readme` from 69.8% to 72.1% (sensitivity 59.6% to 62.9%), and `data_availability` from 83.3% to 84.7%. `code_download`’s own comparable count grew from 342 to 394 papers (50 papers whose code could not previously be evaluated for a download outcome at all — because it had never been detected as archived in the first place — now can be), and agreement across that larger, harder set held at 93.4% (26 disagreement remaining, down from 24 at the point these three fixes began). `zip_not_peeked` — the disagreement cause tracing directly to #424’s own limitation — dropped to 0 as a recorded cause across every construct in this comparison’s disagreement-cause categorization. The “compressed archive never opened” category among `data_archive`’s remaining under-coded papers (Appendix’s own

`data_archive_uncoded_papers.RData` analysis) fell from 70 papers to 12; the remainder now trace to distinct, already-documented causes (a code-only or documentation-only repository, a non-tabular file format `data_check` does not classify as data, or an image/GIS/3D file). Every previously-recorded disagreement in the corpus was re-reviewed the same way as every earlier round: 1,070 of 1,174 prior verdicts carried forward unmodified, because their underlying (paper, column, Cooper value, Metacheck value) triple was unchanged; every disagreement whose Metacheck-derived value changed, or which was newly surfaced, was reviewed from scratch (43 rows immediately after the corpus rerun above; 19 once the 24 papers' `mc_code_download` values were updated by the reprocessing described above and the comparison data rebuilt a second time) against the same evidence-source rule used throughout this manuscript (repository file listing and `repo_metadata` for `data_license/code_archived/any_readme/code_download`, the paper's own text for `data_availability/data_archive`). Every disagreement in the corpus — 819 of 819 — again carries a verdict.

A.9 A full, paper-by-paper review of every “unsupported repository, not detected” disagreement, and a full recheck of Appendix A.6’s original categorization

Appendix A.6 categorized every “repository not detected” disagreement in the corpus, but 40 of the 172 papers whose only Cooper-coded `data_archive` label(s) fell outside Metacheck’s supported-platform categories, and for which `repo_check` consequently found no repository at all, had not yet been individually checked one at a time when an earlier version of this manuscript’s main text was written. This appendix documents that remaining review in full: each paper’s own extracted data-availability text, the specific DOI or URL cited, whether it was live-verified to resolve, whether the cited platform is genuinely supported by Metacheck (checked directly against the vendored source in `code/data/reproducibility_check/vendor/metacheck/R/`, not inferred from this manuscript’s own earlier prose, which had itself gone stale for at least one platform – see below), and the final category assigned. All 40 papers received a verdict at the time of this review; 3 have since had their underlying `repo_check` defect fixed by later work elsewhere in this project and no longer appear in the live disagreement set at all (confirmed: `repo_check` now finds a repository for every one of them), so Table 2 below shows the 37 that remain.

This review also surfaced a further risk worth checking directly rather than assuming away: Appendix A.6’s original 230-row categorization was itself produced before the platform-support defects this review found (the stale 4TU/DANS assumptions, an incomplete dedicated-platform list, a missing DSpace-legacy host count), so any of those original 230 rows that happened to cite one of these same platforms could be mislabeled the same way. Rather than only check the platforms these specific bugs touch, every one of the original 230 rows was rechecked from scratch: each paper’s own text was re-extracted directly from the corpus (paragraph-grouped, not sentence-truncated, avoiding the extraction bug documented earlier in this manuscript), every DOI and URL in that text was identified, and each one’s host was checked directly

against Metacheck’s current supported-platform functions – not against what an earlier comment claimed was or was not supported.

Of the 230, 21 candidate citations resolved to a host now confirmed supported despite the row’s original `UNSUPPORTED_REPO` label; after excluding false positives (a truncated GitHub URL with no repository path, a citation that names a genuinely unrelated, unsupported host despite superficially matching a regex), 9 rows were corrected to `SUPPORTED_PLATFORM_BUG` or `SUPPORTED_PLATFORM_MISSED` at the time of this review. The corrected rows span three recurring patterns: a citation to a genuinely supported host that `repo_check` simply never queried (Dataverse, DataOne/KNB, Zenodo, standard-format Figshare), a supported platform cited via a URL format its matcher does not parse (a Figshare private share-link or collection URL), and a supported platform cited via a non-standard DOI prefix its prefix-matching does not recognise (an alternate Dryad EZID prefix, a non-standard Figshare prefix). None of the corrected rows involve a platform Metacheck genuinely cannot reach at all – every correction is a detection defect on an already-supported platform, consistent with the pattern the original 40-paper review above found. 8 of these 9 papers have since had their underlying defect fixed by later work elsewhere in this project (the same `repo_check-now-finds-it` pattern as above) and no longer appear in the live disagreement set; Table 3 below shows the 1 still remaining in it, with the full original set described narratively here regardless.

Table 4: Table 3. The rows from Appendix A.6’s original 230-paper categorization corrected on full recheck, from `UNSUPPORTED_REPO` to a supported-platform category, still present in the current disagreement set.

Paper	Category	Evidence and verification
10_1111_1365_2435_14709	SUPPORTED_PLATFORM_BUG	Same root cause: the private Figshare share-link format was not recognized by <code>repo_check</code> . RECHECKED: figshare.com is a fully supported platform; the defect is specifically that a private/anonymous Figshare share-link (figshare.com/s/...) uses a different URL format from the standard figshare.com/articles/... pattern <code>repo_check</code> ’s Figshare matcher expects, so the host itself is supported but this specific link format is not recognised -- a supported-platform bug, not an unsupported platform.

Two platforms cited in this set turned out to be genuinely supported by Metacheck despite `repo_check` finding nothing, each traced to a specific, distinct defect rather than a missing backend. **4TU** (`data.4tu.nl`) has dedicated support (`archive-4tu.R`), and the DOI cited (10.4121/uuid:04b25e70-ca3e-489f-9510-36385c1a4104) live-resolves correctly to a real 4TU record. The ID-extractor correctly pulls the UUID from the DOI, but the live 4TU API lookup queries `data.4tu.nl/v2/articles/<uuid>`, which 404s, because 4TU’s own DOI redirect resolves the identical record to a different, numeric article ID that the UUID does not match. This corrects an earlier, incorrect statement elsewhere in this manuscript (Appendix A.6) that this paper’s 4TU fetch failure remained unresolved and should count as `UNSUPPORTED_REPO`; it is a supported-platform bug, not an unsupported platform. **DANS EASY** is one of four separate subject-specific Dataverse installations DANS operates (archaeology/lifesciences/phys-techsciences/ssh), all sharing a single DOI prefix (10.17026).

Metacheck's `.dataverse_host_from_doi()` resolves any DOI under this shared prefix to `archaeology.datastations.nl`, but the paper in this set cites a dataset actually hosted on `lifesciences.datastations.nl` – a genuine Metacheck defect in DOI-to-host resolution for any DOI prefix shared across multiple Dataverse instances, not specific to DANS. A third paper cites the Knowledge Network for Biocomplexity (`knb.ecoinformatics.org`), a confirmed DataOne member node whose DOI live-resolves correctly; the specific reason `repo_check` did not detect it was not further root-caused within the scope of this review, so it is recorded as a supported-platform miss without a fully diagnosed cause, distinct from the two fully diagnosed bugs above.

Table 5: Table 2. Full paper-by-paper review of the 40 previously-unreviewed ‘unsupported repository, not detected’ disagreements: cited platform, live DOI/URL resolution status, and final category.

Paper	Category	Evidence and verification
10_1002_pan3_10149	NO_REPO_DATA_RESTRICTED	Paper text: 'Our raw data are not publicly available for confidentiality reasons, as they contain social, cultural and language data that may not be interpreted outside of their Nyikina context. They will be archived with Griffith University and Walalakoo Aboriginal Corporation.' No repository URL or DOI is given anywhere in the paper -- data are explicitly withheld (Indigenous data sovereignty). There is nothing for <code>repo_check</code> to find; this is not a Metacheck gap of any kind.
10_1111_1365_2745_12930	UNSUPPORTED_REPO	Paper text: 'Data associated with this paper have been deposited in the Environmental Climate Data Sweden depository https://ecds.se/dataset/species-distribution-modelling-data-for-23-dead-wood-inhabiting-bryophytes '. Live-verified the URL resolves (redirects to researchdata.se). researchdata.se is not in any of Metacheck's supported host lists (not Dataverse, DSpace7, DataOne, Figshare, and has no dedicated single-platform module).
10_1002_pan3_10307	UNSUPPORTED_REPO	Paper text: 'Data and code used in the analyses are available at ScienceBase (https://doi.org/10.5066/P9KBZW0G and https://doi.org/10.5066/P9CDA1UN , respectively)'. Live-verified both DOIs resolve to www.sciencebase.gov catalog items. ScienceBase (USGS) is not in any Metacheck supported host list and has no dedicated module.
10_1111_1365_2745_13067	TEXT_EXTRACTION_TRUNCATED	Extracted text is a mid-sentence fragment: '...net online repository (Lopez-Gonzalez, Lewis, Burkitt, & Phillips, 2011; ...)'. No DOI or URL is present in the extracted text at all -- the actual data-availability statement (likely referring to ForestPlots.net , based on the citation) was not captured by GROBID extraction.
10_1111_1365_2435_14226	UNSUPPORTED_REPO	Paper text: 'The data that support the findings of this study are available on the National Tibetan Plateau Data Center https://doi.org/10.11888/Terre.tpd.c.272843 '. Live check: this DOI request failed to resolve (network/redirect issue querying data.tpd.ac.cn from outside China is common). Regardless, tpdc.ac.cn is not in any Metacheck supported host list and has no dedicated module.
10_1002_2688_8319_12315	TEXT_EXTRACTION_TRUNCATED	Paper text: 'The R code, and Google Earth Engine code, used in all analysis and to produce all figures is archived here: https://zenodo.org/record/8347003 '. Zenodo IS a fully supported platform (<code>R/archive-zenodo.R</code>) and the DOI (10.5281/zenodo.8347003, reconstructed from the record ID) live-resolves correctly (HTTP 200, https://zenodo.org/records/10572530). The extracted text inserts a stray space inside the hostname ('zenodo. org'), a GROBID line-wrap artifact, which breaks Metacheck's host-matching regex against the raw citation text. This is a genuine miss of a supported platform caused by text extraction, not a platform-support gap.

Paper	Category	Evidence and verification
10_1111_1365_2745_14315	UNSUPPORTED_REPO	Paper text: 'Novel data collected in this study and compiled from FRED are available through FRED: https://doi.org/10.25581/ornlsfa.014/1459186'. FRED (via ORNL's Sustainable Forestry/SFA data archive) is a niche, domain-specific archive with no entry in any Metacheck supported host list and no dedicated module. Live DOI resolution failed (ornlsfa.014 DOI prefix appears not to resolve through doi.org currently), which is a separate, secondary problem on top of the platform simply being unsupported.
10_1111_2041_210x_14453	UNSUPPORTED_REPO	Paper text: 'FishScale' software ... can be accessed at https://doi.org/10.5066/P13QVR2R. Data available via https://doi.org/10.5066/P9E1DTYB'. Live-verified: the code DOI resolves to code.usgs.gov (a self-hosted GitLab instance on USGS's own domain -- explicitly out of scope per Metacheck's own GitLab module documentation, which matches only gitlab.com); the data DOI resolves to www.sciencebase.gov (unsupported, as elsewhere in this set). Both citations point to unsupported hosts.
10_1111_1365_2745_13529	NCBI_GENBANK	Paper text: 'Data were deposited in the Sequence Read Archive (https://www.ncbi.nlm.nih.gov/bioproject/PRJNA592420)' and 'Sequence data are deposited at NCBI-SRA under the accession number PRJNA592420.' A domain-specific bioinformatics accession (SRA/BioProject), a distinct category from an ordinary data/code repository and outside repo_check's scope by design.
10_1002_pan3_10743	TEXT_EXTRAC- TION_TRUNCATED	Extracted text is entirely a philosophical/conceptual discussion of aesthetics and definitions of 'mess' -- no data-availability statement was captured anywhere in the extracted text for this paper.
10_1111_1365_2664_12991	TEXT_EXTRAC- TION_TRUNCATED	Extraction returned an empty string -- no text was captured for this paper's data-availability section at all.
10_1002_2688_8319_12241	TEXT_EXTRAC- TION_TRUNCATED	Extraction returned an empty string -- no text was captured for this paper's data-availability section at all.
10_1111_1365_2745_13517	TEXT_EXTRAC- TION_TRUNCATED	Extracted text covers only methodological discussion of pollen/charcoal source areas -- no data-availability statement with a URL or DOI was captured anywhere in the extracted text for this paper.
10_1002_pan3_10677	TEXT_EXTRAC- TION_TRUNCATED	Extraction returned an empty string -- no text was captured for this paper's data-availability section at all.
10_1111_1365_2664_13065	UNSUPPORTED_REPO	Paper text: 'Data are available in the United States Department of Agriculture National Agricultural Library (Ag Data Commons) https://doi.org/10.15482/USDA.ADC/1402049'. Live-verified the DOI resolves (HTTP 202) to agdatacommons.nal.usda.gov. Not in any Metacheck supported host list and has no dedicated module.
10_1111_1365_2745_13154	UNSUPPORTED_REPO	Paper text: 'Data available through Digital Commons @ Georgia Southern repository https://doi.org/10.20429/data.2019.01'. Live-verified the DOI resolves (HTTP 200) to digitalcommons.georgiasouthern.edu. A generic bepress Digital Commons instance, not in any Metacheck supported host list.
10_1111_1365_2745_12749	TEXT_EXTRAC- TION_TRUNCATED	Paper text ends mid-sentence: 'Data, including location, growth, and iWUE data are located at ScholarsArchive@OSU' -- the sentence is cut off before any URL or DOI appears, so no identifier was captured for this citation despite a named repository being mentioned.
10_1111_1365_2435_14477	TEXT_EXTRAC- TION_TRUNCATED	Extraction returned an empty string -- no text was captured for this paper's data-availability section at all.

Paper	Category	Evidence and verification
10_1111_1365_2656_13970	UNSUPPORTED_REPO	Paper text: 'Data available from the Apollo-University of Cambridge Repository https://doi.org/10.17863/CAM.97060 '. Live-verified the DOI resolves (HTTP 200) to www.repository.cam.ac.uk . A DSpace-based institutional repository, but not in Metacheck's DSpace7 or DSpace-legacy host lists and has no dedicated module.
10_1111_1365_2664_13892	UNSUPPORTED_REPO	Paper text: 'The entire dataset that supports the findings of this study is available via a permanent open access data repository of the University of Vienna http://phaidra.univie.ac.at/o:1168756 '. Live-verified the URL resolves (HTTP 200). Phaidra (University of Vienna's repository platform) is not in any Metacheck supported host list and has no dedicated module.
10_1111_2041_210x_13110	TEXT_EXTRAC- TION_TRUNCATED	Extracted text covers only general methodological discussion of uncertainty modelling -- no data-availability statement with a URL or DOI was captured anywhere in the extracted text for this paper.
10_1111_1365_2745_14252	TEXT_EXTRAC- TION_TRUNCATED	Extraction returned an empty string -- no text was captured for this paper's data-availability section at all.
10_1111_1365_2664_14700	UNSUPPORTED_REPO	Paper text: 'Data are available from the OFB Digital Repository https://ged.ofb.fr/share/s/XiFQtejFQy28DG6u0TDzoA/folder '. Live-verified the URL returns HTTP 403 (access-restricted share link) on a custom OFB (French biodiversity agency) Nextcloud-style file-share instance. Not in any Metacheck supported host list and has no dedicated module; a bare access-controlled folder link is also not the kind of structured repository API Metacheck's dedicated modules are built to query even where a host is supported.
10_1111_1365_2656_13918	UNSUPPORTED_REPO	Paper text cites four EDI/PASTA data package DOIs (10.6073/PASTA/... and 10.6073/pasta/..., resolving to portal.edirepository.org) for the paper's actual data, plus a properly-formatted Zenodo DOI (10.5281/zenodo.7737746, live-verified HTTP 200) for the analysis code only. EDI/PASTA is not in any Metacheck supported host list and has no dedicated module. Because the data_availability comparison column concerns the data (not the code), and the data are exclusively on EDI/PASTA, this paper is classified as citing an unsupported repository for its data; the Zenodo code citation is correctly formatted and would be expected to be found by Metacheck's code_check separately.
10_1111_2041_210x_13463	UNSUPPORTED_REPO	Paper text: 'The data associated with this manuscript are published on ESS-DIVE, the Department of Energy's data portal...'. No specific DOI is given in the extracted text (the paper states DOIs are 'documented in Supporting Information file 4', not in the main text). ESS-DIVE is not in any Metacheck supported host list and has no dedicated module; Google Earth Engine (also mentioned) is not a structured data repository Metacheck can query as such.
10_1111_1365_2435_13550	NCBI_GENBANK	Paper text: 'The sequences were deposited in NCBI SRA database (accession no. ... SRP218212 and ... SRP068522)'. A domain-specific bioinformatics accession (SRA), outside repo_check's scope by design, consistent with other SRA/GenBank/BioProject cases in this set.
10_1111_1365_2435_12874	TEXT_EXTRAC- TION_TRUNCATED	Extraction returned an empty string -- no text was captured for this paper's data-availability section at all.
10_1111_1365_2664_14787	UNSUPPORTED_REPO	Paper text: 'Data available from the CSIC repository: http://hdl.handle.net/10261/367361 '. Live-verified the URL resolves (HTTP 200) to digital.csic.es . A DSpace-based institutional repository, but not in Metacheck's DSpace7 or DSpace-legacy host lists and has no dedicated module.

Paper	Category	Evidence and verification
10_1111_1365_2656_14227	UNSUPPORTED_REPO	Paper text cites two repositories: 'Guillemot demographic data ... freely available from Heriot Watt research portal https://doi.org/10.17861/8fd6a194-d1a7-469c-8604-e45b0414f8f0 ' (live-verified HTTP 200, researchportal.hw.ac.uk) and 'Puffin and razorbill demographic data ... available upon request ... from the University of Gloucester online repository https://doi.org/10.46289/8FP68KK9 ' (live-verified HTTP 404, likely a stale/incorrect DOI). Neither host is in any Metacheck supported host list, and neither has a dedicated module.
10_1111_1365_2745_13324	UNSUPPORTED_REPO	Paper text: 'Data used in this publication are available in the appendices and archived at the USU Digital Commons: https://doi.org/10.26078/99w7-eb58 '. Live-verified the DOI resolves (HTTP 200) to digitalcommons.usu.edu. A generic bepress Digital Commons instance, not in any Metacheck supported host list.
10_1111_1365_2656_13272	UNSUPPORTED_REPO	Reference-list citation: 'ScholarWorks Data Repository, https://doi.org/10.18122/bio_data/3/boisestate '. Live-verified the DOI resolves (HTTP 200) to scholarworks.boisestate.edu. A generic bepress ScholarWorks instance, not in any Metacheck supported host list.
10_1111_1365_2435_13791	TEXT_EXTRACTION_TRUNCATED	Extraction returned an empty string -- no text was captured for this paper's data-availability section at all.
10_1111_1365_2745_13849	UNSUPPORTED_REPO	Paper text: 'The data are available as a data package on https://www.forestplots.net/ : https://doi.org/10.5521/forestplots.net/2022_1 '. Live-verified the DOI resolves (HTTP 200) to forestplots.net. Forestplots.net is a domain-specific forest-inventory data archive, not in any Metacheck supported host list and has no dedicated module (a separate, unrelated Dryad DOI also appears in this paper's reference list, but only as a citation for a prior, different dataset -- not this paper's own data).
10_1111_1365_2435_14514	UNSUPPORTED_REPO	Paper text: 'The data that support the findings of this study are available from UNL data archive: https://doi.org/10.32873/unl.dr.20230705 '. Live-verified the DOI resolves (HTTP 202) to sandy.unl.edu. Not in any Metacheck supported host list and has no dedicated module.
10_1111_1365_2435_13936	UNSUPPORTED_REPO	Paper text: 'Data are available from the Illinois Data Bank https://doi.org/10.13012/B2IDB-4547091_V1 '. Live-verified the DOI resolves (HTTP 403, access-gated redirect) to databank.illinois.edu. A Hydra/Fedora-based institutional repository, not in any Metacheck supported host list.
10_1111_1365_2664_13845	UNSUPPORTED_REPO	Paper text: 'Data are available in the US Department of the Interior and U.S. Geological Survey's ScienceBase https://doi.org/10.5066/P94B0WUV '. Live-verified the DOI resolves (HTTP 403) to www.sciencebase.gov. ScienceBase is not in any Metacheck supported host list, consistent with the other ScienceBase cases in this set.
10_1111_1365_2664_14229	SOURCE_CITATION_ERROR	Paper text: 'Data available via Unimi Dataverse http://10.0.51.74/RD_UNIMI/5ZXGIV '. Unimi's real Dataverse host (dataverse.unimi.it) IS a supported Dataverse installation (matches .dataverse_host_regex()), but this paper cites a raw internal/private IP address (10.0.51.74, a non-routable RFC 1918 address) instead of the public hostname -- live-verified this address does not resolve externally at all (request failed outright). The paper's own published citation is unusable as given; this is an error in the source text, not a Metacheck platform-support gap or extraction artefact.

Of these 37 papers, 20 genuinely cite an unsupported platform, 2 cite a domain-specific bioinformatics accession outside `repo_check`'s scope, 0 cite a platform Metacheck genuinely supports but failed to detect for the reasons described above, 13 had no usable citation captured

in the paper’s extracted text at all (either empty extraction or a citation cut off before any identifier appeared), 1 explicitly withheld their data with no repository to find, and 1 paper’s own published citation pointed to an internal, non-public address rather than a reachable repository. These per-paper counts feed directly into the corpus-wide reconciling numbers reported in the main text and into Table 1’s totals above.

A.10 Full paper-by-paper detail for every Metacheck limitation identified in the full disagreement-population review

Table 4 lists every one of the 123 papers in the main text’s “Metacheck limitation” cluster individually: its DOI, which of the nine specific defects (grouped into the three broader categories discussed in the main text) applies, and the evidence-based reasoning for that classification, generated directly from the same per-paper audit file (`manual_repo_findability_audit.csv`) the main text’s numbers are computed from. Because the table is built programmatically from that file rather than transcribed by hand, it stays reconciled with the main text’s counts automatically on every re-render (confirmed by the chunk above: the table’s row count is checked against `n_metacheck_limitation` directly).

Table 6: Table 4. Full paper-by-paper detail for every Metacheck limitation identified in the full disagreement-population review.

Paper (DOI)	Category	Why Metacheck misses it
10.1002/2688-8319.12149	File extension not in Metacheck’s lookup table	The deposit’s two files (both .rtf) are the genuine data file (moose density estimates) and its README; both land on ‘documentation’ because .rtf has no entry in .ext_registry, file_category() returns NA for the data file so Tier 2 doesn’t fire either; Tier 3’s coarse crosswalk maps .rtf’s coarse type ‘text’ to ‘documentation’, which correctly identifies the README but incorrectly also catches the actual data file saved in the same format.
10.1002/pan3.10200	File extension not in Metacheck’s lookup table	All 18 files are ESRI/GDAL GeoTIFF raster sidecar files (.tfw, .tif, .tif.aux.xml, .tif.ovr, .tif.vat.cpg, .tif.vat.dbf) for the paper’s 3 conservation-prioritization output layers. None of these extensions appear anywhere in .ext_registry (data_check_helpers.R). tools::file_ext() truncates to the last dot only, so compound extensions like .tif.aux.xml resolve to ‘xml’. Reran data_classify_files() directly and reproduced the exact observed output: .tif crosswalks to ‘materials’ (generic image bucket), .tif.aux.xml crosswalks via ‘xml’ to ‘code’ (false positive), the rest fall to ‘unknown’. No GIS raster family extension exists anywhere in the registry.
10.1002/pan3.10334	File extension not in Metacheck’s lookup table	All 39 files in the OSF repo are .docx anonymized individual letter transcripts, confirmed by the OSF project’s own title and description to be the actual dataset. .docx has no row in .ext_registry, file_category() returns NA for these filenames, and none contain a Tier-2 keyword token (anonymized site codes, e.g. ‘37yam_3_1.docx’). Tier 3’s coarse crosswalk resolves .docx to coarse type ‘text’, mapped to ‘documentation’. Directly reproduced empirically: data_classify_files() on three real filenames returned ‘documentation’ for all three.

Paper (DOI)	Category	Why Metacheck misses it
10.1111/1365-2435.13183	File extension not in Metacheck's lookup table	data_check correctly typed the repo's one Julia source file (wns-simple.jl) as 'code' via metacheck::file_types, where jl->code is registered. However, code_check (whose output is what the manuscript's n_mc_code aggregate actually counts) gates everything through code_lang(), which looks up language ONLY via .ext_registry\$code_lang -- and that table has no row for extension 'jl' at all (Julia appears only as a possible notebook kernel, never as its own .jl entry). So code_lang() returns NA, the file is never flagged as a code-check candidate, and code_check_repo_urls is empty despite data_check correctly flagging the same file. This is an internal inconsistency between two classifier tables that should agree but don't.
10.1111/1365-2435.13315	File extension not in Metacheck's lookup table	The deposit's only file, DataDryad.7z, is a 7-Zip archive. Metacheck's zip-peek.R explicitly documents that only .zip, the .tar family, and single-file .gz/.bz2/.xz can be opened with base R -- .7z is deliberately left unextracted ('left as other, never fetched') since Metacheck has no libarchive/7z dependency. Since the archive is never opened, data_classify_files() only ever sees the archive's own filename. A second, independent factor compounds this: even the Tier-2 keyword fallback that would normally rescue a misnamed data file fails, since the bounded-token regex requires a separator before 'data' and 'DataDryad.7z' has 'data' glued directly to 'dryad' with no boundary character (confirmed by running the exact regex).
10.1111/1365-2435.13848	File extension not in Metacheck's lookup table	ANSYS Fluent simulation files (.cas/.dat.gz) are genuine simulation data/output files with no entry in .ext_registry, a domain-specific engineering/CFD format gap.
10.1111/1365-2435.13921	File extension not in Metacheck's lookup table	Bruker NMR extensionless raw files plus associated .txt result tables are genuine spectroscopy data/output files not recognized by the classifier's extension rules.
10.1111/1365-2435.14337	File extension not in Metacheck's lookup table	Chen_2023_gross.txt is almost certainly the paper's core numeric dataset (gross soil N mineralization/nitrification rates). .txt is deliberately unlocked at Tier 1. The content-based txt_classify_content() upgrade path only recognizes comma/semicolon/tab/pipe delimiters and explicitly, by documented design, does not detect space/whitespace-delimited tables -- plausibly the file's actual format, so it never upgrades past the default 'documentation' guess. README_Chen_2023_gross.txt correctly resolves via the 'readme' keyword.
10.1111/1365-2656.12732	File extension not in Metacheck's lookup table	A genuine data file named with the word 'dataset' (not the bounded token 'data') was missed by Tier 2's keyword regex, which requires 'data' bounded by separators and does not match 'dataset' or similar compound words.
10.1111/1365-2656.12978	File extension not in Metacheck's lookup table	~65 files named 'P#_###_###.txt' are genuine ant-plant interaction-count matrices (confirmed via Dryad's own file description). .txt is unlocked at Tier 1; Tier 2 keyword rules do not match these filenames. txt_classify_content()'s delimiter sniffer only tests comma/semicolon/tab/pipe and explicitly does not detect whitespace-delimited tables -- exactly the format of these species-interaction matrices -- so all fall to the Tier 3 'text'->documentation crosswalk default.
10.1111/1365-2656.13145	File extension not in Metacheck's lookup table	Single file Male_SES_encounter_history_matrix.txt is the paper's primary mark-recapture dataset (34-yr encounter history matrix, 6,245 individuals, 58,177 records), explicitly titled 'Data from: Encounter history matrix' in the paper's own reference citation. .ext_registry leaves .txt's data_type as NA. Tier 2 keyword rules require a bounded 'data' token; the filename contains none. Tier 3 crosswalks the coarse 'text' filetype to 'documentation'.

Paper (DOI)	Category	Why Metacheck misses it
10.1111/1365-2664.12884	File extension not in Metacheck's lookup table	'Supporting info s7.tar' (931.84 KB) genuinely contains MATLAB scripts, R code, and an Excel occupancy data file (confirmed via live Dryad fetch), matching the paper's own description ('Data S7. Occupancy data files and scripts'). The tar's expansion step did not execute for this file in this run despite being well under all size caps (100MB/500MB) and despite .expand_tar() being designed to handle exactly this format.
10.1111/1365-2664.13096	File extension not in Metacheck's lookup table	Sole file 'EEEworm Model V1.0.0 Johnston et al..nlogo' is NetLogo model source code (confirmed as 'programme code' by the paper's own text). .ext_registry has no 'nlogo' entry anywhere. Tier 2 keyword rules find no match in the filename. Tier 3 coarse crosswalk also has no NetLogo mapping. Directly executed: data_classify_files() on this filename returns 'unknown'.
10.1111/1365-2664.13538	File extension not in Metacheck's lookup table	Single .rar archive genuinely contains R scripts and tabular data files (confirmed via Dryad's own file-content description). .rar is absent from .ext_registry entirely. zip-peek.R's own comment states .rar/.7z 'need an external binary or the libarchive-backed archive package, which metacheck does NOT depend on' -- a deliberately-documented limitation, not an oversight. .is_readable_archive() returns FALSE for 'rar', so it is never downloaded or opened.
10.1111/1365-2664.14288	File extension not in Metacheck's lookup table	Sole file Morford_et al-2022_JAE.zip (18.6 GB) genuinely contains 'data products and XGBoost modeling code...summary data, scripts for downloading annual datasets, and processing code' per the Zenodo depositor's own description. .ext_registry has no row for .zip (archives excluded by design). Whether the zip is opened depends on peek_zips, which defaults to FALSE in both data_check() and reproducibility_check(), unlike repo_check()'s default of TRUE -- this is why repo_check could see the file but data_check, run with its own default, could not look inside it.
10.1111/1365-2664.14746	File extension not in Metacheck's lookup table	Three GeoTIFF spatial-raster data layers (microcorr500m.tif, PAconns500m.tif, betweenness500m.tif) are exactly what the paper's Data Availability Statement names as shared data. .ext_registry has no entry for 'tif/'tiff'. Tier 2 finds no keyword match. Tier 3 resolves via metacheck::file_types to coarse type 'image', and .file_type_crosswalk hard-maps image->'materials' -- a rule for ordinary photographs, not disambiguated for georeferenced raster data.
10.1111/1365-2745.13006	File extension not in Metacheck's lookup table	Sole file 'To submission.rar' genuinely contains research data and R code (Figshare's own description: 'contains research materials including a dataset and analytical code'). .ext_registry has no entry for 'rar'. Archive-expansion in data_check.R only unpacks zip/tar-family/gz/bz2/xz; code comments explicitly state .7z/.rar are not carved out and stay excluded, with repo_check instead warning the author to re-upload as .zip.
10.1111/1365-2745.13537	File extension not in Metacheck's lookup table	Sole file BadenEtAlPlantago2020.zip (15.79 GB) genuinely contains the paper's core research data (confirmed via live Dryad fetch: field monitoring, root histology, soil properties, survival/growth/reproduction data). data_check()'s own default is peek_zips=FALSE, differing from repo_check()'s default of TRUE. The production driver script (run_reproducibility_check.R) calls repo_check then data_check without passing peek_zips to either, so repo_check peeked/listed the zip (default TRUE) but data_check never opened it (default FALSE), leaving the file permanently 'unknown'.
10.1111/2041-210X.13124	File extension not in Metacheck's lookup table	Sole .rar file genuinely contains two GIS raster data files (congruence.grd, pr_Litocranius_walleri.grd), confirmed via Dryad's usage notes. .ext_registry has no entry for 'rar' (or '7z'). data_check.R's archive-expansion comment explicitly states '.7z/.rar are NOT here -- base R cannot read them (repo_check warns instead)' -- a deliberate design limitation.

Paper (DOI)	Category	Why Metacheck misses it
10.1111/2041-210X.13332	File extension not in Metacheck's lookup table	The repo's sole source file is committed with no extension at all, literally named 'C++ code' (confirmed via live GitHub fetch to be genuine Arduino/C++ source), data_check's own classifier already gets this right (Tier 2 keyword match on the bare word 'code', data_type='code'). The actual miss is downstream in code_check.R's code_lang(), which determines language purely via tools::file_ext() matched against .ext_registry\$code_lang; an empty extension matches nothing, so code_lang() returns NA and the file never enters code_check's candidate set, leaving code_check_repo_urls empty despite data_check already having typed the file correctly.
10.1111/2041-210X.13459	File extension not in Metacheck's lookup table	Three small tar.gz archives (DogExperiment.tar.gz, FieldExperiment.tar.gz, OTUs.tar.gz, ~1.07MB total) are genuine DNA-sequencing data (confirmed by the paper's own availability statement: 'All reads and OTU sequences...are available from the Dryad Digital Repository'). .is_tar_archive() correctly recognizes the format, but Dryad's file-byte-download endpoints require an OAuth2 bearer token even for fully public datasets (confirmed live: both file- and dataset-level download endpoints returned 401 with no token) -- so the archives are listed but never downloaded, never expanded, and remain stuck at the pre-expansion 'unknown' placeholder.
10.1111/2041-210X.13979	File extension not in Metacheck's lookup table	Sole file 'R code.txt' genuinely is R source code (confirmed via Zenodo's own title/description: implements five richness estimators). data_check's data_classify_files() already correctly tags it 'code' via a Tier-2 keyword match on 'code' bounded by space/period. But code_check()'s separate code_lang() function keys purely on file extension via .ext_registry\$code_lang: .txt's code_lang is NA by design, so code_lang('R code.txt') returns NA, the file is filtered out of code_check's relevant-languages set, and code_check_repo_urls comes back empty -- a cross-module inconsistency between data_check's keyword logic and code_check's extension-only logic.
10.1002/pan3.6	Recognised extension assigned to the wrong category	All 33 genuine interview-transcript files (CP01.pdf...CP35.pdf, redacted/anonymized, exactly matching the paper's stated raw data) are classified 'documentation'. .pdf has no row in .ext_registry (deliberately left unlocked, since a PDF can be materials/documentation/preprint depending on content). Filenames like 'CP01.pdf' contain no Tier-2 keyword token, so Tier 3's coarse crosswalk fires: metacheck::file_types maps .pdf to coarse type 'text', and the crosswalk maps text->'documentation'. This is a real miss on genuinely primary raw data (anonymized qualitative transcripts), not materials/questionnaires/preprints, with no keyword signal to rescue it.
10.1111/1365-2435.13500	Recognised extension assigned to the wrong category	NOTE: this paper does not exactly fit the stated premise of 'classified none of the files as data or code' -- 5 of 6 files WERE classified 'code', just wrongly. The 5 real data files (SitesMassBins.Rd etc., each 0.9-1.7MB, MIME application/gzip per the Figshare API, i.e. genuine serialized data objects, not R documentation) were saved with extension '.Rd', colliding with R's own .Rd package-documentation-markup extension. .ext_registry has no Tier-1 entry for 'rd'; Tier 3's coarse crosswalk hardcodes ext=='rd' -> type 'code' (confirmed directly), with no check against file size or binary content that would flag these multi-megabyte gzip blobs as implausible R documentation source.
10.1111/1365-2435.13735	Recognised extension assigned to the wrong category	18 genuine simulation-output .txt files (e.g. Coord_row_sim_s0000.txt, Trait_sim_s0000.txt) inside the deposited ZIP are simulation-generated DATA, confirmed by the paper's text and Dryad's own listing, but none were ever assigned data_type=='data'. .txt is deliberately not format-locked in .ext_registry since it can be prose or data; Tier 2 keyword rules require tokens like 'data'/output/results', none of which appear in these filenames; Tier 3's crosswalk routes generic 'text' filetype to 'documentation'. NOTABLY, the RDS snapshot actually used in the manuscript run shows these files as 'code' (not 'documentation' as today's vendored code produces when rerun), indicating the classifier code changed between when the snapshot was generated and now -- a reproducibility concern flagged separately from the root classification gap. Either way, genuine simulation-result data is mistyped into a non-data bucket.

Paper (DOI)	Category	Why Metacheck misses it
10.1111/1365-2435.14060	Recognised extension assigned to the wrong category	8 of 9 files in the Dryad repo (aggression_rate.txt, group_size.txt, N_groups_day.txt, N_individuals_day.txt, steepness.txt, strength_associations.txt, time_groups.txt, time_individuals.txt) are genuine tabular research data but all classified data_type='documentation'. .ext_registry leaves .txt's data_type as NA (deliberately not format-locked). Tier 2 keyword rules only override for bounded name tokens 'data'/code/etc; none of these 8 filenames contain such a token. All 8 fall to Tier 3's coarse crosswalk: filetype() returns 'text', and .file_type_crosswalk['text'] resolves to 'documentation'.
10.1111/1365-2656.12671	Recognised extension assigned to the wrong category	.doc files on Dryad containing genuine research data were classified under a non-data category by the generic document-extension handling.
10.1111/1365-2656.13222	Recognised extension assigned to the wrong category	CTs.txt is the primary data file (critical thermal limits for juvenile Rana temporaria, confirmed via Dryad's own dataset description). .txt is left unlocked at Tier 1. Tier 2 keyword rules find no match in the filename 'CTs' (short for 'critical thermals'). Tier 3 crosswalks the coarse 'text' filetype to 'documentation'. Verified directly: data_classify_files('CTs.txt') returns 'documentation'.
10.1111/1365-2664.14504	Recognised extension assigned to the wrong category	Four GeoTIFF raster layers (DailyAbsorpWinter_AllConflict.tif, DailyAbsorpWinter_Livestock.tif, FidelityWinter.tif, ResistanceWinter.tif) are genuine SAMC model input layers for a wolf-recolonization connectivity analysis. .ext_registry has no row for '.tif'/.tiff. Tier 2 finds no keyword match. Tier 3's coarse crosswalk buckets .tif under coarse type 'image', and .file_type_crosswalk hardcodes image->'materials' unconditionally -- a rule built for photographs/figures, not GIS raster data (unlike .nii/.dcm/.mnc/.mgz, which the registry already special-cases to 'data').
10.1111/2041-210X.14457	Recognised extension assigned to the wrong category	64 .JPG drone photographs are the paper's primary measurement data for 3D structure-from-motion photogrammetry (explicitly: 'Data required to recreate analysis with ground truthing images'), not participant-facing stimuli. .ext_registry has no entry for 'jpg'/'jpeg'; classification falls to the coarse crosswalk, which hard-codes image->'materials' -- a rule correct for the common case (a stimulus photo) but wrong for this domain where photographs are literally the quantitative primary data.
10.1111/1365-2435.14350	Archive never downloaded/opened (pipeline configuration mismatch)	The repo's Wieczynski_etal_2023_FE_Master.tar.gz genuinely contains a Mathematica model-code notebook plus 30 .csv data files (confirmed by downloading and unpacking it directly). .tar.gz IS recognized as a readable archive by .is_tar_archive(), but under the production download='data' mode, the baseline want formula never sets want=TRUE for an 'unknown'-typed row, and the rescue line 'want <- want & (!(data_type %in% never_fetch) keep_archive)' is an AND that can only narrow an already-true want, never promote a false one to true -- confirmed by re-executing the formula in R. So the archive is never downloaded or expanded.
10.1111/1365-2664.13795	Repository file listing did not walk into subfolders	Paper's text states 'scripts used to calculate functional indices are available at' the biomonitoR GitHub repo. data_check_files lists only 9 root-level files; live GitHub API shows the real repo root also has DESCRIPTION, NAMESPACE, and 96+ real .R source files under R/ (Pi.R, Tree.R, abundance.R, etc.), none enumerated by Metacheck. .is_r_package_file()'s own documented logic says a root-level package is NOT excluded when no other data/code-shaped content exists outside it -- exactly this case -- so the R/ files should have been kept and classified 'code' via the existing .ext_registry rule for .R, but the crawl never descended into subdirectories at all.

Paper (DOI)	Category	Why Metacheck misses it
10.1111/2041-210X.12784	Repository file listing did not walk into subfolders	data_check_files lists only 5 root-level files for the frair GitHub repo. Live GitHub Trees API (recursive) shows the real repository has 45+ additional files including DESCRIPTION, NAMESPACE, 24 .R source files under R/, and data/gammarus.RData + data/bythotrephes.RData -- exactly the two datasets the paper's text names as 'data used in this manuscript.' .ext_registry's rules for .RData->'data' and .R->'code' are correct and sufficient; one telling detail (rhub.yaml appears as a bare top-level name in Metacheck's list but actually lives at .github/workflows/rhub.yaml in the live repo) confirms the enumeration was shallow/non-recursive, not a classifier-rule bug.
10.1111/2041-210X.13009	Repository file listing did not walk into subfolders	data_check_files lists only 11 root-level files of the tRophicPosition GitHub repo. Live fetch shows the real repo also contains R/ (46 .R source files) and data/ (5 .rda example datasets), none enumerated. .ext_registry already correctly maps .r->'code' and .rda/.rdata/.rds->'data' -- the classification rules are not at fault; the file-enumeration step simply never walked into subdirectories.
10.1111/2041-210X.13503	Repository file listing did not walk into subfolders	data_check_files lists only 7 files for the actel GitHub repo. Live GitHub API tree (recursive) shows the real repository has 332 files including DESCRIPTION, NAMESPACE, and 19 genuine .R source files under R/, none enumerated. .ext_registry already correctly maps .r->'code', and a dedicated function .is_r_package_file() is specifically designed to NOT exclude a root-level R package's own files when the package IS the deposit -- the classification rules are correct; the crawl simply never walked into R/, man/, tests/, or vignettes/.
10.1111/2041-210X.13750	Repository file listing did not walk into subfolders	data_check_files lists only 13 root-level files for the RevGadgets GitHub repo. Live GitHub API shows the real repo root also has DESCRIPTION, NAMESPACE, and subdirectories R/ (39 genuine .R source files), tests/, inst/, man/, revdep/, none reflected in Metacheck's listing. .ext_registry's format-lock for .R->'code' is correct and sufficient; the files were simply never in the enumerated list.
10.1002/pan3.10756	Supported platform, but a specific detection defect	figshare_links() correctly detects the DOI citation via the institutional vanity-host/DOI-prefix table (repository.mmu.ac.uk / 10.23634, both explicitly registered in .figshare_vanity_hosts() and .figshare_doi_prefix_hosts()), and live, direct invocation of the installed package's repo_check() on this single paper today DOES detect the repo (repo_n = 2, two duplicate entries differing only by a trailing citation-punctuation period) -- directly contradicting the repo_n = 0 stored in res_repo_check.RData. The live run's gated_repos shows repo_error = 'arguments imply differing number of rows' for this paper's own figshare id lookup (one live attempt reported '00636962 could not be found on api.figshare.com'), suggesting the same batch-level tryCatch/data.frame-construction fragility identified for 10.1111/1365-2745.13728 is also implicated here, compounded by apparent intermittency in api.figshare.com's own response reliability for this specific institutional-DOI lookup path (a second live attempt via a different code path, querying the article id 32616504 directly, succeeded). The stored repo_n = 0 is best explained as a stale/incorrect batch-run artifact layered on top of a real but intermittent figshare-institutional-DOI lookup fragility, not a permanent inability to detect or resolve this repository.

Paper (DOI)	Category	Why Metacheck misses it
10.1111/1365-2435.12975	Supported platform, but a specific detection defect	The prior audit's claim that 'metacheck_any_detector_hit was true, no gap' is true only at the level of link extraction/host recognition, not file enumeration -- and files_n = 0 proves no files were ever listed for this paper. Re-verified against metacheck's actual installed source (.figshare_id(), figshare_links(), and inst/modules/repo_check.R): the URL shape figshare.com/s/<token> carries an opaque share token with no numeric Figshare article ID anywhere in it, and Figshare's public API has no endpoint that resolves a third party's private-link token to article metadata/files (the private_links API family is for the owning account managing its own links). metacheck's figshare_links() correctly extracts the URL and repo_check.R correctly records repo_error = 'private share link (figshare.com/s/...) cannot be resolved without a browser' via its gated_repos table rather than silently reporting zero links -- this is a documented, by-design limitation of the /s/ share-link mechanism itself, not a missing config entry or an extraction failure. It still means real data that a human reader can view is entirely invisible to metacheck's downstream file counts, so it is classified as a platform-support bug/limitation (SUPPORTED_PLATFORM_BUG) rather than a false 'no gap'.
10.1111/1365-2435.13103	Supported platform, but a specific detection defect	files_n = 0 directly contradicts the prior audit's assumption that detection succeeding meant the gap was closed. The resolvable DOI (10.4225/...) is a genuine Bridges/Figshare DOI that metacheck's dataverse/figshare DOI-prefix tables should in principle be able to query for file listings, but because the concatenated /s/ share-link token was also extracted as a second repo_url and is unresolvable by design (opaque token, no API endpoint for third-party resolution -- same limitation as .figshare_id() documents elsewhere), and because the batch-level figshare DOI retrieval call that would otherwise fetch files for the whole corpus is itself broken by a separate code defect (see 10.1111/1365-2664.14575 below, which shares the identical 'arguments imply differing number of rows' failure signature across the entire figshare-DOI batch), no files were enumerated for either repo_url belonging to this paper. This combines two real, separate code-level gaps -- the /s/ token's unresolvable-by-design limitation, and the corpus-wide figshare-DOI batch crash -- both of which are genuine platform bugs, not a non-issue.
10.1111/1365-2435.13360	Supported platform, but a specific detection defect	metacheck fully supports Dataverse as a platform, and dataverse.nl specifically is one of its individually-verified known hosts (R/archive-dataverse.R, .dataverse_hosts() line 64), with its DOI prefix 10.34894 even already mapped to it in .dataverse_doi_prefix_hosts() (line 303). The live chain confirms this paper's dataset genuinely resolves to dataverse.nl under DOI 10.34894/UUNBQQ, which is exactly the registered prefix. The gap is that the paper's text never contains either the host string "dataverse.nl" or the DOI string "10.34894" -- it cites the dataset only as the bare handle-resolver URL "https://hdl.handle.net/10411/6Q1YDC". dataverse_links() has three matching passes (real-hyperlink host match, bare-text host match, bare-DOI-prefix match), and all three require either the dataverse.nl hostname or a 10.34894-prefixed DOI to literally appear in the extracted text/href -- neither does, since handle.net URLs use their own handle-prefix namespace (here 10411, DANS's generic handle prefix) which is unrelated to the target repository's DOI prefix. Metacheck does have precedent for resolving a specific hdl.handle.net prefix to a platform (archive-psycharchives.R maps handle prefix 20.500.12034 to PsychArchives), but no equivalent mapping exists for DANS's 10411 prefix pointing at DataverseNL (or at DANS's other dataverse-based Data Stations). Resolving this would require either (a) adding 10411 as a recognized handle prefix that redirects to a Dataverse-family host, or (b) following the live HTTP redirect chain for any hdl.handle.net citation before applying the existing host/DOI-prefix matchers -- neither of which the current static-text extraction does for this handle prefix. [RECLASSIFIED from SUPPORTED_PLATFORM_CONFIG_GAP: not a host/DOI-prefix table entry -- requires a new resolution mechanism (handle-prefix redirect, shortlink expansion, or DOI-to-host resolution) or the host runs an API shape metacheck does not support.]

Paper (DOI)	Category	Why Metacheck misses it
10.1111/1365-2435.13850	Supported platform, but a specific detection defect	metacheck fully supports Figshare and even has explicit, dedicated support for exactly this host and DOI prefix: R/archive-figshare.R's .figshare_vanity_hosts() lists "smithsonian.figshare.com" by name (line 97), and .figshare_doi_prefix_hosts() maps DOI prefix "10.25573" to "smithsonian.figshare.com" (line 136). The citation text is also cleanly extractable with no garbling. The actual bug is in ID-type resolution: the cited DOI, https://doi.org/10.25573/data.c.5073065.v1, is a Figshare COLLECTION DOI (identifiable by the literal ".c." segment after the DOI prefix, the same marker .figshare_collection_id() uses for the generic 10.6084/m9.figshare.c.<id> case), but .figshare_id()'s institutional-prefix pattern (archive-figshare.R lines 327-344) does not exclude or recognize this ".c." collection marker for institutional prefixes -- it simply strips any non-digit-containing sub-segment and returns the trailing digits, "5073065", as if it were an ARTICLE id. I confirmed this live: devtools::load_all() then .figshare_id(url) returns "5073065", but .figshare_collection_id(url) returns NA (its own pattern only checks the literal 10.6084 prefix, not institutional ones like 10.25573, per the comment at lines 460-463 of the same file: "Only the literal 10.6084 prefix is checked here... no institutional-collection-DOI citation has been found yet"). Because figshare_info() checks ids\$figshare_id first and only falls through to collection-expansion logic when figshare_id is NA (line 591: "unresolved <- is.na(ids\$figshare_id)"), it never tries the collection path for this URL -- it instead calls GET https://api.figshare.com/v2/articles/5073065 directly, which I confirmed live returns HTTP 404 (while GET https://api.figshare.com/v2/collections/5073065 returns HTTP 200). Running figshare_info() on this exact URL in the live package reproduces the failure: it returns a row with error = "unfound" and a warning "5073065 could not be found on api.figshare.com". This directly contradicts the payload's metacheck_any_detector_hit = true for this paper: the one link actually cited in the text is not resolvable by the current code, because institutional-prefix DOIs for Figshare collections are not yet recognized as collections the way bare 10.6084/m9.figshare.c.<id> DOIs are. CONFIRMED CODE BUG, still present. dryad_links() DOES find the link (datadryad.org/stash/share/<hash> matches the host filter in R/archive-dryad.R's found_href block), but .dryad_doi() -- the function that extracts a resolvable Dryad DOI from the URL -- has NO regex pattern at all for Dryad's '/stash/share/<opaque-hash>' reviewer-link URL shape (confirmed by direct test: .dryad_doi() on this exact share-link URL returns NA). Its two patterns only cover '10.5061/dryad.<id>' and the alternate-prefix bare-DOI form; neither matches a share-link hash. Downstream, dryad_info() short-circuits with 'No valid Dryad links' whenever every dryad_doi in the batch is NA, so zero files are ever listed even though the link itself was detected and the dataset is live. Live-reran module_run() on this exact paper just now reproduces repo_n=1 (link detected) but files_n=0, matching this exact mechanism.
10.1111/1365-2435.14022	Supported platform, but a specific detection defect	

Paper (DOI)	Category	Why Metacheck misses it
10.1111/1365-2435.14284	Supported platform, but a specific detection defect	Zenodo is a fully supported platform (R/archive-zenodo.R) and the Zenodo DOI is extractable: I located the exact string "https://doi.org/10.5281/zenodo.7573829" in full_text, clean and un-garbled, matching the coder's recorded link exactly. The data-availability sentence itself reads "All data and associated scripts used in this study are available on" with the actual link/citation marker stripped out (a GROBID artifact from resolving an inline citation) -- the literal DOI string only survives later, embedded in the formatted bibliography as the self-citation entry "Shah, A. A., Hotaling, S., ... (2023). ... Zenodo. https://doi.org/10.5281/zenodo.7573829". This is why metacheck's own detector misses it despite the platform and pattern being fully supported: zenodo_links() in R/archive-zenodo.R calls text_search(paper, zen_bare_regex, return = "match", perl = TRUE) without passing include_refs = TRUE, and text_search() in R/text_search.R defaults include_refs to FALSE, which explicitly filters out any text row where section_type == "references" (lines 104-109). Since the only textual occurrence of this DOI in the paper is inside a reference-list entry (and the paper's extracted url/href table apparently has no corresponding hyperlink either, since found_href would otherwise have caught it regardless of section), the bare-DOI body-text fallback in zenodo_links() never sees it, and the detector returns no hit even though both the platform support and the raw text content are present and intact. This is a genuine config/argument gap in a supported detector, not a missing platform or a garbled citation.
10.1111/1365-2435.14709	Supported platform, but a specific detection defect	Re-verified against metacheck's installed R/archive-figshare.R: figshare_links() does extract this URL (hence a detector hit), but .figshare_id() returns NA for /s/<token> share links because the token carries no numeric Figshare article ID, and no plain-HTTP endpoint exists for resolving a third party's private share link. metacheck therefore cannot enumerate this repository's files -- files_n = 0 confirms no file listing was ever produced for this paper, contradicting the prior audit's 'no gap' conclusion, which treated a mention-level detector hit as equivalent to successful file enumeration. The repo_error is correctly surfaced in gated_repos rather than silently hidden, so this is a known, by-design limitation of the Figshare /s/ share-link mechanism (unresolvable by any third party without a JS browser), classified as a platform-support bug/limitation rather than a false non-gap.
10.1111/1365-2656.13005	Supported platform, but a specific detection defect	The data-availability sentence in the paper's full text reads: "The data used in this study are available from the Dryad Digital [Repository]... Repository: https ://doi.org/10.5061/dryad.m20hf10 (Naniwadekar, Chaplod, Datta, Rathore, & Sridhar, 2019)." metacheck has a dedicated Dryad detector (R/archive-dryad.R), and 10.5061 is the FIRST and primary entry in .dryad_doi_prefixes(), with its own purpose-built body-text regex "(?:https?:/)(?:doi\ .org/)?10\ .5061/dryad\ \ .[A-Za-z0-9]+" in dryad_bare_regex. I ran this exact regex in R against the exact text string from the paper (including the stray "https :/" space artifact from PDF extraction) and it matched successfully, returning "doi.org/10.5061/dryad.m20hf10" -- the leading optional "https?:/" group simply fails to consume the malformed "https :/" and the engine falls through to matching starting at "doi.org/10.5061/...", which is literally present with no intervening whitespace. Since the regex is not the point of failure on the text I was given, yet metacheck_any_detector_hit is reported as False for this paper, the gap must lie elsewhere in the pipeline between GROBID text extraction and the detector call (e.g. the real GROBID output differing from this harness's full_text in whitespace/hyphenation around the DOI, or a failure in the live dryad_info() lookup/repo_check step rather than in link detection itself). This is a supported-platform bug rather than a missing-text or platform-gap issue: the correct, resolvable DOI is present in the text in a form the detector's own regex is proven to match, and the live Dryad page is fully public with the exact matching dataset.

Paper (DOI)	Category	Why Metacheck misses it
10.1111/1365-2656.13986	Supported platform, but a specific detection defect	The prior audit correctly identified this as a clean, unambiguous figshare.com URL with a genuine detector hit, but incorrectly concluded 'no gap' on that basis alone. <code>files_n = 0</code> shows no files were enumerated. As with the other /s/-token cases in this batch, metacheck's <code>.figshare_id()</code> cannot extract a resolvable article ID from an opaque share token, and Figshare provides no third-party API path to resolve it -- a documented, by-design platform limitation (not a missing config entry, not an extraction failure) that nonetheless means real, likely-public data is invisible to metacheck's file counts for this paper.
10.1111/1365-2656.14169	Supported platform, but a specific detection defect	The prior audit's own reasoning already correctly described the mechanism -- <code>.figshare_id()</code> returns NA for /s/ tokens, <code>figshare_links()</code> sets a <code>figshare_unsupported</code>-style flag, and <code>repo_check.R</code> surfaces <code>repo_error = 'private share link ... cannot be resolved without a browser'</code> -- but then wrongly filed this under <code>CORRECTLY_DETECTABLE_NO_GAP_FOUND</code>. <code>files_n = 0</code> is definitive: no files were enumerated, so real data potentially visible to any human with a browser is completely invisible in metacheck's output for this paper. This is the correct <code>SUPPORTED_PLATFORM_BUG</code> case: figshare is a supported platform, the repo is detected, but a specific, well-documented limitation (opaque share tokens have no third-party-resolvable article ID) prevents file enumeration.
10.1111/1365-2664.13980	Supported platform, but a specific detection defect	CONFIRMED CODE BUG, same mechanism as the 14022 Dryad case above (different paper, same root cause). The citation reads 'Data available via the Dryad Digital Repository' followed by a 'datadryad.org/stash/share/<hash>' link -- another reviewer share-link. <code>dryad_links()</code> detects the href (host match), but <code>.dryad_doi()</code> again returns NA for this URL shape (no stash/share pattern exists in its regex), so <code>dryad_info()</code> lists zero files. Live-reran <code>module_run()</code> on this exact paper just now reproduces <code>repo_n=1/files_n=0</code>, matching this mechanism exactly. Note: the prior audit's text described a DIFFERENT, permanent DOI form (10.5061/dryad.t76hdr81n) as an 'additional link' supposedly also present; that exact string is NOT found anywhere in the currently extracted text (confirmed by direct search) -- only the stash/share link is. Whether that permanent DOI ever appeared in an earlier extraction or the prior audit's claim was simply inaccurate could not be determined from current data; it does not change this paper's classification, since the citation actually present now is the undetectable share-link form either way.
10.1111/1365-2664.14704	Supported platform, but a specific detection defect	This is the same corpus-wide batch-poisoning defect identified for 10.1111/1365-2664.14575, affecting the Dataverse platform instead of Figshare. Checking <code>res_repo_check\$gated_repos</code> shows 'arguments imply differing number of rows' (with varying row counts) attached to all 103 dataverse-flagged <code>repo_url</code> rows across the entire corpus -- i.e., every single Dataverse citation in the manuscript's full dataset failed in this pipeline run. Reading <code>inst/modules/repo_check.R</code>'s Dataverse block confirms <code>dataverse_info()</code> is called once for the whole corpus's dataverse URL batch inside a single <code>tryCatch</code> (lines ~861-927), with the Dataverse module's own code comments explicitly anticipating and partially guarding against exactly this failure mode ('a field can come back as a length-zero value...which would...break this <code>data.frame()</code> call with arguments imply differing number of rows'), but the guard (<code>%empty_or%</code>) is not applied everywhere a field could be malformed across the full corpus of real-world API responses. Calling <code>.dataverse_info()</code> directly for THIS paper's host/DOI in isolation succeeds cleanly with no error, proving the defect is not specific to this paper's own data -- some other paper's Dataverse record elsewhere in the same batch triggered the crash, and the caught error message was then stamped onto every Dataverse <code>repo_url</code> in the corpus via <code>repos\$repo_error[repos\$repo_url %in% dv_urls] <- conditionMessage(e)</code>. This paper's data is fully public, free, and CC-BY-licensed, and metacheck's own underlying retrieval code can read it successfully when called directly -- the failure is purely an artifact of the module's all-or-nothing batch error handling.

Paper (DOI)	Category	Why Metacheck misses it
10.1111/1365-2745.12882	Supported platform, but a specific detection defect	Two distinct issues compound here. First, a narrower, previously-identified limitation (documented in the prior audit and still accurate): the citation uses a bare urn:uuid identifier with no DOI anywhere in the data-availability statement, and metacheck's <code>.dataone_pid()</code> (R/archive-dataone.R) only has regex branches for doi:10.x/... and bare-DOI forms, not for a top-level urn:uuid citation PID -- so <code>dataone_info()</code>'s validity filter (requiring both a non-NA host AND a non-NA pid) would exclude this specific citation from its live per-record API lookup regardless of any other bug. Second, and more decisively given that <code>files_n = 0</code> for every one of these 8 papers: <code>res_repo_check\$gated_repos</code> shows the SAME 'arguments imply differing number of rows' error (with varying counts) attached to every dataone-flagged <code>repo_url</code> across the corpus (4 rows total, from only 2-3 papers using DataONE), exactly mirroring the dataverse and figshare-DOI batch-poisoning defect in <code>inst/modules/repo_check.R</code>'s DataONE block (<code>dataone_info()</code> called once per corpus batch inside one <code>tryCatch</code>; one bad record anywhere in the batch crashes the shared call, and the error message is stamped onto every dataone <code>repo_url</code> via <code>repos\$repo_error[repos\$repo_url %in% dataone_urls] <- conditionMessage(e)</code>). Because this paper's own citation would already be filtered out of the 'valid' (<code>pid-resolvable</code>) set before ever reaching the API call, the specific batch crash could in principle have originated from one of the OTHER dataone-citing papers' (the two 10.5063/f1... DOIs) records. Either way, <code>files_n = 0</code> is confirmed, and the correct classification is a genuine code-level platform bug (both the urn:uuid PID-extraction gap and the batch-wide crash handling), not a non-issue. metacheck fully supports Dryad (R/archive-dryad.R), and its bare-DOI regex <code>`10\..5061/dryad\.[A-Za-z0-9]+`</code> (<code>dryad_links()</code>, line 78) is written correctly and would match the text as printed, extracting "10.5061/dryad.745" as a Dryad hit -- which is why <code>metacheck_any_detector_hit</code> is true and the sanity check is not a clean pass. The real problem lies upstream of the regex, in GROBID's PDF text extraction: the paper's actual DOI is 10.5061/dryad.745g830, but GROBID inserted a spurious internal artifact, "file://localhost/opt/grobid/grobid-home/tmp/g830", directly between "745" and the true suffix "g830" in the extracted <code>full_text</code>, so the DOI appears split into "10.5061/dryad.745" followed by an unrelated local file path that happens to end in the string "g830" that truly belongs to the DOI. metacheck's regex, which has no way to know a PDF-extraction artifact sits in the middle of a DOI, correctly matches only the truncated "10.5061/dryad.745" -- which independently resolves to an HTTP 404, i.e. a dead DOI. So metacheck's detector registers a syntactic hit, but on a non-resolving link, not on the real, live, public dataset (10.5061/dryad.745g830) that a manual check finds. This is a text-extraction/GROBID artifact bug feeding a correctly-implemented Dryad detector, not a missing platform or missing regex case, and the human coder independently transcribed the exact same garbled DOI from the same corrupted source text, so this is not an extraction-pipeline issue unique to metacheck's own parsing either -- it is inherited directly from the GROBID output both relied on.
10.1111/1365-2745.13085	Supported platform, but a specific detection defect	metacheck's <code>figshare_links()</code> correctly detects this URL as a figshare mention (<code>figshare_unsupported = FALSE</code> at the detection stage), but <code>.figshare_id()</code> (inspected directly in the installed package source) has no regex pattern that can extract a numeric article id from a figshare.com/s/<hash> private share-link token -- only DOI-based, <code>/articles/</code>, and <code>ndownloader.figshare.com</code> paths are handled. This is a genuine, structural inability, but it is explicitly known and intentionally handled: <code>repo_check.R</code> (lines ~169-182) contains a dedicated comment block stating this exact case is flagged via <code>repo_error</code> ("private share link (figshare.com/s/...) cannot be resolved without a browser") specifically so it is distinguishable from 'no link found' or 'malformed link', because the hash is opaque and figshare's own resolution page sits behind an AWS WAF bot challenge with no documented plain-HTTP API alternative. Confirmed live and reproduced today by directly invoking the installed package's <code>figshare_id()/repo_check()</code> functions.
10.1111/1365-2745.13653	Supported platform, but a specific detection defect	

Paper (DOI)	Category	Why Metacheck misses it
10.1111/1365-2745.14098	Supported platform, but a specific detection defect	IMPORTANT CORRECTION to an earlier, unverified pass that claimed the GitHub link was 'genuinely inaccessible, a dead link' -- this is live-verified FALSE. Two citations exist: (1) a tinyurl.com shortlink to a GitHub blob URL, which Metacheck's <code>github_links()</code> cannot match because it only regex-matches literal github.com strings and performs no shortlink expansion anywhere in the codebase; (2) a bare DOI (<code>doi.org/10.3929/ethz-b-000597256</code>) resolving to a fully-supported DSpace-7 host (<code>research-collection.ethz.ch</code>, explicitly listed in <code>.dspace7_hosts()</code>), but <code>dspace7_links()</code> only matches literal hostname strings and never resolves a bare DOI to its target host first -- no DOI-resolution-to-host step exists anywhere in <code>repo_check.R</code> or <code>doi.R</code>. [RECLASSIFIED from <code>SUPPORTED_PLATFORM_CONFIG_GAP</code>: not a host/DOI-prefix table entry -- requires a new resolution mechanism (handle-prefix redirect, shortlink expansion, or DOI-to-host resolution) or the host runs an API shape metacheck does not support.] Metacheck has a dedicated Dryad module (<code>R/archive-dryad.R</code>) with <code>dryad_links()</code>/<code>dryad_info()</code>, wired into <code>inst/modules/repo_check.R</code>, and the module's own bare-DOI regex <code>("(?:https?://)?(?:doi \.org)/?10\\.5061/dryad\\.[A-Za-z0-9]+")</code> is a perfect, standard match for this citation -- I ran the exact regex from the source file against both the short body-text fragment ("Dryad Digital Repository, <code>https://doi.org/10.5061/dryad.g66247g</code>") and the full reference-list sentence from this paper, and in both cases it correctly extracted <code>"https://doi.org/10.5061/dryad.g66247g"</code> verbatim. The live DOI resolves (302) to a fully public, non-login <code>datadryad.org</code> dataset page with three files listed, so the repository itself is genuinely live and accessible -- this is not a dead link. Since the regex match succeeds in isolation but <code>metacheck_any_detector_hit</code> was reported false for the real pipeline run on this paper, the gap is not that Dryad is unsupported or that this DOI prefix/pattern is unregistered (<code>10.5061/dryad.</code> is exactly the pattern <code>archive-dryad.R</code> is built around); it must instead be that <code>dryad_links()</code>'s two input sources -- <code>paper_table(paper, "url")</code> (real anchor-tag hrefs extracted from the source PDF/XML) and <code>text_search(paper, ...)</code> over metacheck's own parsed-paper text representation -- did not surface this particular sentence to the regex in the actual extraction pipeline (e.g., a GROBID/PDF-parsing segmentation artifact placing the citation in a block <code>text_search</code> does not scan, or the real paper's extracted text differing from the flat <code>full_text</code> string used here). Because the detector logic is demonstrably correct and the platform is supported, but the real run still missed a citation its own regex can parse, I classify this as a <code>SUPPORTED_PLATFORM_BUG</code> (an upstream text/URL-extraction issue feeding the Dryad detector) rather than a missing config or missing platform.
10.1111/2041-210X.13073	Supported platform, but a specific detection defect	

Paper (DOI)	Category	Why Metacheck misses it
10.1111/2041-210X.13262	Supported platform, but a specific detection defect	metacheck does have a GitLab detector (R/archive-gitlab.R), but I read its source directly and its own docstring states plainly: "Only gitlab.com is matched; a self-hosted GitLab instance's own domain is" not matched (the hostname-matching logic at the top of the file hardcodes "gitlab.com" throughout, e.g. the regex construction and every API call target "https://gitlab.com/..."). The DOI in this paper (https://doi.org/10.5066/P99JX2ND) resolves via a 302 redirect to code.usgs.gov, a self-hosted GitLab instance on USGS's own domain, not gitlab.com. Since metacheck's GitLab backend is hardcoded to gitlab.com only, it would not recognize code.usgs.gov URLs as GitLab at all, even though the live page is unambiguously running GitLab software. This is a config/host-registration gap in an otherwise-supported platform family, not a missing backend. Separately, even if metacheck matched this host, it would be checking a repository that genuinely requires authentication for non-member access (confirmed live), so flagging it as inaccessible would actually be correct, not a detector failure. The paper's second, arguably more important, repository link (the NPS IRMA DataStore page, confirmed publicly accessible via its REST API with the actual dataset .zip file) sits on a platform (IRMA DataStore) that is not on metacheck's supported list at all (dataone, dataverse, dryad, dspace/dspace7, figshare, github, gitlab, mendeley, osf, psycharchives, rbox, researchdata4tu, reshare, zenodo, aspredicted) — IRMA DataStore is a bespoke NPS Ext JS application with its own REST API shape (irmaservices.nps.gov/datastore/v7/rest/...), not a thin wrapper around any of those listed platforms, so that part of the gap is better described as a second, additional config/platform gap rather than a text-extraction problem, since both links are clearly present and correctly worded in the extracted full_text. [RECLASSIFIED from SUPPORTED_PLATFORM_CONFIG_GAP: self-hosted GitLab is out of scope by explicit design (gitlab_links() hardcodes gitlab.com only), not a missing host-list entry.] Paper cites 'https://code.usgs.gov/vtcfwru/amonitor' -- a self-hosted GitLab instance run by USGS. Metacheck's archive-gitlab.R implements a full GitLab backend but its gitlab_links() regex is scoped to the gitlab.com hostname only, by explicit documented design, since self-hosted GitLab instances cannot be distinguished from arbitrary institutional websites by URL pattern alone. A secondary GitHub link for a sub-workflow is also present but is extracted with inserted spaces ('cbala ntic/tempo rally -adapt ivesampling'), a compounding issue. [RECLASSIFIED from SUPPORTED_PLATFORM_CONFIG_GAP: self-hosted GitLab is out of scope by explicit design (gitlab_links() hardcodes gitlab.com only), not a missing host-list entry.]
10.1111/2041-210X.13397	Supported platform, but a specific detection defect	CONFIRMED CODE BUG, still present. The paper's data-availability sentence ('...has been made publicly available on http://Zenodo.org') cites only the generic Zenodo homepage; the ACTUAL resolvable citation, the Zenodo DOI (10.5281/zenodo.3826426), appears only in the reference list entry for Arje et al. (2020). zenodo_links() DOES include a working regex for exactly this DOI shape (confirmed by direct test: .zenodo_id() correctly resolves this exact DOI string to id 3826426 when tested directly), but it calls text_search(paper, zen_bare_regex, return='match', perl=TRUE) WITHOUT passing include_refs=TRUE. text_search()'s default (include_refs=FALSE, R/text_search.R) explicitly filters out any text row whose section_type is 'references' before searching -- confirmed directly: the row containing the DOI has section_type='references', and calling text_search() on this paper with the same regex returns 0 rows, while testing the identical regex directly against the raw row string (bypassing text_search()'s section filter) matches it immediately. So the one genuinely resolvable Zenodo citation in this paper is systematically invisible to the detector because it sits in the reference list rather than the body text, and only the unresolvable bare homepage mention is ever found instead.
10.1111/2041-210X.13428	Supported platform, but a specific detection defect	

Paper (DOI)	Category	Why Metacheck misses it
10.1111/2041-210X.13768	Supported platform, but a specific detection defect	Paper's formal DAS text is scrambled by PDF-extraction reflow, but the code deposit self-citation (DOI 10.5281/zenodo.5562809) survives intact in the reference list. Metacheck's <code>zenodo_links()</code> has a documented bare-DOI fallback regex specifically built for exactly this citation style, and the DOI is neither garbled nor truncated. Live-verified the deposit is public, the zip supports HTTP range requests, and <code>data_check_helpers.R</code> explicitly classifies <code>.py</code> files as <code>'code'</code> -- every mechanism needed should work. No specific bug could be pinpointed from static source review alone; this would need an actual pipeline rerun to confirm where/if it diverges.
10.1111/2041-210X.13785	Supported platform, but a specific detection defect	Zenodo is a fully supported platform in metacheck, implemented in <code>R/archive-zenodo.R</code>. The availability statement reads "Code and data for all analyses (Xirocostas et al., 2021) and the ZAX Herbivory Trainer (Debono, 2021) are archived on Zenodo," and the reference list spells out both DOIs cleanly and completely: "https://doi.org/10.5281/zenodo.5746183" and "https://doi.org/10.5281/zenodo.5746655" -- neither is garbled, split, or missing whitespace. metacheck's own <code>zenodo_links()</code> bare-text fallback regex is <code>'(?:https?://)?(?:doi\.(org)?10\.[0-9]+)'</code>, which matches these two strings exactly as they appear in the extracted text (I checked this against the regex by inspection: "https://doi.org/10.5281/zenodo.5746183" satisfies every token of the pattern with no mismatch). Both DOIs resolve live to real, public, open Zenodo records with listed files. Since the platform is supported, the regex should match the citation form actually used, and the repositories are genuinely live and accessible, yet <code>metacheck_any_detector_hit</code> is reported as false, the most likely explanation is a bug somewhere in the detection pipeline between text/URL extraction and the regex match or the live Zenodo API lookup (e.g., the GROBID-extracted paper object's <code>'url'</code> table or body text segmentation for this specific PDF, or a failure in the <code>zenodo_info()</code> API call for these two specific records) rather than a missing platform or a textual extraction gap, since the text I was handed contains the DOIs in good, unambiguous condition.
10.1111/2041-210X.13865	Supported platform, but a specific detection defect	No DAS exists in the extracted text; the only repository mention is a self-citation reference-list entry for Zenodo DOI 10.5281/zenodo.6400498, which appears verbatim and unmangled. Metacheck's <code>zenodo_links()</code> bare-text fallback regex is explicitly designed to catch exactly this citation style. Live-verified the Zenodo API record returns a complete file listing normally. Given the DOI is clean, the regex matches it, the API record is live, and the backend code exists and works for hundreds of other Zenodo papers in this corpus, this is best explained as an execution-time pipeline bug for this specific paper rather than a missing capability or config/URL-pattern gap.

Paper (DOI)	Category	Why Metacheck misses it
10.1111/2041-210X.14104	Supported platform, but a specific detection defect	metacheck has two separate DSpace detectors: archive-psycharchives.R's <code>.dspace_legacy_hosts()</code> (6 hardcoded hosts answering the legacy DSpace <=6 /rest/ API) and archive-dspace7.R's <code>.dspace7_hosts()</code> (64 hardcoded hosts answering the DSpace 7+ /server/api/ HAL-JSON API). Both lists are closed enumerations of specific institutional hostnames, built by manually checking DataCite's client registry; neither "dataworks.iupui.edu" nor its current canonical name "dataworks.indianapolis.iu.edu" is in either list. Live-checking the actual item confirms why it would not fit either detector even if the host were added: this installation runs DSpace 5.2 with the old Cocoon/XMLUI stack, and both /rest/test and /server/api return 404 there, so it exposes neither of the two REST shapes metacheck's two DSpace modules know how to query. <code>dspace7_links()</code> does have a bare-DOI live-resolution fallback (added for issue #435) that resolves any doi.org citation found in a paper and checks whether the redirect target's host is in <code>.dspace7_hosts()</code> -- but since dataworks.indianapolis.iu.edu is absent from that list, this specific DOI would still not match even with that fallback engaged, and even if the host were added to the list, DSpace 5.2's lack of a /server/api endpoint means <code>.dspace7_info()</code> calls would simply fail (404) against it, so the correct fix is adding IUPUI/Indianapolis DataWorks to <code>.dspace_legacy_hosts()</code> (since DSpace 5 uses the legacy /rest/ shape family) only after confirming that host actually answers /rest/test -- which, per the curl check above, it currently does not either, meaning this specific DSpace 5.2 installation may need its own handling rather than simply being added to either existing host list. [RECLASSIFIED from SUPPORTED_PLATFORM_CONFIG_GAP: not a host/DOI-prefix table entry -- requires a new resolution mechanism (handle-prefix redirect, shortlink expansion, or DOI-to-host resolution) or the host runs an API shape metacheck does not support.] <code>dataone_links()</code> correctly detects the bare doi.org/10.5063/f1mc8x92 citation verbatim and correctly routes it to the DataONE/KNB detector (host knb.ecoinformatics.org, doi_prefix 10.5063, both explicitly registered in <code>.dataone_hosts()</code>). However, direct inspection and live invocation of <code>.dataone_pid()</code> shows its regex captures the DOI suffix exactly as it appears in the citation text (case-preserved; ignore.case = TRUE only affects matching 'doi:/'doi.org/' prefixes, not the captured identifier itself) -- producing pid 'doi:10.5063/f1mc8x92' in lowercase. KNB's Metacat v2 API is case-sensitive on this pid: querying the meta endpoint with the lowercase pid returns HTTP 404 live right now, while querying the identical endpoint with the pid case-corrected to doi:10.5063/F1MC8X92 returns a full, valid EML metadata record with 1 file (confirmed by direct invocation of the installed package's <code>.dataone_info()</code> today). This is a genuine, deterministic code defect: <code>.dataone_pid()</code> should normalize/preserve the DOI's registered canonical case (or try both cases) rather than passing through whatever case happens to appear in the citing paper's text. Read the paper's full text myself: Data Availability statement reads 'Repository. https://doi.org/10.5522/04/14484084.v1 (Pritchard et al., 2021)'; also repeated in the reference list -- correctly present in metacheck's own text extraction (confirmed separately: <code>link_in_url_table=TRUE</code>, <code>link_in_extracted_text=TRUE</code>). metacheck's <code>figshare_links()/figshare_id()</code> correctly parse the generic 10.6084/m9.figshare.* DOI shape and any already-registered institutional prefix, but UCL's own institutional prefix (10.5522, used by <code>rdr.ucl.ac.uk</code>) is NOT currently in <code>.figshare_doi_prefix_hosts()</code> in archive-figshare.R. This is a genuine, fixable platform-coverage gap -- not an extraction problem, not a repository-access problem. Confirmed the underlying repository is real, public, and CC0-licensed via a direct live check of the Figshare API. [CONFIG_GAP_TIER: confirmed_simple_fix -- live-verified today, a single host/DOI-prefix table entry resolves this with no further ambiguity.]
10.1111/2041-210X.14387	Supported platform, but a specific detection defect	
10.1002/2688-8319.12070	Supported platform family, but this host/DOI-prefix not registered	

Paper (DOI)	Category	Why Metacheck misses it
10.1002/pan3.10406	Supported platform family, but this host/DOI-prefix not registered	metacheck does support the DSpace platform family in general: it has two separate DSpace detectors, archive-dspace7.R (DSpace 7+'s HAL-JSON API under /server/api/) and the legacy DSpace REST API handling in archive-psycharchives.R (/rest/, DSpace <=6). Both detectors only work against a hardcoded allowlist of specific hostnames (.dspace7_hosts() and .dspace_legacy_hosts() respectively), because there is no general way to recognize an arbitrary domain as a DSpace installation from a URL alone. I checked both lists in the current R source and confirmed "www.repository.cam.ac.uk" (Apollo, University of Cambridge) is present in neither .dspace7_hosts() (archive-dspace7.R, 64 hosts) nor .dspace_legacy_hosts() (archive-psycharchives.R, 6 hosts). I also grepped the whole R/ directory for "cam.ac.uk", "apollo", and "cambridge" and found no matches anywhere, confirming this host has never been added to either detector. The live page itself (https://www.repository.cam.ac.uk/handle/1810/340870) uses the classic DSpace handle-based URL pattern (/handle/1810/340870) and standard DSpace item-page structure, so this is a straightforward DSpace installation that metacheck's existing DSpace code could in principle handle once the host is added. This is a configuration gap in an already-supported platform family, not a missing platform and not an extraction failure: the DOI link is present, verbatim and unambiguous, in the paper's own data availability statement ("All data are available in the Apollo repository held by the University of Cambridge at https://doi.org/10.17863/CAM.88300"), and metacheck_any_detector_hit=false correctly reflects that no current detector's host allowlist covers this repository. Note also that the coder's platform label said "Apollo ... DSpace-based" and the live page confirms this is indeed DSpace, so there is no mislabeling issue here, only a missing host entry. [CONFIG_GAP_TIER: confirmed_simple_fix -- live-verified today, a single host/DOI-prefix table entry resolves this with no further ambiguity.]
10.1002/pan3.10491	Supported platform family, but this host/DOI-prefix not registered	metacheck's Figshare detector (R/archive-figshare.R) already explicitly supports this exact host: .figshare_vanity_hosts() lists "rdr.ucl.ac.uk" by name (comment: "University College London (UK)"), and the detector's host regex and /articles/ path matching would catch any citation that includes that hostname literally. The gap is that this paper's data-availability statement cites a bare DOI, "https://doi.org/10.5522/04/23217572.v1", with no "rdr.ucl.ac.uk" substring anywhere in the text -- the hostname only appears after HTTP redirect resolution, which metacheck's text-based detectors do not perform. Figshare-hosted institutional DOIs are handled instead via .figshare_doi_prefix_hosts(), a DOI-prefix-to-host lookup table used specifically so a bare institutional DOI with no visible hostname can still be matched (this is exactly the mechanism built for e.g. Monash's 10.26180 prefix). However, UCL's own prefix, 10.5522, is not present in that table (the table currently has 10.17633, 10.25383, 10.23641, 10.25573, 10.23640, 10.25452, 10.57831, 10.25444, 10.25422, 10.17045, 10.83399, 10.17608, 10.60696, 10.60580, and 10.26188, but not 10.5522). Because neither the host regex nor the DOI-prefix regex matches, the figshare detector's regex search in .figshare_id()/found_href/other_fs finds nothing for this citation, even though the platform family (Figshare) and even the specific institutional vanity host (rdr.ucl.ac.uk) are both already supported by the module -- this is a one-line missing table entry (adding "10.5522" = "rdr.ucl.ac.uk" to .figshare_doi_prefix_hosts()), not a missing platform or a missing vanity host. [CONFIG_GAP_TIER: confirmed_simple_fix -- live-verified today, a single host/DOI-prefix table entry resolves this with no further ambiguity.]

Paper (DOI)	Category	Why Metacheck misses it
10.1002/pan3.10689	Supported platform family, but this host/DOI-prefix not registered	metacheck has a full, working Figshare backend (R/archive-figshare.R) that correctly resolves articles via <code>api.figshare.com</code>, including dozens of institutional "vanity" Figshare instances and their own DOI prefixes (<code>.figshare_vanity_hosts()</code> and <code>.figshare_doi_prefix_hosts()</code>). But this paper's data link is <code>https://doi.org/10.7294/25914451.v2</code>, hosted at VTechData (<code>data.lib.vt.edu</code>) under DOI prefix 10.7294. Neither the host "<code>data.lib.vt.edu</code>" nor the DOI prefix "10.7294" appears in <code>.figshare_vanity_hosts()</code> or <code>.figshare_doi_prefix_hosts()</code>, and the bare-DOI fallback regex in <code>figshare_links()</code> (<code>figshare_id()</code>) only recognizes the literal 10.6084 prefix plus whatever prefixes are registered in that list -- so a bare citation of "10.7294/25914451.v2" (as this paper's availability statement uses, with no "figshare.com" or "data.lib.vt.edu" string anywhere near the DOI) is invisible to <code>figshare_links()</code>, and the resolved redirect host <code>data.lib.vt.edu</code> is also absent from the host regex, so even a real hyperlink to the landing page would not match either. I independently confirmed live via <code>api.figshare.com/v2/articles/25914451</code> that this is genuinely the same Figshare API backend already supported (<code>is_public: true</code>, matching title, files <code>PAN_MSV_Dataset.xlsx</code>/<code>Survey Instrument.pdf</code>/<code>README.rtf</code>, license CC0 1.0) -- so this is a missing host/DOI-prefix registration for an existing, working platform integration, not a missing platform and not a broken or inaccessible repository. I also note archive-dryad.R's own comments record that DOI prefix 10.7294 (VTechData) was previously checked and confirmed NOT to be Dryad, but it was apparently never cross-registered as Figshare despite that check. Fix: add "10.7294" -> "data.lib.vt.edu" to <code>.figshare_doi_prefix_hosts()</code> and "data.lib.vt.edu" to <code>.figshare_vanity_hosts()</code>. [CONFIG_GAP_TIER: confirmed_simple_fix -- live-verified today, a single host/DOI-prefix table entry resolves this with no further ambiguity.]
10.1002/pan3.14	Supported platform family, but this host/DOI-prefix not registered	metacheck has a working DSpace 7+ detector (R/archive-dspace7.R) that calls the exact HAL-JSON API this installation runs. I confirmed live that GET <code>https://api.repository.cam.ac.uk/server/api</code> returns a HAL root document with "<code>dspaceVersion</code>": "DSpace 8.1" and "<code>dspaceName</code>": "Apollo - University of Cambridge Repository" -- i.e. Apollo is DSpace 8 (a DSpace 7+-family API), squarely within what <code>archive-dspace7.R</code> is built to handle (its own module comment already documents verifying this exact API shape live against another installation, <code>repository.gatech.edu</code>). The failure is purely a missing host registration: <code>.dspace7_hosts()</code> in <code>R/archive-dspace7.R</code> hardcodes 64 specific DSpace 7+ domains (e.g. <code>aura.abdn.ac.uk</code>, <code>repository.gatech.edu</code>, <code>scholarbank.nus.edu.sg</code>) and neither of Apollo's two relevant domains -- the public UI host <code>www.repository.cam.ac.uk</code> (where the DOI 10.17863/CAM.35449 redirects, and which holds the <code>/handle/1810/288430</code> item page <code>.dspace7_parse()</code> would need to recognise) nor its separate REST API host <code>api.repository.cam.ac.uk</code> -- appears in that list. Because <code>.dspace7_links()</code>'s own bare-mention fallback (<code>ds7_bare_regex</code>) is built from the same host list, and the DOI-resolution fallback only keeps a DOI if it redirects to a host already in <code>.dspace7_hosts()</code>, the CAM.35449 DOI is silently dropped at every stage even though it resolves live to a fully supported DSpace-8 API. This is a config gap, not a missing capability or an extraction problem: the link is clearly present and correctly extractable from the paper's data-availability statement, and the detector code itself would succeed against this API if the two Apollo hostnames were simply added to <code>.dspace7_hosts()</code>. [CONFIG_GAP_TIER: confirmed_simple_fix -- live-verified today, a single host/DOI-prefix table entry resolves this with no further ambiguity.]

Paper (DOI)	Category	Why Metacheck misses it
10.1111/1365-2435.13523	Supported platform family, but this host/DOI-prefix not registered	metacheck has a working DSpace 7 detector (R/archive-dspace7.R) and its host allowlist, .dspace7_hosts(), already includes "unsworks.unsw.edu.au" (line 106) -- so if the paper had cited the unsworks.unsw.edu.au URL directly, dspace7_links() would find it. The actual problem is that this paper's data-availability statement cites a different, un-registered domain: "http://handle.unsw.edu.au/1959.4/resource/collection/resdatac_1003/1". UNSW operates handle.unsw.edu.au as a separate CNRI Handle-System redirector in front of UNSWorks; it 302-redirects to the unsworks.unsw.edu.au item page but is a distinct hostname that never appears in .dspace7_hosts(). dspace7_links() checks three paths: (1) real hyperlink hrefs matching a known host -- fails, handle.unsw.edu.au isn't in the list; (2) a bare-text mention of a known host string -- fails, same reason, the text only contains "handle.unsw.edu.au", never "unsworks.unsw.edu.au"; (3) a bare DOI-shaped string ("10.xxxx/yyyy") resolved live and checked against the known-host list -- also fails, because the citation in the paper's text is a Handle-resolver URL with no "10." DOI prefix anywhere in it (the item's real DOI, 10.26190/5e1ce7350483e, only appears on the landing page itself, never in the citing paper's text). So all three of dspace7_links()'s detection paths miss this specific citation even though the underlying repository (unsworks.unsw.edu.au, DSpace 7) is fully supported. This is a configuration gap, not a missing platform or a code bug: the fix would be adding "handle.unsw.edu.au" (or a generic handle.* / hdl.handle.net resolver pattern) to the host recognition logic, analogous to how issue #435 already added live DOI-redirect resolution for the no-host-string-at-all case. [CONFIG_GAP_TIER: confirmed_simple_fix -- live-verified today, a single host/DOI-prefix table entry resolves this with no further ambiguity.]
10.1111/1365-2435.14271	Supported platform family, but this host/DOI-prefix not registered	metacheck has no detector logic anywhere in R/ that matches the host research-data.westernsydney.edu.au or the DOI prefix 10.26183 (confirmed by grep across the whole repository: zero matches for "westernsydney", "research-data.westernsydney", or "26183"). The repository files R/archive-figshare.R, R/archive-dspace7.R, R/archive-dataverse.R etc. implement detection by matching specific known hostnames/URL shapes for each platform family, and none of them recognize this host. I attempted to determine whether research-data.westernsydney.edu.au is itself a thin wrapper around one of metacheck's 14 supported platforms (as the task suggested institutional repositories sometimes are), but the entire domain is behind Akamai edge bot-protection that returned HTTP 403 to every automated request I made (WebFetch and curl alike, with a realistic browser User-Agent), including to /robots.txt and to a Figshare-style /articles/ test path, so I could not read the page's actual markup, meta tags, or API surface to confirm the backend software. The DNS CNAME points to Akamai's edgesuite.net rather than to a known SaaS repository provider's CNAME (e.g. figshare.com), which is at least weak evidence this is a self-hosted/custom Western Sydney University system rather than a rebranded Figshare or DSpace instance, but this is inconclusive given the blocked access. Given that (a) the DOI (10.26183/bf6d-kh20) and its registering prefix are not referenced anywhere in metacheck's codebase, and (b) I could not confirm the underlying platform is secretly one of the 14 supported families due to the access block, I classify this as a platform the detector list simply does not cover for this specific institutional host -- i.e. a configuration/coverage gap for an unsupported institutional repository, rather than a bug in an existing detector or a text-extraction failure (the link is clearly present and correctly extractable in the paper's data-availability statement, exactly where the coder found it). [CONFIG_GAP_TIER: still_uncertain -- host could not be live-verified (blocked, timed out, or inconclusive), or the fix is a code change rather than a table entry.]

Paper (DOI)	Category	Why Metacheck misses it
10.1111/1365-2435.14670	Supported platform family, but this host/DOI-prefix not registered	researchdata.se (the Swedish National Data Service's catalogue) is a Dataverse installation — confirmed live by the <code>api.researchdata.se</code> endpoint pattern serving dataset files, which is Dataverse's standard REST API shape (<code>/api/access/datafile</code> or <code>/dataset/.../file</code> paths), consistent with the file family it belongs to. metacheck's <code>R/archive-dataverse.R</code> already implements full Dataverse support (<code>dataverse_links()</code>, <code>dataverse_info()</code>, file download) via two mechanisms: an explicit host allowlist (<code>dataverse_hosts()</code>) and a verified per-host DOI-prefix allowlist (<code>dataverse_doi_prefix_hosts()</code>), because Dataverse is open-source software run by hundreds of independent, non-uniform installations with no single hostname or DOI prefix. Neither list currently contains this installation: 'researchdata.se' is absent from <code>dataverse_hosts()</code>, and the DOI prefix actually used here, 10.5878, is absent from <code>dataverse_doi_prefix_hosts()</code>. In fact the file's own header comment explicitly documents having considered and excluded this exact host ('snd.gu.se (Sweden; SND appears to have moved to researchdata.se)' — line 39-40 of <code>archive-dataverse.R</code>) during a prior verification pass, apparently because the live-check at that time targeted the old <code>snd.gu.se</code> domain rather than <code>researchdata.se</code> itself (or because <code>researchdata.se</code>'s <code>/api/info/version</code> check failed/was not retried). Because <code>dataverse_links()</code> matches only against these two allowlists (plus a bare-DOI fallback keyed to the same prefix list) and neither one recognizes 'researchdata.se' or prefix '10.5878', the citation 'https://doi.org/10.5878/kr3s-6v85' in this paper's Data Availability Statement would not be matched by metacheck's Dataverse detector, consistent with <code>metacheck_any_detector_hit</code> being reported false for this paper. This is a configuration gap in an already-supported platform family, not a missing platform family and not a text-extraction problem — the link is present verbatim and unambiguous in the full text. [CONFIG_GAP_TIER: confirmed_simple_fix -- live-verified today, a single host/DOI-prefix table entry resolves this with no further ambiguity.]
10.1111/1365-2656.13268	Supported platform family, but this host/DOI-prefix not registered	digital.CSIC is a DSpace-based institutional repository (confirmed live: the DOI <code>http://dx.doi.org/10.20350/digitalCSIC/12516</code> redirects via HTTP 302 to <code>https://digital.csic.es/handle/10261/211267</code>, and the <code>/handle/{prefix}/{suffix}</code> path is DSpace's canonical permanent-URL scheme), and <code>dspace/dspace7</code> is on metacheck's list of supported platforms. Yet <code>metacheck_any_detector_hit</code> is false. The paper cites the resource only via its DataCite DOI, "10.20350/digitalCSIC/12516" (in both the data-availability sentence and the reference list), never writing out the <code>digital.csic.es</code> hostname or a <code>/handle/</code> URL anywhere in the text. If metacheck's <code>DSpace/dspace7</code> detector pattern-matches on known DSpace hostnames or URL shapes (e.g. a string match for "digital.csic.es" or <code>"/handle/"</code>) rather than resolving arbitrary DOIs and checking the resulting redirect target, it would correctly recognize a <code>digital.csic.es</code> URL if one were written in the paper, but would miss this exact citation because the paper never exposes that host string -- only the DOI prefix 10.20350, which the detector's host/pattern list evidently does not include. This is a configuration gap for a specific DOI-prefix-to-host mapping within an otherwise-supported platform family, not a missing detector for DSpace generally. [CONFIG_GAP_TIER: still_uncertain -- host could not be live-verified (blocked, timed out, or inconclusive), or the fix is a code change rather than a table entry.]

Paper (DOI)	Category	Why Metacheck misses it
10.1111/1365-2656.13439	Supported platform family, but this host/DOI-prefix not registered	metacheck's R/archive-figshare.R implements full, working Figshare v2 API support (host regex matching, DOI-prefix table, article retrieval, file download), and I confirmed live that VTechData (data.lib.vt.edu) is genuinely built on the shared Figshare platform: the DOI https://doi.org/10.7294/WSDX-AJ44 resolves to a data.lib.vt.edu/articles/dataset/.../14096975 URL (the classic Figshare "/articles/<type>/<slug>/<id>" path shape), and the article's full metadata (title, license, 8 files, DOI) is served live and unauthenticated from the standard api.figshare.com/v2/articles/14096975 endpoint — exactly the behavior every other Figshare vanity host exhibits. However, neither the host "data.lib.vt.edu" nor the DOI prefix "10.7294" is registered in .figshare_vanity_hosts() or .figshare_doi_prefix_hosts() in R/archive-figshare.R. Consequently figshare_links()'s host_regex/doi_prefix_regex (built from .figshare_host_regex() and .figshare_id()'s pattern list never match this citation, so metacheck would not recognize "https://doi.org/10.7294/WSDX-AJ44" as a Figshare link at all, consistent with metacheck_any_detector_hit being false. Notably, R/archive-dryad.R line 27-28 already contains a comment stating "10.7294 ... confirmed NOT Dryad (VTechData...)" — i.e. a prior investigation had already identified that this exact DOI prefix belongs to VTechData/Figshare, but the corresponding entry was never added to archive-figshare.R's vanity-host or DOI-prefix lists to close the gap. This is a config gap (missing list entries for an already-confirmed Figshare-family host), not a missing detector or a broken one. [CONFIG_GAP_TIER: confirmed_simple_fix -- live-verified today, a single host/DOI-prefix table entry resolves this with no further ambiguity.]
10.1111/1365-2656.13478	Supported platform family, but this host/DOI-prefix not registered	The repository is a genuine legacy-DSpace installation — the URL path "/jspui/" is DSpace's Java Server Pages UI, and the item answers with a standard DSpace handle (1007/3374) and bitstream structure, which I confirmed live (public item, "Acceso Abierto," Attribution license, two CSV bitstreams). metacheck does have a DSpace-legacy backend (R/archive-psycharchives.R, built around the DSpace <=6 /rest/ API) that is platform-capable of handling exactly this kind of installation. However, its host-matching is not pattern-based; it is a hardcoded allowlist in .dspace_legacy_hosts(): www.psycharchives.org, bonndoc.ulb.uni-bonn.de, dspace.ut.ee, qsar.db.org, repositorio-digital.cide.edu, and ses.library.usyd.edu.au. There is a separate, equally hardcoded allowlist for DSpace 7+ installations in .dspace7_hosts() (R/archive-dspace7.R). "cicese.repositorioinstitucional.mx" appears in neither list, so .dspace_legacy_host_regex() and .dspace7_host_regex() never match this host, and dspace_links()/dspace7_links() therefore never pick up the citation even though the underlying detector code is fully able to query this exact DSpace REST shape once the host is registered. This is a configuration gap (an unlisted host for an otherwise-supported platform family), not a missing detector or an extraction failure — the link is present verbatim in the paper text and matches the coder's recorded URL exactly. [CONFIG_GAP_TIER: still_uncertain -- host could not be live-verified (blocked, timed out, or inconclusive), or the fix is a code change rather than a table entry.]

Paper (DOI)	Category	Why Metacheck misses it
10.1111/1365-2656.13643	Supported platform family, but this host/DOI-prefix not registered	The data statement ("Data available from UQ eSpace https://doi.org/10.48610/088a9c5") is clearly present and extractable, and metacheck does have a general-purpose mechanism designed for exactly this situation: R/archive-dspace7.R's <code>dspace7_links()</code> live-resolves every bare DOI found in a paper's text and checks whether it redirects to a host in <code>.dspace7_hosts()</code> (added specifically per issue #435 to catch papers that cite only a DOI with no recognizable host string in the visible text, as is the case here). I resolved the DOI live myself and it redirects to <code>espace.library.uq.edu.au</code>, which is not in <code>.dspace7_hosts()</code> (64 hosts) nor in <code>.dspace_legacy_hosts()</code> (the 6-host legacy-DSpace list in <code>archive-psycharchives.R</code>, which covers the <code>/rest/</code> API shape). I also confirmed that the landing page is a custom React single-page application whose server-rendered HTML carries no API fingerprint I could detect (every probed path, including <code>/server/api</code> and <code>/rest/test</code>, returned 403), so it does not even clearly expose a DSpace7 HAL-API or legacy DSpace REST API the way the other 64+6 verified hosts do -- it appears to be a bespoke/Fez-descended frontend rather than a stock DSpace install, meaning even adding this host to the existing list might not work without separately verifying what, if any, machine-readable API it answers. This is therefore a configuration/host-registry gap in an already-supported platform family (DSpace7-style resolution plumbing exists and executes against every DOI in the paper) rather than a missing platform category or an extraction failure. [CONFIG_GAP_TIER: still_uncertain -- host could not be live-verified (blocked, timed out, or inconclusive), or the fix is a code change rather than a table entry.]
10.1111/1365-2656.13671	Supported platform family, but this host/DOI-prefix not registered	metacheck has two separate, working DSpace detector implementations: <code>archive-psycharchives.R</code> for legacy DSpace (the <code>/rest/</code> API) and <code>archive-dspace7.R</code> for DSpace 7+ (the <code>/server/api/</code> HAL-JSON API). Both work by matching the citing URL's host against a hardcoded allowlist: <code>.dspace_legacy_hosts()</code> (6 installations: <code>psycharchives.org</code>, <code>bonndoc.ulb.uni-bonn.de</code>, <code>dspace.ut.ee</code>, <code>qsardb.org</code>, <code>repositorio-digital.cide.edu</code>, <code>ses.library.usyd.edu.au</code>) and <code>.dspace7_hosts()</code> (64 installations, e.g. <code>repository.gatech.edu</code>, <code>conservancy.umn.edu</code>). I read both lists directly from <code>R/archive-psycharchives.R</code> and <code>R/archive-dspace7.R</code> and confirmed "aurora.auburn.edu" appears in neither. The paper's data-availability DOI (https://doi.org/10.35099/aurora-91) correctly 302-redirects to https://aurora.auburn.edu/handle/11200/50021 — a <code>/handle/<prefix>:<suffix></code> URL shape that is the classic DSpace permalink pattern (the same shape <code>.dspace7_parse()</code>'s <code>handle_pat</code> regex already looks for). <code>dspace7_links()</code> does have a fallback mechanism added for issue #435 that live-resolves any bare DOI found in text and checks whether the resolved URL lands on a KNOWN <code>dspace7</code> host — but since <code>aurora.auburn.edu</code> is not in that host list, this fallback would also fail to recognize it; the resolved URL is simply never compared against anything matching. This is therefore a configuration gap, not a missing backend: metacheck's DSpace 7 (or possibly legacy DSpace, pending live confirmation of which API version AUrora runs) code would almost certainly detect and extract <code>files/license/DOI</code> correctly if "aurora.auburn.edu" were simply added to one of the two host lists, exactly as the file's own header comments describe doing for new hosts found via DataCite's client registry. I was unable to independently confirm live which DSpace version AUrora runs (connection to <code>aurora.auburn.edu</code> timed out from this environment on repeated <code>WebFetch</code> and <code>curl</code> attempts, while the DOI redirect itself resolved fine), so I cannot say with certainty whether the fix belongs in <code>.dspace7_hosts()</code> or <code>.dspace_legacy_hosts()</code>, but it unambiguously belongs in one of metacheck's existing DSpace host registries rather than requiring new platform code. [CONFIG_GAP_TIER: still_uncertain -- host could not be live-verified (blocked, timed out, or inconclusive), or the fix is a code change rather than a table entry.]

Paper (DOI)	Category	Why Metacheck misses it
10.1111/1365-2656.13970	Supported platform family, but this host/DOI-prefix not registered	Apollo (www.repository.cam.ac.uk) is a genuine DSpace installation: a direct request to https://www.repository.cam.ac.uk/server/api returns HTML whose head contains <code><meta name="Generator" content="DSpace 8.1"></code> and <code><title>DSpace</title></code>, i.e. the modern DSpace Angular frontend running the same <code>/server/api/</code> HAL-JSON REST architecture that metacheck's <code>R/archive-dspace7.R</code> is specifically built to query (DSpace 8 is the direct successor of DSpace 7 and uses the identical API shape). However, metacheck's <code>dspace7_links()/dspace7_info()</code> only recognize a fixed allowlist of hosts returned by <code>.dspace7_hosts()</code> in <code>archive-dspace7.R</code>, and I confirmed by reading that file and by grepping the whole <code>R/</code> directory for "cam.ac.uk" (no matches) that "www.repository.cam.ac.uk" is not among the 64 listed hosts. Because <code>dspace7_links()</code> matches on this host allowlist (plus a same-paper DOI-resolution fallback that only keeps DOIs whose live redirect lands on an already-listed host), a citation of https://doi.org/10.17863/CAM.97060 is invisible to the detector even though the DOI resolves live to a page the detector's own REST logic could fully parse once pointed at the right host. This is a configuration gap (an unlisted host for an already-supported platform family), not a missing detector, and not an extraction problem -- the link is present, correctly formed, and findable in the text exactly where the coder recorded it. [CONFIG_GAP_TIER: confirmed_simple_fix -- live-verified today, a single host/DOI-prefix table entry resolves this with no further ambiguity.] metacheck has full DSpace support in two forms: <code>archive-dspace7.R</code> (<code>.dspace7_hosts()</code>, 64 confirmed DSpace-7+ installations, REST API at <code>/server/api/</code>) and <code>archive-psycharchives.R</code> (<code>.dspace_legacy_hosts()</code>, 6 confirmed legacy DSpace <=6 installations, REST API at <code>/rest/</code>). The paper's DOI (10.20350/digitalCSIC/15508) redirects to https://digital.csic.es/handle/10261/334537 -- the path shape <code>"/handle/<prefix>/<suffix>"</code> is the classic DSpace handle-resolution URL, which independently corroborates the human coder's platform label (DIGITAL.CSIC, DSpace-based, run by CSIC). However, the host "digital.csic.es" does not appear in either <code>.dspace7_hosts()</code> or <code>.dspace_legacy_hosts()</code>. Both registries were built by manually verifying individual hosts found via DataCite's public client registry filtered by <code>software=dspace</code>; <code>digital.csic.es</code> was apparently never checked and added to either list. Because <code>dspace7_links()/dspace_links()</code> match only against these explicit host allowlists (via <code>.dspace7_host_regex()/dspace_legacy_host_regex()</code>), and because this paper's text never gives a bare DOI that could trigger the <code>dspace7_links()</code> DOI-resolution fallback path for an unlisted host (that fallback only keeps DOIs that resolve to an ALREADY-known host in <code>.dspace7_hosts()</code>), so an unlisted host like <code>digital.csic.es</code> would be filtered out by <code>.dspace7_resolve_doi_url()</code> even if reached), metacheck's detector will not recognize this citation as a repository link at all, regardless of whether <code>digital.csic.es</code> turns out to run legacy DSpace or DSpace 7+. This is a configuration gap in an otherwise-supported platform family, not a missing platform or an extraction failure -- the fix would be adding "digital.csic.es" (after confirming live whether it answers <code>/rest/test</code> or <code>/server/api</code>) to the appropriate hosts list. [CONFIG_GAP_TIER: still_uncertain -- host could not be live-verified (blocked, timed out, or inconclusive), or the fix is a code change rather than a table entry.]
10.1111/1365-2656.14007	Supported platform family, but this host/DOI-prefix not registered	Because <code>dspace7_links()/dspace_links()</code> match only against these explicit host allowlists (via <code>.dspace7_host_regex()/dspace_legacy_host_regex()</code>), and because this paper's text never gives a bare DOI that could trigger the <code>dspace7_links()</code> DOI-resolution fallback path for an unlisted host (that fallback only keeps DOIs that resolve to an ALREADY-known host in <code>.dspace7_hosts()</code>), so an unlisted host like <code>digital.csic.es</code> would be filtered out by <code>.dspace7_resolve_doi_url()</code> even if reached), metacheck's detector will not recognize this citation as a repository link at all, regardless of whether <code>digital.csic.es</code> turns out to run legacy DSpace or DSpace 7+. This is a configuration gap in an otherwise-supported platform family, not a missing platform or an extraction failure -- the fix would be adding "digital.csic.es" (after confirming live whether it answers <code>/rest/test</code> or <code>/server/api</code>) to the appropriate hosts list. [CONFIG_GAP_TIER: still_uncertain -- host could not be live-verified (blocked, timed out, or inconclusive), or the fix is a code change rather than a table entry.]

Paper (DOI)	Category	Why Metacheck misses it
10.1111/1365-2664.13035	Supported platform family, but this host/DOI-prefix not registered	CONFIRMED CONFIG GAP, still present. The citation ('Data in this study are available from DataSTORRE... http://hdl.handle.net/11667/102') uses a bare Handle System URL, not a DOI, so it can only be recognised via a host match. datastorre.stir.ac.uk is absent from BOTH <code>.dspace7_hosts()</code> (R/archive-dspace7.R, 63 verified installations) and <code>.dspace_legacy_hosts()</code> (R/archive-psycharchives.R, 6 installations) -- confirmed by direct search, zero matches for 'datastorre' or 'stir.ac.uk' in either file. <code>dspace7_links()</code>'s DOI-resolution fallback (which CAN catch a bare DOI whose host isn't separately listed, by live-resolving it and checking where it redirects) only matches DOI-shaped strings ('10.xxxx/...'); a hdl.handle.net URL is not DOI-shaped, so even that fallback never fires. The repository platform (DSpace) is generally supported, but this specific host was simply never added to either allowlist -- a straightforward config/allowlist omission, not a logic defect. [CONFIG_GAP_TIER: confirmed_simple_fix -- live-verified today, a single host/DOI-prefix table entry resolves this with no further ambiguity.] metacheck's archive-figshare.R fully supports Figshare, including an extensive, actively-maintained list of 37 institutional "vanity" Figshare hosts in <code>.figshare_vanity_hosts()</code> (e.g. aura.american.edu, dro.deakin.edu.au, rdr.ucl.ac.uk) plus a parallel <code>.figshare_doi_prefix_hosts()</code> mapping institutional DOI prefixes (10.17633, 10.26180, etc.) to their hosts, specifically to catch bare citations and vanity-domain links that are not under figshare.com itself. The paper's repository, USDA's Ag Data Commons, is confirmed live to be exactly this kind of vanity Figshare instance: its DOI (prefix 10.15482) redirects through data.nal.usda.gov to a final landing page at <a href="http://agdatacommons.nal.usda.gov/articles/<title>/<numeric-id>">agdatacommons.nal.usda.gov/articles/<title>/<numeric-id> — the identical <a href="http://articles/<slug>/<id>">articles/<slug>/<id> path shape documented throughout archive-figshare.R as diagnostic of a genuine Figshare-backed institution (matching the coder's own platform label, 'Ag Data Commons ... built on the Figshare platform'). However, neither 'agdatacommons.nal.usda.gov' nor 'data.nal.usda.gov' appears in <code>.figshare_vanity_hosts()</code>, and the DOI prefix '10.15482' does not appear in <code>.figshare_doi_prefix_hosts()</code>. Consequently <code>figshare_links()</code>'s <code>host_regex</code> (built from <code>.figshare_host_regex()</code>, which is derived entirely from these two lists plus the literal figshare.com/figsh.com domains) does not match this paper's URL at all, and the DOI-prefix fallback used for bare mentions (which also keys off <code>names(.figshare_doi_prefix_hosts())</code>) likewise never fires for 10.15482. The link is present verbatim in the extracted text and cleanly extractable; the architecture to resolve it (<code>figshare_info()</code>, <code>figshare_id()</code>, the files/license API fields) is fully built and already handles this exact URL shape for dozens of other institutions. This is purely a missing registry entry — adding 'agdatacommons.nal.usda.gov' to <code>.figshare_vanity_hosts()</code> (and/or '10.15482' to <code>.figshare_doi_prefix_hosts()</code>) would make this link detectable with no other code change. [CONFIG_GAP_TIER: confirmed_simple_fix -- live-verified today, a single host/DOI-prefix table entry resolves this with no further ambiguity.]
10.1111/1365-2664.13065	Supported platform family, but this host/DOI-prefix not registered	CONFIRMED CONFIG GAP, still present. The citation ('Data in this study are available from DataSTORRE... http://hdl.handle.net/11667/102') uses a bare Handle System URL, not a DOI, so it can only be recognised via a host match. datastorre.stir.ac.uk is absent from BOTH <code>.dspace7_hosts()</code> (R/archive-dspace7.R, 63 verified installations) and <code>.dspace_legacy_hosts()</code> (R/archive-psycharchives.R, 6 installations) -- confirmed by direct search, zero matches for 'datastorre' or 'stir.ac.uk' in either file. <code>dspace7_links()</code>'s DOI-resolution fallback (which CAN catch a bare DOI whose host isn't separately listed, by live-resolving it and checking where it redirects) only matches DOI-shaped strings ('10.xxxx/...'); a hdl.handle.net URL is not DOI-shaped, so even that fallback never fires. The repository platform (DSpace) is generally supported, but this specific host was simply never added to either allowlist -- a straightforward config/allowlist omission, not a logic defect. [CONFIG_GAP_TIER: confirmed_simple_fix -- live-verified today, a single host/DOI-prefix table entry resolves this with no further ambiguity.] metacheck's archive-figshare.R fully supports Figshare, including an extensive, actively-maintained list of 37 institutional "vanity" Figshare hosts in <code>.figshare_vanity_hosts()</code> (e.g. aura.american.edu, dro.deakin.edu.au, rdr.ucl.ac.uk) plus a parallel <code>.figshare_doi_prefix_hosts()</code> mapping institutional DOI prefixes (10.17633, 10.26180, etc.) to their hosts, specifically to catch bare citations and vanity-domain links that are not under figshare.com itself. The paper's repository, USDA's Ag Data Commons, is confirmed live to be exactly this kind of vanity Figshare instance: its DOI (prefix 10.15482) redirects through data.nal.usda.gov to a final landing page at <a href="http://agdatacommons.nal.usda.gov/articles/<title>/<numeric-id>">agdatacommons.nal.usda.gov/articles/<title>/<numeric-id> — the identical <a href="http://articles/<slug>/<id>">articles/<slug>/<id> path shape documented throughout archive-figshare.R as diagnostic of a genuine Figshare-backed institution (matching the coder's own platform label, 'Ag Data Commons ... built on the Figshare platform'). However, neither 'agdatacommons.nal.usda.gov' nor 'data.nal.usda.gov' appears in <code>.figshare_vanity_hosts()</code>, and the DOI prefix '10.15482' does not appear in <code>.figshare_doi_prefix_hosts()</code>. Consequently <code>figshare_links()</code>'s <code>host_regex</code> (built from <code>.figshare_host_regex()</code>, which is derived entirely from these two lists plus the literal figshare.com/figsh.com domains) does not match this paper's URL at all, and the DOI-prefix fallback used for bare mentions (which also keys off <code>names(.figshare_doi_prefix_hosts())</code>) likewise never fires for 10.15482. The link is present verbatim in the extracted text and cleanly extractable; the architecture to resolve it (<code>figshare_info()</code>, <code>figshare_id()</code>, the files/license API fields) is fully built and already handles this exact URL shape for dozens of other institutions. This is purely a missing registry entry — adding 'agdatacommons.nal.usda.gov' to <code>.figshare_vanity_hosts()</code> (and/or '10.15482' to <code>.figshare_doi_prefix_hosts()</code>) would make this link detectable with no other code change. [CONFIG_GAP_TIER: confirmed_simple_fix -- live-verified today, a single host/DOI-prefix table entry resolves this with no further ambiguity.]

Paper (DOI)	Category	Why Metacheck misses it
10.1111/1365-2664.13397	Supported platform family, but this host/DOI-prefix not registered	metacheck fully supports Figshare (R/archive-figshare.R), including a mechanism built specifically for institutional Figshare instances cited only as a bare DOI with no "figshare.com" substring anywhere in the text: .figshare_doi_prefix_hosts() maps DOI prefixes (e.g. 10.26180 for Monash, 10.17633 for Brunel) to their figshare.com vanity subdomain, so figshare_links()/figshare_id() can resolve such bare DOIs. This paper's citation is exactly that shape - the only link anywhere in the text is the bare DOI https://doi.org/10.25909/5b581a5a151da, with no "figshare" string anywhere nearby. I confirmed live that this DOI redirects to adelaide.figshare.com (University of Adelaide's institutional Figshare instance) and that DataCite's own registry lists prefix 10.25909 under client "adelaide.repo" / provider "adelaide". I then grepped R/archive-figshare.R directly: neither "adelaide" nor "25909" appears anywhere in the file, confirming adelaide.figshare.com is absent from .figshare_vanity_hosts() and 10.25909 is absent from .figshare_doi_prefix_hosts(). I verified the practical consequence by running the actual function in R: .figshare_id("https://doi.org/10.25909/5b581a5a151da") returns NA, while .figshare_id() on the equivalent full article URL (adelaide.figshare.com/articles/dataset/.../6859511/1) correctly returns "6859511" - the exact same code path works once the host is known, it is simply not registered for this institution's DOI prefix. This is the identical gap category the file's own comments describe closing for roughly 35 other institutions already listed; University of Adelaide (adelaide.figshare.com, prefix 10.25909) is just not yet one of the entries. Adding "10.25909" = "adelaide.figshare.com" to .figshare_doi_prefix_hosts() (and the host to .figshare_vanity_hosts() for completeness) would make this citation detectable with no architectural change - a configuration/data gap, not a code bug or missing platform support. [CONFIG_GAP_TIER: confirmed_simple_fix -- live-verified today, a single host/DOI-prefix table entry resolves this with no further ambiguity.] DAS: 'Data available via Apollo https://doi.org/10.17863/CAM.50140.' This resolves to Apollo (University of Cambridge, DSpace 8.1). Metacheck's archive-dspace7.R already has a working DSpace-7 backend, but its .dspace7_hosts()/dspace_legacy_hosts() allowlists do not include repository.cam.ac.uk/api.repository.cam.ac.uk. Separately, Apollo's REST/HAL API is served from a different subdomain (api.repository.cam.ac.uk) than the citable UI host (www.repository.cam.ac.uk), which Metacheck's parser assumes are the same host -- this would need explicit handling even after adding the host. [CONFIG_GAP_TIER: confirmed_simple_fix -- live-verified today, a single host/DOI-prefix table entry resolves this with no further ambiguity.]
10.1111/1365-2664.13614	Supported platform family, but this host/DOI-prefix not registered	

Paper (DOI)	Category	Why Metacheck misses it
10.1111/1365-2664.14198	Supported platform family, but this host/DOI-prefix not registered	metacheck has real, working DSpace support split across two code paths: R/archive-dspace7.R (DSpace 7+, HAL-JSON /server/api/) and R/archive-psycharchives.R's .dspace_legacy_hosts() (DSpace <=6, legacy /rest/ API). Both paths work only against an explicit allowlist of installation hostnames rather than detecting DSpace generically from the resolved URL. I read both host lists directly from source: .dspace7_hosts() in archive-dspace7.R (64 hosts, e.g. repository.gatech.edu, aura.abdn.ac.uk) and .dspace_legacy_hosts() in archive-psycharchives.R (6 hosts: www.psycharchives.org, bonndoc.ulb.uni-bonn.de, dspace.ut.ee, qsar.db.org, repositorio-digital.cide.edu, ses.library.usyd.edu.au). The paper's repository link resolves to datastorre.stir.ac.uk (University of Stirling's dataSTORRE, confirmed live to be running DSpace), which appears in neither list. Because both dspace7_links() and dspace_legacy_links() match only against their respective host allowlists (via .dspace7_host_regex()/dspace_legacy_host_regex()), a citation to http://hdl.handle.net/11667/195 -- which does carry a resolvable handle and does redirect to a real, public, CC-BY-licensed DSpace item with a listed data file -- is invisible to metacheck simply because datastorre.stir.ac.uk was never added to either registry. This is exactly the SUPPORTED_PLATFORM_CONFIG_GAP case: the platform family (DSpace) is supported, extraction of the citation is trivial (a single unambiguous handle link appears twice in the text), but this specific host is unregistered. [CONFIG_GAP_TIER: confirmed_simple_fix -- live-verified today, a single host/DOI-prefix table entry resolves this with no further ambiguity.]
10.1111/1365-2664.14415	Supported platform family, but this host/DOI-prefix not registered	metacheck has two DSpace detectors: archive-psycharchives.R's legacy .dspace_legacy_hosts() (REST API under /rest/) and archive-dspace7.R's .dspace7_hosts() (HAL-JSON API under /server/api/). I checked both host allow-lists directly in the source, and ecommons.cornell.edu is registered in NEITHER list (I grepped archive-psycharchives.R for \ "ecommons\ " ^ "cornell\ " with no match, and read the full 64-entry .dspace7_hosts() vector in archive-dspace7.R -- Cornell's eCommons is absent from both). Because the DOI 10.7298/3dbp-t721 does not resolve via a doi.org URL pattern that is pre-registered the way dataverse/dryad have DOI-prefix host maps, and because dspace7_links()/dspace_legacy_links() only match a literal host-string from their fixed lists (or, per the dspace7_links() code comment on issue #435, resolve a bare DOI live and check if it lands on an "already-known" host) -- a DOI resolving to a DSpace host that simply isn't in the list yields no match at all, live-resolution included. This is exactly the \ "SUPPORTED_PLATFORM_CONFIG_GAP\ " case: metacheck's dspace7/legacy-dspace detector architecture and API-calling code are fully built and would work immediately against ecommons.cornell.edu (I confirmed the live page is plain DSpace with /handle/ and bitstream URLs, consistent with the other registered DSpace installations), but the specific host was never added to either hardcoded host list, so metacheck_any_detector_hit=false is explained purely by a missing registry entry, not a missing capability or a text-extraction failure (the link itself is present and correctly formed in the full text). [CONFIG_GAP_TIER: confirmed_simple_fix -- live-verified today, a single host/DOI-prefix table entry resolves this with no further ambiguity.]

Paper (DOI)	Category	Why Metacheck misses it
10.1111/1365-2664.14787	Supported platform family, but this host/DOI-prefix not registered	The link in the paper is a Handle redirect (http://hdl.handle.net/10261/367361) that resolves to digital.csic.es/handle/10261/367361. DIGITAL.CSIC is CSIC's institutional repository and runs on DSpace, which uses the canonical DSpace URL shape <code>"handle/{prefix}/{suffix}"</code> (this matches the URL exactly, unlike Figshare's <code>"articles/..."</code>, Dataverse's <code>"dataset.xhtml?..."</code>, or EPrints' <code>"/{numeric-id}/"</code>). metacheck does list <code>dspace/dspace7</code> among its supported platforms, so the general backend capability exists. The likely gap is not an absent detector family but a configuration/registration gap: metacheck's dspace detector very likely needs to recognize either the raw <code>hdl.handle.net</code> redirect pattern and/or the specific <code>digital.csic.es</code> host as a known DSpace installation (its DOI prefix 10261 or its domain may simply not be in whatever host/prefix list the dspace detector module consults), since <code>metacheck_any_detector_hit</code> was reported false for this paper even though the citation text itself ("Data available from the CSIC repository: http://hdl.handle.net/10261/367361") is clearly present and extractable, not garbled or missing. I was unable to live-verify the landing page's accessibility, license, or file listing today because every request to <code>digital.csic.es</code> was intercepted by an Anubis bot-protection challenge page rather than reaching the real DSpace item page; this is independent of, and does not change, the detection root-cause analysis of why metacheck's own detector did not fire. [CONFIG_GAP_TIER: still_uncertain -- host could not be live-verified (blocked, timed out, or inconclusive), or the fix is a code change rather than a table entry.]
10.1111/1365-2745.12985	Supported platform family, but this host/DOI-prefix not registered	metacheck does have a DataONE detector (<code>R/archive-dataone.R</code>), and EDI (Environmental Data Initiative) is architecturally a real DataONE member node running the Metacat/PASTA software family that <code>archive-dataone.R</code>'s own file header describes (same "each host answers at a different base path" federation model as Arctic Data Center and KNB). However, <code>.dataone_hosts()</code> in <code>archive-dataone.R</code> hardcodes only five verified member-node entries: <code>arcticdata.io</code> (DOI prefix 10.18739), <code>knb.ecoinformatics.org</code> (10.5063), <code>metacat.tfri.gov.tw</code>, <code>smithsonian.dataone.org</code>, and <code>data.piscoweb.org</code> (10.6085). EDI's own host, <code>portal.edirepository.org</code> (with its PASTA API host <code>pasta.lternet.edu</code>), is not among them, and the paper's DOI prefix 10.6073 (the <code>"10.6073/pasta/..."</code> PASTA-issued DOI used in this citation) is not registered in <code>doi_prefix</code> for any listed host either. Consequently <code>dataone_links()</code>'s <code>host_regex</code> and <code>doi_regex</code> both fail to match this citation: neither the href-based match nor the bare-DOI-in-text fallback can fire, so metacheck's detector produces no hit for this real, live, properly-cited DataONE-family dataset. This is a configuration gap (an unlisted host/DOI-prefix within an already-supported platform family), not a missing detector, and not a text-extraction problem — the citation text itself is complete, unambiguous, and appears twice in the paper (data-availability statement and reference list), exactly matching what the human coder recorded. [CONFIG_GAP_TIER: confirmed_simple_fix -- live-verified today, a single host/DOI-prefix table entry resolves this with no further ambiguity.]

Paper (DOI)	Category	Why Metacheck misses it
10.1111/1365-2745.13306	Supported platform family, but this host/DOI-prefix not registered	metacheck has a DataONE detector (R/archive-dataone.R, function <code>dataone_links()</code>), and EDI/PASTA is architecturally a DataONE member node, so this is the closest-matching platform family rather than a wholly unsupported one. But <code>dataone_links()</code> only matches a fixed allowlist of hosts and DOI prefixes returned by <code>.dataone_hosts()</code>: <code>arcticdata.io</code> (10.18739), <code>knb.ecoinformatics.org</code> (10.5063), <code>metacat.tfri.gov.tw</code>, <code>smithsonian.dataone.org</code>, and <code>data.piscoweb.org</code> (10.6085). Neither the host in this citation's redirect target (<code>portal.edirepository.org</code>, or the PASTA API host <code>pasta.lternet.org/pasta.edirepository.org</code>) nor its DOI prefix (10.6073, EDI's registered DataCite prefix) appears in that list. Consequently the detector's <code>host_regex</code> and <code>doi_regex</code> (built in <code>dataone_links()</code> from <code>.dataone_hosts()</code>) will not match either the literal string <code>"https://doi.org/10.6073/pasta/c2ba534cb03a7d1ef203a3af63bfe6bc"</code> in the paper's text or any resolved EDI/<code>portal.edirepository.org</code> URL, so <code>metacheck_any_detector_hit</code> being false is explained precisely by this missing registry entry, not by any absence of DataONE support in general or by the link being missing/garbled in the text (the DOI is present, correctly formatted, and verbatim-matches the coder's recorded link). Adding EDI's host (<code>portal.edirepository.org</code>, and ideally the PASTA API host for metadata retrieval) plus the 10.6073 DOI prefix to <code>.dataone_hosts()</code> would close this specific gap without requiring a new platform module. [CONFIG_GAP_TIER: confirmed_simple_fix -- live-verified today, a single host/DOI-prefix table entry resolves this with no further ambiguity.]
10.1111/1365-2745.13864	Supported platform family, but this host/DOI-prefix not registered	metacheck does have a DataONE detector (R/archive-dataone.R), and DataONE is exactly the right platform family for this citation: <code>www.lter.uaf.edu</code>'s data-detail pages for ids 704/705/706 are LTER PASTA data packages with Network Data IDs (<code>knbn-lter-bnz.704.4</code> etc.) and DOIs under the <code>10.6073/pasta/</code> prefix -- both are DataONE-standard identifier shapes, and PASTA is a registered DataONE member node. However, <code>.dataone_hosts()</code> in <code>archive-dataone.R</code> hardcodes only 5 specific, individually live-verified member-node hosts (<code>arcticdata.io</code>, <code>knb.ecoinformatics.org</code>, <code>metacat.tfri.gov.tw</code>, <code>smithsonian.dataone.org</code>, <code>data.piscoweb.org</code>) with their own <code>doi_prefix</code> values (10.18739, 10.5063, NA, NA, 10.6085). Neither the host <code>www.lter.uaf.edu</code> nor the DOI prefix 10.6073 appears anywhere in that list or anywhere else in the codebase (confirmed by grep). Consequently <code>dataone_links()</code>'s <code>host_regex/doi_regex</code> match on this citation fails entirely -- the citation is neither a recognized host substring nor a recognized DOI prefix -- so <code>.dataone_host()</code> returns NA and the dataset is silently dropped before any API call is attempted. This is precisely the scenario the file's own header comment anticipates: "a host neither listed by domain nor identifiable by its DOI prefix is not one this package can currently query, even if it is a real DataONE member node." The fix is a configuration addition (one more verified host/<code>doi_prefix</code> entry, e.g. <code>host = "lter.uaf.edu"</code> or more generally a PASTA/LTER-wide <code>doi_prefix = "10.6073"</code>), not a new detector module or a bug fix to existing logic. [CONFIG_GAP_TIER: confirmed_simple_fix -- live-verified today, a single host/DOI-prefix table entry resolves this with no further ambiguity.]

Paper (DOI)	Category	Why Metacheck misses it
10.1111/1365-2745.14006	Supported platform family, but this host/DOI-prefix not registered	metacheck's <code>dataone_links()</code> (R/archive-dataone.R) implements DataONE support as a hard-coded allow-list of specific member-node hosts and their DOI prefixes via <code>.dataone_hosts()</code>: <code>arcticdata.io</code> (10.18739), <code>knb.ecoinformatics.org</code> (10.5063), <code>metacat.tfri.gov.tw</code> (no prefix), <code>smithsonian.dataone.org</code> (no prefix), and <code>data.piscoweb.org</code> (10.6085). The citation in this paper is to the Environmental Data Initiative (EDI) repository, whose DOI prefix is 10.6073 and whose domains are <code>portal.edirepository.org</code> / <code>pasta.lternet.edu</code> -- none of which appear in <code>.dataone_hosts()</code> or <code>.dataone_host_regex()</code>. EDI is itself a real, independently-run DataONE member node (same federation architecture the file's own header comment describes), so this is not a platform metacheck lacks support for in principle -- it is the DataONE detector's registered host/DOI-prefix table missing this specific, very common member node. Because <code>dataone_links()</code> builds its matching regex strictly from that host list and DOI-prefix list (lines 44-47, 77-80 of <code>archive-dataone.R</code>), a bare <code>10.6073/pasta/...</code> DOI or a <code>portal.edirepository.org</code> / <code>pasta.lternet.edu</code> URL will never match, regardless of how the link is phrased in the text. This is a config-table gap in an already-supported platform family (DataONE), analogous to how <code>archive-dataverse.R</code> or <code>archive-figshare.R</code> maintain their own host lists -- adding EDI (host <code>portal.edirepository.org</code>, API base presumably <code>/package/</code> on <code>pasta.lternet.edu</code> or <code>pasta-s.lternet.edu</code>, DOI prefix 10.6073) to <code>.dataone_hosts()</code> would close this specific gap, following the same verification discipline the file's comments describe for its other entries. [CONFIG_GAP_TIER: confirmed_simple_fix -- live-verified today, a single host/DOI-prefix table entry resolves this with no further ambiguity.]
10.1111/1365-2745.14047	Supported platform family, but this host/DOI-prefix not registered	metacheck has a DataONE detector (R/archive-dataone.R), and EDI (Environmental Data Initiative) is a genuine DataONE member node -- its DOI shape <code>"doi:10.6073/pasta/<suffix>"</code> matches exactly the <code>"doi:<prefix><suffix>"</code> PID convention that <code>archive-dataone.R</code>'s own <code>.dataone_pid()</code> regex and comments describe for DataONE member nodes generally (e.g. the file's own worked example for PISCO: <code>"doi:10.6085/AA/marine_ltm.20.1"</code>). However, <code>.dataone_hosts()</code> in R/archive-dataone.R hard-codes an allowlist of only five explicitly verified member-node hosts: <code>arcticdata.io</code>, <code>knb.ecoinformatics.org</code>, <code>metacat.tfri.gov.tw</code>, <code>smithsonian.dataone.org</code>, and <code>data.piscoweb.org</code>, each paired with its own <code>api_base</code> and (where known) <code>doi_prefix</code>. Neither EDI's web host (<code>portal.edirepository.org</code>), its API host (<code>pasta.lternet.edu</code>), nor its DOI prefix (10.6073) appears anywhere in that list or anywhere else in R/ (confirmed via <code>grep</code>: zero matches for <code>"10.6073"</code>, <code>"edirepository"</code>, or <code>"pasta.lternet"</code> in the whole R/ directory). Consequently <code>dataone_links()/dataone_host()/dataone_pid()</code> would not recognize this URL as a DataONE host at all -- the <code>host_regex</code> and <code>doi_regex</code> built from <code>.dataone_hosts()</code> simply would not match <code>"pasta/7f004957c6cc224d84e61bf1f4bb76a4"</code> or the 10.6073 prefix, so the citation would never even be classified as a candidate DataONE link, let alone queried. This is precisely a config gap in an otherwise-supported platform family: the detector code and API-querying logic for DataONE member nodes already exist and work (as documented by the file's own comments about adding new verified hosts), but EDI specifically has not yet been added to the verified-host list the way <code>arcticdata.io</code>, <code>KNB</code>, etc. were. [CONFIG_GAP_TIER: confirmed_simple_fix -- live-verified today, a single host/DOI-prefix table entry resolves this with no further ambiguity.]

Paper (DOI)	Category	Why Metacheck misses it
10.1111/2041-210X.13547	Supported platform family, but this host/DOI-prefix not registered	Zenodo is a fully supported platform in metacheck (<code>R/archive-zenodo.R</code>, <code>zenodo_links()</code>). That function already implements a deliberate two-tier search: real hyperlinks from the paper's url table, plus a body-text fallback regex for bare DOIs like "10.5281/zenodo.NNNNNNN" that appear without a hyperlink. The problem for this specific paper is that the fallback regex is run through <code>text_search(paper, zen_bare_regex, return = "match", perl = TRUE)</code> without passing <code>include_refs = TRUE</code>, and <code>text_search()</code> defaults to <code>include_refs = FALSE</code>, which explicitly filters out any text row whose <code>section_type</code> is "references" (see <code>text_search.R</code>, the <code>`section_filter <- !text\$section_type %in% "references"</code> line, used whenever <code>include_refs</code> is not set to <code>TRUE</code>). In this paper, the inline data-availability sentence is a bare author-year citation with no DOI or URL at all ("The Paris 12 dataset is available on Zenodo (Machon, 2020)."; similarly for the code, "available on Zenodo (Louvet, 2020), as well as on a GitHub repository"), and the only place the actual resolvable Zenodo DOIs (10.5281/zenodo.3770339 and 10.5281/zenodo.4314518) appear anywhere in the extracted text is inside the References section's bibliography entries. Since <code>zenodo_links()</code> never searches the references section for its bare-DOI fallback, it finds no hyperlink and no in-body bare DOI, so it returns zero Zenodo URLs for this paper even though both repositories are live, public Zenodo records with files. This is not a missing backend (Zenodo is fully implemented) and not a garbled DOI (the DOIs are intact and correctly formatted) -- it is a configuration/search-scope gap in how the existing Zenodo detector decides which part of the document to scan for bare identifiers, which this specific paper's citation style (name-year in text, DOI only in the reference list) exposes. [CONFIG_GAP_TIER: still_uncertain -- host could not be live-verified (blocked, timed out, or inconclusive), or the fix is a code change rather than a table entry.] (Second, independent audit of this paper reached the same conclusion as the earlier one in the 54-paper batch.) Four <code>hdl.handle.net</code> links under prefix 21.15109/CONCORDA/" resolve to a Dataverse installation hosted at <code>researchdata.hu</code>. Metacheck's <code>archive-dataverse.R</code> fully supports Dataverse generically, but its hardcoded host <code>allowlist (.dataverse_hosts())</code> does not include <code>researchdata.hu/science-data.hu</code> (confirmed zero <code>grep</code> matches). The DAS body text also has a duplicated 'CONCORDA' token spliced into the URLs, a compounding but secondary extraction defect, since the same handles appear correctly (only whitespace-broken) in the reference list. [CONFIG_GAP_TIER: confirmed_simple_fix -- live-verified today, a single host/DOI-prefix table entry resolves this with no further ambiguity.]
10.1111/2041-210X.14074	Supported platform family, but this host/DOI-prefix not registered	<code>metacheck</code> fully supports this exact platform and host: <code>rdr.ucl.ac.uk</code> is explicitly hardcoded in <code>R/archive-figshare.R</code>'s <code>.figshare_vanity_hosts()</code> (line 88, labeled "University College London (UK)"), the <code>figshare.com</code> DOI prefix 10.5522 is not even needed here because the underlying numeric article id (13469022) appears directly, and <code>.figshare_id()</code>'s regex patterns are written to match exactly this "10.6084/m9.figshare.N" or vanity-host "/articles/.../N" shape. The reason <code>metacheck_any_detector_hit</code> is false is almost certainly that the GROBID-extracted <code>full_text</code> corrupted all four DOIs identically: every occurrence reads "10.5522/04/13469 022.v1" (and 13469 085, 13469 139, 13469 166) with a literal space inserted between the 5th and 6th digit of the article id, splitting what should be one contiguous numeral "13469022" into "13469" + " " + "022". <code>figshare_links()</code>'s bare-text regex fallback (and <code>.figshare_id()</code>'s digit-capture groups) require an unbroken run of digits (<code>[0-9]+</code>); a stray embedded space breaks that run, so the regex simply does not match this citation as extracted, even though a human reader (and the live DOI resolver) has no trouble with it. This is a text-extraction artifact (GROBID/PDF-to-text splitting digit groups, the same class of bug already fixed elsewhere for effect-size subscripts per the recent commit history on this branch) rather than any gap in metacheck's Figshare/UCL-RDR platform support, which is otherwise a fully correct, already-registered configuration.
10.1002/pan3.10185	Citation corrupted by the PDF-to-text extraction step	

Paper (DOI)	Category	Why Metacheck misses it
10.1002/pan3.10623	Citation corrupted by the PDF-to-text extraction step	Paper states 'All data used in the analysis can be found freely online on Figshare,' but the adjoining reference-list DOI is corrupted by PDF column-reflow (an unrelated reference interleaved mid-title; no Figshare DOI/URL survives extraction). Independently located the real dataset via Figshare's search API: DOI 10.6084/m9.figshare.25066136.v1, verified live (is_public TRUE, MIT license, 11 files). Metacheck's archive-figshare.R fully supports resolving this DOI via the documented api.figshare.com/v2 REST API -- it never had a chance to act because no identifier survived extraction.
10.1111/1365-2435.14477	Citation corrupted by the PDF-to-text extraction step	metacheck already fully supports this exact host and DOI prefix: R/archive-dataverse.R's .dataverse_hosts() explicitly lists "phys-techsciences.datastations.nl" (one of the four DANS Data Stations, which are Dataverse installations, not a generic institutional repository as the coder's label "DANS (Data Archiving and Networked Services, Netherlands)" implies), and .dataverse_doi_prefix_hosts() explicitly maps that host to the "10.17026" DOI prefix used by all DANS Data Stations. So this is not a platform gap or a config gap -- the detector's host/prefix registry already covers it. The reason metacheck's detector did not fire (metacheck_any_detector_hit = false) is that the GROBID/PDF text extraction garbled the one and only DOI in the paper: it appears as "https://doi. org/ 10.17026/ dans-zqs-g6cn" with spurious internal whitespace splitting "doi.org/", "10.17026/", and "dans-zqs-g6cn" into separate tokens, inside a reference-list citation rather than in the data-availability sentence itself (which is truncated mid-word and has no link at all: "...are stored in DANS (Data archiving and networked...)."). A regex expecting a contiguous "10.17026/dans-zqs-g6cn" or "doi.org/10.17026/..." string would not match this whitespace-split form, so the extraction/garbling is the actual failure point, not any lack of platform or host support.
10.1111/1365-2435.14696	Citation corrupted by the PDF-to-text extraction step	Metacheck fully supports Zenodo via a dedicated R/archive-zenodo.R module, which explicitly anticipates bare "10.5281/zenodo.NNNNNNN" citations without a URL scheme (see the module's own comment: "10.5281/zenodo.1234567 is routinely cited without any URL scheme at all"), and this paper's DOI (10.5281/zenodo.13986711) is exactly that standard Zenodo prefix pattern, resolving live to a public CC-BY-4.0 Zenodo record (zenodo.org/records/13986711) with a CSV dataset and README. So this is not a platform-support gap. The failure is that the in-text data-availability sentence itself is truncated by the PDF text extraction right after "Data reported in the manuscript are available from the Zenodo" -- the DOI never appears in that sentence at all in full_text. The DOI does appear later, in the separate bibliography entry for Lu, Wenglein, Bluhm, Stuckenberg, Potapov & Scheu (2024), but there it is corrupted by stray whitespace injected during extraction: "https://doi. org/ 10.5281/ ZENODO.13986711" (spaces after "doi.", after "org/", and after "10.5281/"). A DOI/URL regex that requires a contiguous "doi.org/10.5281/zenodo.NNNNNNN" string without internal whitespace tolerance would fail to match this specific occurrence, even though the Zenodo backend itself would handle the ID correctly once given a clean DOI. This is best classified as a garbled-DOI extraction gap rather than a missing platform or a genuinely absent link.

Paper (DOI)	Category	Why Metacheck misses it
10.1111/1365-2656.13345	Citation corrupted by the PDF-to-text extraction step	Metacheck fully supports Dryad via R/archive-dryad.R, whose <code>dryad_links()</code> function matches bare Dryad DOIs with the regex <code>"(?:https?://)?(?:doi\\.org/)?10\\.5061/dryad\\.[A-Za-z0-9]+"</code> (and <code>.dryad_doi()</code> re-applies the equivalent capturing pattern <code>"(10\\.5061/dryad\\.[A-Za-z0-9]+)"</code>). In this paper's extracted <code>full_text</code>, the DOI is corrupted by what looks like a PDF line-wrap or column-break artifact that inserted a literal space inside the Dryad suffix: "https://doi.org/10.5061/dryad.hmgqn k9f3" in the data-availability statement, and again as "https://doi.org/10.5061/dryad.hmgqn k9f3" in the reference list (La Sorte & Graham, 2020 citation). Because "[A-Za-z0-9]+" cannot match across a whitespace character, the regex engine stops at "hmgqn" and either extracts the truncated/incorrect DOI "10.5061/dryad.hmgqn" or otherwise fails to recover the real, correct suffix "hmgqnk9f3". The live repository itself is a completely ordinary, fully supported, publicly accessible Dryad dataset with files present -- this is not a platform gap or a configuration gap, but a text-extraction corruption of the DOI string that defeats the existing, otherwise-correct Dryad regex.
10.1111/1365-2656.13893	Citation corrupted by the PDF-to-text extraction step	metacheck fully supports Dryad (R/archive-dryad.R), and its <code>dryad_links()</code> regex <code>(10\\.5061/dryad\\.[A-Za-z0-9]+, plus a doi\\.org/ prefix variant)</code> would ordinarily match this exact DOI (10.5061/dryad.7d7wm3806, under the explicitly allow-listed 10.5061 prefix) without any config change needed. The reason <code>metacheck_any_detector_hit</code> is false is that the PDF-extracted <code>full_text</code> corrupts the DOI at both of its two occurrences with stray whitespace introduced during text extraction, not because the platform or DOI prefix is unsupported. In the data-availability statement, the text reads <code>"https://doi.org/10.5061/dryad.7d7wm 3806"</code>, with a space inserted between <code>"7d7wm"</code> and <code>"3806"</code> that splits the DOI suffix into two tokens; since the regex's <code>[A-Za-z0-9]+</code> character class does not match whitespace, it would at most capture the truncated <code>"dryad.7d7wm"</code> and miss the real dataset DOI <code>"dryad.7d7wm3806"</code>. In the reference-list citation (Riecke, Hegelbach, & Schaub, 2023), the second occurrence instead has the space between <code>"doi."</code> and <code>"org"</code> (<code>"https://doi. org/10.5061/dryad.7d7wm3806"</code>), which breaks the literal <code>"doi\\.org\\."</code> substring the regex's <code>doi.org/-prefix</code> alternative requires, though the bare <code>"10.5061/dryad\\."</code> alternative in the pattern could still match this one correctly since the DOI digits themselves are intact there. Because the live dataset is confirmed real, public, and directly reachable at exactly the DOI the human coder recorded, this is a text-extraction/OCR-layer corruption gap specific to this PDF's rendering of the DOI (likely a kerning or line-wrap artifact in the original Wiley PDF), not a platform-support gap, a configuration gap, or a detector logic bug against clean input.

Paper (DOI)	Category	Why Metacheck misses it
10.1111/1365-2656.14000	Citation corrupted by the PDF-to-text extraction step	metacheck fully supports this exact case at the code level: R/archive-figshare.R already registers "10.26180" => "bridges.monash.edu" in .figshare_doi_prefix_hosts(), and .figshare_id() (lines ~307-344) carries a comment explicitly citing this very paper's DOI as a confirmed-working worked example ("confirmed live 2026-09-19: https://doi.org/10.26180/19095317.v1 resolves to Monash article id 19095317"). The failure is not in the detector's logic but in the upstream PDF text extraction: in full_text, the DOI string is rendered with two stray internal spaces -- "https://doi.org/10.26180/ 19095 317.v1" -- splitting the contiguous digit string "19095317" into "19095" and "317" (apparently a line/column-wrap artifact from the two-column PDF layout, identical in both the reference-list occurrence and the implied data-availability-statement occurrence). The very first scanning stage that finds Figshare candidate strings, figshare_links()'s fs_bare_regex (line 220 of archive-figshare.R), requires the DOI suffix to match "[A-Za-z0-9_-]+" contiguously with no whitespace, so this regex never matches the garbled string and no candidate URL is even extracted for .figshare_id() to later resolve. The coder's own recorded repo_link_verbatim has the identical "19095 317" space-split, confirming this is a property of the underlying text extraction itself, not an error either the coder or I introduced. The platform (Figshare via Monash Bridges), its DOI-prefix registration, and its id-resolution regex are all already correctly built and tested for the clean form of this exact DOI; only the whitespace-corrupted rendering of this one citation defeats the first-stage text-search regex, which is a text-extraction/garbling problem, not a missing or buggy detector.
10.1111/1365-2656.14054	Citation corrupted by the PDF-to-text extraction step	Code citation appears only in the reference list (no DAS exists): 'Amarasekare, P. (2024)... Zenodo https://doi. org/ 10.5281/ ze- nodo.10497888' -- PDF extraction inserted spaces after 'doi.', after 'org/', after '10.5281/', and hyphen-broke the word 'zenodo' itself. Metacheck's zenodo_links() regex (archive-zenodo.R) requires a contiguous string and cannot match text containing embedded spaces/hyphen-breaks. Live-verified the Zenodo record resolves and is open, CC-BY-4.0, linking to the companion GitHub repo.
10.1111/1365-2664.13890	Citation corrupted by the PDF-to-text extraction step	metacheck_any_detector_hit was reported as false, but this looks like it should actually be detectable as a dspace/dspace7 case, so I lean toward SUPPORTED_PLATFORM_BUG/CONFIG_GAP rather than a clean no-gap -- see caveat below. The repository DOI 10.20350/digitalCSIC/13825 resolves to digital.csic.es/handle/10261/236117, and the "/handle/<naan>/<suffix>" path is the canonical DSpace resolver URL pattern, confirming DIGITAL.CSIC is a generic DSpace installation (the coder's platform label "DSpace-based" is correct and matches the URL evidence, independent of the live Anubis block). Since metacheck's supported-platform list explicitly includes "dspace/dspace7 (generic DSpace installations)," the platform family itself is supported. The DOI is also present, unambiguously, in the paper's reference list, though the raw PDF-extracted text has it corrupted by stray inserted spaces ("https://doi.org/10.20350/ digit alCSI C/13825" instead of ".../digitalCSIC/13825") -- a GROBID-style extraction artifact that could defeat a naive strict-pattern link extractor looking for an unbroken "digitalCSIC" token or path segment, even though the DOI prefix "10.20350" and the surrounding citation context are intact and a human (or a more tolerant regex/whitespace-stripping extractor) would still recognize it. Given metacheck_any_detector_hit=false, the most defensible classification is that this is a DOI-level extraction problem (stray whitespace inside the handle/suffix string) rather than a missing-platform gap, so I classify it as EXTRACTION_GAP_DOI_GARBLED: the link is present in the text but corrupted by split/inserted whitespace in a way that likely prevents metacheck's extractor from recognizing it as the clean "10.20350/digitalCSIC/13825" DOI string, even though the DSpace/DIGITAL.CSIC platform itself is on metacheck's supported list and would otherwise be detectable.

Paper (DOI)	Category	Why Metacheck misses it
10.1111/1365-2664.13961	Citation corrupted by the PDF-to-text extraction step	metacheck has a dedicated, well-built Dryad detector (R/archive-dryad.R) whose regex matches <code>datadryad.org/dataset(s)/doi:</code> URLs as well as bare <code>"10.5061/dryad.XXXXX"</code> DOIs with or without a URL scheme, so the platform itself is fully supported and this exact DOI shape would normally be caught. The reason <code>metacheck_any_detector_hit</code> is false here is that the real Dryad link is corrupted at both places it appears in the extracted <code>full_text</code>: (1) in the data-availability statement, the URL token has been replaced by the author's ORCID link (<code>https://orcid.org/0000-0002-9304-228X</code>) -- almost certainly a PDF-to-text extraction artifact where a hyperlinked DOI field was resolved to the wrong embedded link target; and (2) in the reference-list citation, the DOI itself is printed one character short as <code>"10.5061/dryad.0pngf20h"</code> (missing the digit "2"), which I verified live returns HTTP 404 rather than resolving to the dataset. Only the correct, non-garbled DOI <code>"10.5061/dryad.0p2ngf20h"</code> -- which does not appear anywhere in the extracted text -- actually resolves to the real public dataset. Since neither occurrence in the text metacheck would see contains a matchable, correct Dryad identifier, no amount of improving the Dryad detector logic would fix this particular case; the gap is upstream, in the source PDF/text extraction corrupting the DOI before metacheck ever sees it.
10.1111/1365-2664.14229	Citation corrupted by the PDF-to-text extraction step	metacheck already fully supports this exact repository: <code>dataverse.unimi.it</code> is explicitly listed in <code>.dataverse_hosts()</code> (R/archive-dataverse.R, line 137, "Università degli Studi di Milano (Italy)"), and its DOI prefix <code>10.13130</code> is explicitly mapped to it in <code>.dataverse_doi_prefix_hosts()</code> (line 314: <code>"dataverse.unimi.it" = "10.13130"</code>), individually verified against DataCite per that function's header comment. This is a correctly-configured, supported Dataverse installation with no gap in metacheck's platform or host/prefix registry. The failure is purely in the paper's extracted text: the data-availability sentence gives the link as <code>"http://10.0.51.74/RD_UNIMI/5ZXGIV"</code> -- the real DOI resolver domain was corrupted into what looks like an internal/private IPv4 address during PDF extraction, so it contains neither the <code>"dataverse.unimi.it"</code> hostname string (which <code>dataverse_links()</code>'s host-regex path requires) nor the <code>"10.13130"</code> DOI prefix (which the bare-DOI-prefix regex path requires). The second occurrence, in the reference list (<code>"https://doi.org/10.13130/ RD_UNIMI/ 5ZXGIV"</code>), does carry the correct, verified <code>10.13130</code> prefix, but PDF line-wrap extraction inserted literal spaces after each forward slash, splitting the DOI into three fragments (<code>"10.13130/"</code>, <code>"RD_UNIMI/"</code>, <code>"5ZXGIV"</code>). metacheck's DOI-prefix regex (an unbroken <code>"[A-Za-z0-9/._-]*"</code> after the prefix) cannot match across that whitespace, so this occurrence also fails to resolve. Both failures are text-extraction corruption of the same underlying, already-registered link, not a gap in platform coverage or host/prefix configuration -- this is a GROBID/PDF-extraction defect (domain substitution in one place, inserted whitespace in another) rather than anything metacheck's Dataverse module needs to be extended to handle.

Paper (DOI)	Category	Why Metacheck misses it
10.1111/1365-2745.13577	Citation corrupted by the PDF-to-text extraction step	Dryad is a platform metacheck explicitly supports, and the dataset is genuinely live and public (confirmed by following the DOI redirect to datadryad.org). The reason metacheck's detector did not fire (metacheck_any_detector_hit: false) is that in the extracted full_text, every occurrence of the Dryad DOI -- both in the in-text data-availability statement ("Data on plant phenology are available at the Dryad Digital Repository: https://doi.org/10.5061/dryad.c59zw_3r69") and in the reference-list entry (Bucher & Römermann, 2020a) -- contains a stray internal space that splits "c59zw3r69" into "c59zw" and "3r69". This is almost certainly an artifact of PDF line-wrapping/hyphenation during text extraction (GROBID or similar), not a typo in the original PDF, since a genuinely broken DOI would not be something the original authors or Wiley typesetting would publish twice identically. Because the space breaks the DOI into two whitespace-separated tokens, any regex-based or URL-pattern detector looking for a well-formed <code>`10.5061/dryad.xxxxx`</code> DOI or a contiguous <code>`doi.org/10.5061/dryad...`</code> URL would fail to match it as written in this extracted text, even though the platform (Dryad) is fully supported by metacheck. This is a text-extraction corruption problem upstream of the detector logic, not a gap in platform coverage or a config issue -- the fix would belong in the extraction/normalization step (e.g. stripping/collapsing whitespace inside DOI-shaped strings before pattern matching), not in adding a new platform detector. Separately, a second repository link in the paper (TRY database, Dataset 416, a plant-trait database) is on an unsupported platform and was not something the coder recorded either, but it is incidental to this specific gap -- the primary data-sharing link the coder flagged is the Dryad one, and that one is a garbled-DOI extraction issue.
10.1111/1365-2745.13606	Citation corrupted by the PDF-to-text extraction step	Dryad is a fully supported platform in metacheck (R/archive-dryad.R), including a specific fallback in <code>dryad_links()</code> for exactly this scenario: a bare, non-hyperlinked DOI mention in body/reference text, matched via the regex <code>`(?:https?:/)?(?:doi\..org/)?10\.5061/dryad\.[A-Za-z0-9]+`</code> (plus a <code>found_href</code> branch for real anchor-tag links, which does not apply here since the coder's own note confirms no clickable URL exists in the <code>extracted_urls</code> table). The detector still misses this citation because the text extraction (likely GROBID, given the paper's Wiley/BES PDF source) corrupted the DOI across what was probably a line wrap in the original PDF: the extracted full_text reads <code>"https://doi.org/10.5061/ dryad.f7m0c fxv3"</code> (verified by direct substring inspection of full_text) — note two inserted spaces, one between <code>"10.5061/"</code> and <code>"dryad."</code>, and a second one splitting the identifier itself into <code>"f7m0c"</code> and <code>"fxv3"</code>. Because <code>`[A-Za-z0-9]+`</code> cannot match across embedded whitespace, neither the <code>dryad_bare_regex</code> pattern nor the DOI-extraction helper <code>`dryad_doi()`</code> (which uses the same <code>`10\.5061/dryad\.[A-Za-z0-9]+`</code> pattern) can recognize this as a single token, so <code>dryad_links()</code> returns no match for this paper. This is a text-extraction/garbling problem upstream of metacheck's Dryad detector logic, not a missing platform or a misconfigured host: live-checking the DOI confirms it resolves correctly to a real, public, two-file Dryad dataset that exactly matches the paper's subject matter, so the citation itself is intact in the published paper — only the extracted text metacheck works from is corrupted at this one point.

Paper (DOI)	Category	Why Metacheck misses it
10.1111/1365-2745.13694	Citation corrupted by the PDF-to-text extraction step	Metacheck has a full, working Zenodo detector (R/archive-zenodo.R), and its design is explicitly built to catch bare DOI citations of the exact form "10.5281/zenodo.NNNNNNN" with or without a URL scheme or doi.org prefix (see the regex and comment at R/archive-zenodo.R lines 5 and 28: "10.5281/zenodo.1234567" is routinely cited without any URL scheme at all"). The gap here is not a missing platform but a text-extraction corruption: in the extracted full_text, the citation reads "Zenodo, https://doi.org/10.5281/ zenodo.3756781" with a literal space inserted between "10.5281/" and "zenodo.3756781" (almost certainly a PDF line-wrap artifact from GROBID/PDF-to-text conversion at a hyphenation or column break). Metacheck's R/text-extractors.R extract_urls() function, which is the upstream source of link candidates that archive-zenodo.R and other archive-*.R detectors work from, uses a \ b-bounded regex requiring a contiguous domain-like token (([\ w.-]+) \ .[a-z]{2,}) with no internal whitespace. That regex matches "https://doi.org/10.5281/" and "zenodo.3756781" only as two separate, incomplete fragments across the inserted space, neither of which satisfies the Zenodo-ID pattern (10 \ .5281/zenodo \ .[0-9]+) that archive-zenodo.R looks for. The result is that no valid Zenodo candidate URL is ever produced for the detector to test, which is consistent with metacheck_any_detector_hit being false even though the paper does cite a fully public, CC-BY-4.0, live Zenodo dataset with three files, and even though Zenodo is a fully supported platform. This is a text-extraction/garbling problem (a split DOI token), not a missing detector and not a dead repository.
10.1111/1365-2745.13875	Citation corrupted by the PDF-to-text extraction step	metacheck fully supports Dryad (R/archive-dryad.R), and the DOI prefix 10.7280 used by this exact dataset is explicitly in .dryad_doi_prefixes() (archive-dryad.R line 37), with dryad_links()'s regex built to match "10.7280/[A-Za-z0-9]+" as a bare in-text DOI even without a scheme. The failure is not a platform gap but a text-extraction corruption: the paper's primary data-availability sentence reads "Data available from the Dryad Data Repository http://10.0.28.112/D1D99J (Campbell et al., 2022)" -- the DOI prefix "10.7280" has been replaced by what looks like a private IP address "10.0.28.112" (almost certainly a PDF/GROBID hyperlink-extraction artifact mangling the embedded doi.org URL), so neither the datadryad.org host string nor the 10.7280 prefix pattern is intact at that location for dryad_bare_regex to match, and it is unlikely the href/url table preserved a clean target either, since the same corrupted string appears in the flattened text. A second, well-formed copy of the same DOI (https://doi.org/10.7280/D1D99J) exists in the reference list under the Campbell et al. (2022) entry; in principle this clean occurrence should satisfy the regex, but the reported metacheck_any_detector_hit=false indicates it did not fire in the live run, most plausibly because GROBID-style extraction of the actual PDF introduces a line-wrap or whitespace split around the DOI (e.g. "doi. org/10.7280/D1D99J") that defeats the regex even though the flattened full_text I was given displays it as one unbroken string. I manually confirmed the DOI itself is completely valid and live: doi.org/10.7280/D1D99J redirects cleanly to a real, public, accessible Dryad dataset page with six CSV files and four READMEs, so this is a genuine false negative caused by extraction-level corruption/formatting of the DOI text, not a missing platform or a dead repository.

Paper (DOI)	Category	Why Metacheck misses it
10.1111/1365-2745.13882	Citation corrupted by the PDF-to-text extraction step	Metacheck fully supports this exact platform and host: R/archive-dataone.R's .dataone_hosts() already registers host = "arcticdata.io" with doi_prefix = "10.18739" (line 33), which is precisely the DOI prefix all four of this paper's Arctic Data Center citations use, and dataone_links()'s bare-DOI regex (lines 105-109) is built to catch exactly this "https://doi.org/10.18739/<suffix>" citation shape. Live-checking confirmed the dataset is real, public, CC0-licensed, and resolvable through the identical Metacat object endpoint metacheck's own .dataone_info() queries. The actual failure is upstream, in the PDF-to-text extraction that produced full_text: within the data-availability statement itself, three of the four DOI suffixes are split by a spurious whitespace character in the middle of the identifier -- 'https://doi.org/ A2GT5 FG86' instead of 'https://doi.org/A2GT5FG86', 'http://doi. org/10.18739/ A23R0 PV43' instead of '.../A23R0PV43', and 'http://doi. org/10.18739/A2GX44W2M' with a space inserted inside 'doi.org'. Only the second DOI (A2000018G) survived extraction intact. dataone_links()'s matching regex for the bare-DOI fallback is '(?:doi\ \.org/)?(?:10\ \.18739)/[A-Za-z0-9._-]+', which requires the suffix to be one unbroken token of those characters with no whitespace; a space inside 'A2GT5 FG86' terminates the match after 'A2GT5', producing an incomplete, non-resolvable PID rather than no match at all, and likewise breaks the 'doi.org' host token itself in the third and fourth citations ('doi. org'). This is a genuine upstream text-extraction corruption (apparently from the PDF layout, where a line-wrap or font-kerning artifact inserted a space mid-token), not a gap in metacheck's own DataONE detector logic, host list, or regex design for the well-formed case. This matches the EXTRACTION_GAP_DOI_GARBLED category: the link is present in the source text but corrupted in a way that defeats token-based regex matching.
10.1111/1365-2745.13999	Citation corrupted by the PDF-to-text extraction step	Dryad is a fully supported platform in metacheck (R/archive-dryad.R exists and the live link resolves cleanly to a standard public Dryad landing page with no login and a "no restrictions" license), so this is not a platform-support gap. The failure is a text-extraction corruption: in full_text the DOI appears as "https://doi.org/10.5061/dryad.7sqv9 s4sh" -- a literal space has been inserted between "7sqv9" and "s4sh," splitting the Dryad suffix in two. Any DOI/URL-matching regex metacheck runs over the extracted text (whether matching "10.5061/dryad\ \. \ S+" or a general doi.org URL pattern) would stop at the whitespace and either fail to match the full suffix or extract "dryad.7sqv9" as a truncated/invalid DOI, which would not resolve via the Dryad API. The same kind of mid-token space corruption is visible elsewhere in this paper's extracted references (e.g. "foreco.2012.02" rendered as ". foreco.2012.02", "isprs jprs.2020.01.012" instead of "isprsjprs..."), indicating this is a systematic PDF-extraction artifact (likely from line-wrap/hyphenation handling in GROBID or whatever extractor produced full_text) rather than something specific to Dryad's platform support. Additionally, no distinct "Data Availability Statement" section/heading was present anywhere in the extracted full_text at all -- the only occurrence of the link is embedded in the reference list citing the companion Dryad data-descriptor record, suggesting the statement box itself (common in BES/Journal of Ecology articles as a separate panel) may not have been captured by extraction, compounding the problem for any heading-based statement-locator logic upstream of the regex match.
10.1111/2041-210X.12918	Citation corrupted by the PDF-to-text extraction step	DAS cites two locations: 'The version 1.0.0 of ecolottery is available on CRAN... The development version is available on Github (https://github.com/frmunoz/ [extraction cuts off here, jumping into Authors' Contributions]).' The repo-name path segment ('ecolottery') is entirely absent from the extracted text due to a PDF extraction defect -- Metacheck's github_links() regex requires both org AND repo segments, so with only 'frmunoz' present it cannot match even in principle. The adjacent CRAN URL in the same sentence is independently also garbled, corroborating sentence-level extraction corruption.

Paper (DOI)	Category	Why Metacheck misses it
10.1111/2041-210X.13350	Citation corrupted by the PDF-to-text extraction step	Paper cites two links: CRAN (extracted cleanly, twice) and GitHub ("https ://github.com/ achub aty/grain scape", extracted with spaces inserted inside the URL and inside both the owner and repo names). Any regex-based URL matcher fails on the garbled GitHub form even though GitHub has a complete, working backend in Metacheck. CRAN is cited cleanly but is moot since no CRAN backend exists at all.
10.1111/2041-210X.13681	Citation corrupted by the PDF-to-text extraction step	Both repository DOIs ARE present in the extracted full_text, and both platforms (Dataverse, confirmed by the "Powered by The Dataverse Project v6.8.4-SP" footer at borealisdata.ca, and Zenodo, confirmed by the zenodo.org landing page) are on metacheck's supported-platform list, so this is not a coverage gap. The failure is a text-extraction/OCR artifact: the PDF-to-text conversion injected a stray space exactly at the point where each DOI line-wrapped in the original PDF, producing "https://doi. org/10.7939/DVN/QYMAYV" (space between "doi." and "org") for the Dataverse reference and "https://doi.org/10.5281/ zenodo.5070812" (space between "10.5281/" and "zenodo.5070812") for the Zenodo reference. Any DOI-matching regex or URL tokenizer that expects a contiguous string (no embedded whitespace) will fail to recognize either as a valid, resolvable DOI/URL, even though a human reader can see and manually reconstruct both. This matches metacheck_any_detector_hit=false: the detector presumably could not match either garbled string. Separately, the in-text Data Availability Statement itself uses named-citation style ("available on UAL Dataverse Derocher (2021)"; "available on Zenodo Togunov et al. (2021)") with no inline clickable URL -- the resolvable DOIs exist only in the back-matter reference list, so even a perfect-whitespace extraction would still require cross-referencing the in-text citation key to the bibliography entry to recover the link, which is an extra step beyond simple statement-adjacent URL scanning. The coder's recorded link (https://doi.org/10.7939/DVN/QYMAYV, acquired via "a targeted web search" rather than from the text) is confirmed correct and live, but the coder missed the second repository (the Zenodo replication code) entirely.
10.1111/2041-210X.13821	Citation corrupted by the PDF-to-text extraction step	metacheck fully supports this exact platform and host: R/archive-dspace7.R's .dspace7_hosts() explicitly lists "conservancy.umn.edu" (annotated "Data Repository for the University of Minnesota (USA)"), and dspace7_links() already contains code specifically added for this exact failure mode (documented inline as issue #435, citing the near-identical sibling DOI 10.13020/q3y6-h459 on the same host): when a paper cites only a bare DOI with the host domain never appearing in the surrounding text, the function live-resolves every DOI-shaped string found via text_search() and keeps any that redirect to a known DSpace7 host. That mechanism would correctly catch this paper's citation -- except that the DOI as it literally appears in paper's full_text is garbled by a stray inserted space: "https://doi.org/10.13020/ 1f35-rn20" instead of "https://doi.org/10.13020/1f35-rn20" (verified by extracting the exact sentence from full_text). I tested R/archive-dspace7.R's doi_bare_regex directly against both strings: it matches the clean DOI but returns zero matches against the garbled one, because the regex requires an unbroken token after "10.xxx/" and the inserted space breaks it mid-string. This is almost certainly a PDF-text-extraction artifact (the reference list line likely wrapped at that point in the original PDF, and the extraction tool inserted a space at the line break) rather than a metacheck regex design flaw -- the same gap would equally affect any other DOI-based detector in metacheck that depends on this shared doi_bare_regex pattern if a citation's DOI happens to be split by extraction at the same point. I did not attempt to fix this since the task calls for an audit and root-cause classification only.

Paper (DOI)	Category	Why Metacheck misses it
10.1111/2041-210X.13850	Citation corrupted by the PDF-to-text extraction step	CORRECTION to an earlier, unverified pass that claimed the GitHub repo was 'genuinely inaccessible... not a metacheck defect' -- this is live-verified FALSE; both the GitHub org and the actual repo (ConScape/ConScape.jl) return HTTP 200 via the GitHub API. Primary cause: the reference-list Zenodo citation has a space inserted between '10.5281/' and 'zenodo.5810881' (confirmed via grep on raw full_text), defeating zenodo_links()'s regex (tested directly in R, returns no match). Secondary: the body-text GitHub citation is a bare org URL with no repo segment, so github_repo() correctly returns NULL for that specific citation form -- not itself a bug, since no specific repo was actually named in that URL.
10.1111/1365-2656.12794	No trace of the citation in the extracted text	Dryad is one of metacheck's supported platforms (dataone, dataverse, dryad, dspace/dspace7, figshare, github, gitlab, mendeley, osf, psycharchives, rbox, researchdata4tu, reshare, zenodo, aspredicted), so this is not a platform-support gap. The actual problem is upstream of detection: the full_text metacheck extracted for this paper (apparently a PMC manuscript version, judging by the trailing "Refer to Web version on PubMed Central for supplementary material" and "J Anim Ecol. Author manuscript; available in PMC 2018 June01" notices embedded in the body) simply does not contain the journal's "Data Accessibility" statement or the Dryad DOI (10.5061/dryad.95486) anywhere in the extracted string — I read the entire full_text field end to end and confirmed no occurrence of "Dryad," "data availability," "data accessibility," or any doi.org/5061 string. Since metacheck's Dryad detector can only act on text it is given, and the availability statement text was never captured by the upstream extraction pipeline (likely because the PMC author-manuscript rendering omits or relocates the data-accessibility statement relative to the original journal HTML/PDF), the detector correctly reports metacheck_any_detector_hit = false through no fault of the Dryad-detection logic itself. The repository itself is independently verified live and public (confirmed via WebFetch to datadryad.org), so this is a text-extraction/input gap, not a detector-capability or configuration gap.
10.1111/1365-2656.12900	No trace of the citation in the extracted text	metacheck does have a dedicated Dryad detector (R/archive-dryad.R in the metacheck-new source tree), and the cited repository — https://doi.org/10.5061/dryad.mb64cr8 , redirecting live to datadryad.org/dataset/doi:10.5061/dryad.mb64cr8 — is a real, public, directly-matching Dryad dataset for this exact paper, with no login barrier. The detector nonetheless could not have fired because the actual DOI/URL string is simply absent from the full_text that was extracted from the PDF: the data-availability sentence reads only "Data for this publication are available online at the Dryad Digital" and is then cut off, running straight into the unrelated author-contributions paragraph ("S.P. conceived the ideas..."). No "Repository" word, no "10.5061" string, and no "doi.org" string appear anywhere in full_text. This is a pure text-extraction gap (likely the original PDF rendered the DOI as a hyperlink or footnote-style reference that the extraction pipeline dropped or failed to inline as visible text) — not a missing platform, not a configuration gap for an unrecognized host, and not a bug in the Dryad-matching logic itself, since that logic never receives the string it needs to match against.

Paper (DOI)	Category	Why Metacheck misses it
10.1111/1365-2656.13052	No trace of the citation in the extracted text	metacheck has a dedicated, working detector for Dryad (R/archive-dryad.R exists alongside the other platform modules such as archive-figshare.R, archive-dataverse.R, etc.), so this is not a platform-support gap. The failure is purely at the text-extraction stage: the DOI string "10.5061/dryad.vc80r13" that identifies this specific Dryad record is absent from the full_text metacheck extracted from the PDF. It appears the source PDF (a Wiley/Journal of Animal Ecology article) had the data-availability sentence and the corresponding reference-list entry each truncated mid-string during text extraction -- the data-availability statement stops at "...available from the Dryad Digital" (no "Repository", no URL), and the reference-list citation "Gibb, H., Grossman, B. F., Dickman, C. R., Decker, D., & Wardle, G. (2019). Data from... Dryad Digital Repository, https://doi.org/10.5061/dryad." is cut off immediately after "dryad." before the record suffix "vc80r13" that makes the DOI resolvable. Since no detector, regex, or URL-matcher can find a string that isn't present in the text it is given, metacheck's extracted_urls list correctly came up empty for this paper even though its Dryad backend would very likely have correctly identified and verified doi:10.5061/dryad.vc80r13 as a live, public dataset had the complete string been present in the extracted text.
10.1111/1365-2656.14093	No trace of the citation in the extracted text	metacheck does support Figshare as a platform, so if a resolvable Figshare URL or DOI (e.g. a "10.6084/m9.figshare.XXXXXXX" string or a figshare.com/articles/... link) were present in the text, the existing Figshare detector should be able to find and validate it. The actual problem here is upstream of detection: the paper's extracted full_text never contains such a string. The data-availability sentence reads only "Data are archived on Figshare(Mahoney & Pasch, 2024)," citing a reference-list entry that is itself corrupted by the same GROBID-style extraction/OCR process that produced this full_text -- "Mahoney, S. M., & Pasch, B. (2024). , 55(6), 1483-1487." is missing the dataset title and the words "Figshare" or any repository name, leaving only a volume/issue/page triplet. I independently searched Figshare and the open web for a Mahoney & Pasch 2024 dataset matching this paper's topic (mammalian food-caching evolution / Mt. Graham red squirrels) and could not locate one; the only Figshare DOI associated with these two authors (10.6084/m9.figshare.5285635) belongs to their unrelated 2017 "Caching for where and what" paper. Because no machine-resolvable identifier for the actual 2024 dataset exists in the extracted text (and I could not find one by searching, so I cannot even confirm the live dataset is reachable), metacheck's Figshare detector has nothing to act on; this is a text-extraction gap (the citation/DOI was either never present in the PDF's data statement as a hyperlink, or was stripped during PDF-to-text conversion), not a platform-coverage or detector-logic failure. metacheck_any_detector_hit=false is consistent with this: there is no link in the text for any detector, Figshare or otherwise, to find.

Paper (DOI)	Category	Why Metacheck misses it
10.1111/1365-2664.12927	No trace of the citation in the extracted text	Metacheck's figshare.R has full detector support for this exact institutional instance -- "ore.exeter.ac.uk" (Open Research Exeter's branded Figshare front end) is explicitly listed in .figshare_vanity_hosts(), and the figshare detection regex (.figshare_host_regex()) also matches any literal "figshare.com" substring, which covers the newer "uoe.figshare.com" subdomain this item now lives under. So this is not a platform gap: both the old branded host and the current host are things metacheck's figshare backend already knows how to query via api.figshare.com. The actual gap is at the extraction stage. The paper's data-availability sentence cites only a bare hdl.handle.net/10871/27138 URL -- no figshare.com domain, no figshare DOI (10.6084/m9.figshare... or an institutional DOI prefix), and no bare numeric article ID appear anywhere in the full text. figshare_links() works by regex-matching the paper's own hyperlink table and body text for figshare-looking strings (a figshare host in the URL, or a figshare DOI pattern); it does not follow live HTTP redirects to discover that a handle.net link currently resolves to Figshare. Separately, metacheck has no generic handle.net resolver at all -- the only handle.net handling in the codebase is in archive-psycharchives.R, scoped to PsychArchives' own handle prefix (20.500.12034), which has no bearing on Exeter's prefix (10871). Because the citation text itself never names figshare or exposes a figshare-recognizable identifier, no existing detector -- correctly implemented as each one is for the identifiers it is given -- has anything to match against, so detection fails for a text-extraction reason, not a missing-platform-support reason or a bug in the figshare detector itself.
10.1111/1365-2664.14138	No trace of the citation in the extracted text	Figshare DOI (10.6084/m9.figshare.16436379.v2) is cited exactly once, inside a formatted 'Data from:' entry in the References list; no DAS section exists in the extracted text at all. Metacheck's repo_check calls text_search() with default include_refs=FALSE, which filters out all rows tagged section_type=='references' (text_search.R lines 104-110) before pattern-matching -- so the citation is structurally invisible to the default search regardless of host support. Metacheck's Figshare backend is otherwise fully functional and would resolve this DOI if the text reached it (confirmed live via api.figshare.com/v2/articles/16436379, HTTP 200, CC BY 4.0).
10.1111/1365-2435.14431	On-disk cache entry predates a fix or wrongly cached a failed lookup	STALE CACHE, not a current code bug. A fresh, uncached, live rerun of module_run(paper,"repo_check") just now correctly finds repo_n=1/files_n=96 for OSF node fwjh9 -- the detector itself works. The stored res_repo_check.RData shows repo_n=0 because the pipeline is called with cache=TRUE (01_run_metacheck.R), and code/data/.metacheck_repo_info_cache/osf_fwjh9.rds (dated 2026-09-09, nearly a month stale) holds a cached data.frame(osf_id='fwjh9', osf_type='error') -- i.e. a FAILED lookup from whenever that cache entry was written (OSF's API is documented in the package itself as frequently down). This should never have been cached as a success: R/archive-osf-helpers.R calls .repo_info_cache_put() only when .repo_info_ok(result) is TRUE, and .repo_info_ok() (repo-info-cache.R) checks for a column literally named 'error' -- but OSF's own error frame uses the column name 'osf_type' with value 'error', not a column named 'error'. Because 'error' is not among that frame's column names, .repo_info_ok() falls through to its default TRUE, so a transient OSF failure was permanently cached as if it had succeeded, and cache=TRUE has silently replayed that false negative on every subsequent pipeline run since. This IS a real code bug (the .repo_info_ok() check should also recognise OSF's osf_type=='error' shape), separate from the immediate cause of this paper's repo_n=0 (the stale cache entry itself, fixable today by clearing that one cache file, independent of whether the .repo_info_ok() bug is also patched).

Paper (DOI)	Category	Why Metacheck misses it
10.1111/1365-2745.13708	On-disk cache entry predates a fix or wrongly cached a failed lookup	STALE CACHE, not a current code bug. A fresh, uncached, live rerun of <code>module_run(paper,"repo_check")</code> just now correctly finds <code>repo_n=1/files_n=2</code> (data CSV plus R script) for KNB record doi:10.5063/PG1Q4B -- the detector itself works correctly today. The stored <code>res_repo_check.RData</code> shows <code>repo_n=0</code> because the pipeline runs with <code>cache=TRUE</code>, and <code>code/data/.metacheck_repo_info_cache/dataone_knb.ecoinformatics.org_doi_10.5063_PG1Q4B.rds</code> (dated 2026-09-13) holds a genuinely empty <code>files=list()</code> result with NO error column (so <code>.repo_info_ok()</code> correctly, but prematurely, treated it as a valid 'ok' result worth caching permanently). Git history shows <code>R/archive-dataone.R</code> was patched for exactly this failure mode in commit <code>ba661c7</code> ('Fix DSpace7/DataONE metadata-and-file gaps...'), confirmed live 2026-09-29, issue #435 follow-up); some KNB EML records describe a file as an otherEntity with no physical child, which the OLD parser (in effect when the Sept 13 cache entry was written) returned zero files for; the fix added a fallback branch reading the entity's own id attribute. The Sept 13 cache entry predates that fix by over two weeks and has been silently replayed as a false negative on every pipeline run since, masking that the underlying bug was already fixed in the source.
10.1111/2041-210X.13766	Citation points to an organisation/user page, not one enumerable repository	<code>github_links()</code> correctly finds the literal availability-statement URL <code>github.com/OneZoom/</code> verbatim. However, inspecting <code>github_tree_files()/github_repo()</code> directly in the installed package shows <code>github_repo()</code> requires a regex match of the shape 'owner/repo'; a bare organization-root URL like <code>github.com/OneZoom</code> has no second path segment to match as a 'repo', so <code>github_repo('github.com/OneZoom')</code> deterministically returns NULL (confirmed by direct invocation today), producing 'invalid or inaccessible GitHub repository'. This is not a bug in the usual sense -- a GitHub organization page is not itself a single repository with an enumerable file tree; the 12 individual repos under it are the actual archives, and the paper's own availability statement never names one of them specifically. None of the existing categories cleanly describe 'the cited link points to a multi-repository org/user root rather than one repository'; flagging as <code>SUPPORTED_PLATFORM_CONFIG_GAP</code> is the closest existing-category fit, but a more precise <code>NEW</code> label would be <code>GITHUB_ORG_ROOT_NOT_A_SINGLE_REPO</code> (flagged as new).
10.1111/2041-210X.14080	Citation points to an organisation/user page, not one enumerable repository	Same underlying mechanism as 10.1111/2041-210X.13766: <code>github_repo('github.com/gobbleturk')</code> deterministically returns NULL (confirmed by direct invocation today) because the citation is a bare GitHub user-root URL with no repository segment, not <code>'github.com/gobbleturk/SpectralRangeMaps'</code> (confirmed: <code>github_repo()</code> DOES correctly resolve the full repo path when tested with the complete owner/repo string). More importantly for this paper, the data/code availability statement's primary, citable archive is the Zenodo DOI 10.5281/zenodo.7569498 (a versioned, permanent archive of the same code); the bare GitHub mention is a secondary, non-versioned pointer. That this paper's <code>repo_n</code> reflects only the unresolvable GitHub mention and not the Zenodo citation suggests the Zenodo link is not being captured as a distinct <code>repo_url</code> for this paper at all, which merits separate investigation of <code>zenodo_links()</code> against this paper's extracted text; flagged as the same new <code>GITHUB_ORG_ROOT_NOT_A_SINGLE_REPO</code> category as 13766 for the GitHub portion of the gap.

References

Cooper, N., & the BES Data and Code Hackathon Group. (2026a). Data and code for MEE paper v1.0. Zenodo. <https://doi.org/10.5281/zenodo.19737911>

- Cooper, N., & the BES Data and Code Hackathon Group. (2026b). Data- and code-archiving in the British Ecological Society journals: Present status and recommendations for future improvements. *Methods in Ecology and Evolution*, 17, 1954–1966. <https://doi.org/10.1111/2041-210x.70338>
- DeBruine, L., Mesquida, C., Werner, J., & Lakens, D. (2026). Metacheck: Check research outputs for best practices. <https://github.com/scienceverse/metacheck>
- Riedel, N., Kip, M., & Bobrov, E. (2020). ODDPub – a Text-Mining Algorithm to Detect Data Sharing in Biomedical Publications. *Data Science Journal*, 19(1). <https://doi.org/10.5334/dsj-2020-042>
- Vilhuber, L., & Cavanagh, J. (2025). Report of the AEA Data Editor. *AEA Papers and Proceedings*, 115, 944–957. <https://doi.org/10.1257/pandp.115.944>
- Wilkinson, M. D., Dumontier, M., Aalbersberg, I. J., Appleton, G., Axton, M., Baak, A., Blomberg, N., Boiten, J.-W., Da Silva Santos, L. B., Bourne, P. E., Bouwman, J., Brookes, A. J., Clark, T., Crosas, M., Dillo, I., Dumon, O., Edmunds, S., Evelo, C. T., Finkers, R., ... Mons, B. (2016). The FAIR Guiding Principles for scientific data management and stewardship. *Scientific Data*, 3(1), 160018. <https://doi.org/10.1038/sdata.2016.18>